

# ECE 66100 Homework #6

by  
Adrien Dubois (dubois6@purdue.edu)

October 18, 2024

## Contents

<b>1 Theory Questions</b>	<b>1</b>
1.1 Question 1:	1
1.2 Question 2:	2
<b>2 Otsu's Algorithm:</b>	<b>2</b>
2.1 Implementation:	3
2.1.1 Otsu's Algorithm:	3
2.1.2 Calling coode:	3
<b>3 Texture Based Otsu's Algorithm:</b>	<b>4</b>
3.1 Explanation of the algorithm:	4
3.2 Implementation:	4
<b>4 Contouring:</b>	<b>5</b>
4.1 Contouring Implementation:	5
4.2 Closing Implementation:	5
<b>5 Results:</b>	<b>6</b>
5.1 Original Images:	6
5.2 RGB Otsu's Results:	6
5.2.1 Histograms:	6
5.2.2 Iteration Masks:	9
5.2.3 BGR Final Masks:	11
5.2.4 Segmentation Result:	12
5.2.5 Contouring Results:	13
5.2.6 What worked and what did not work:	15
5.3 Texture Based Otsu's Results:	16
5.3.1 Optimal hyperparamters:	16
5.3.2 Histograms:	16
5.3.3 Iteration Masks:	16
5.3.4 BGR Final Masks:	18
5.3.5 Final Contouring Result:	20
5.3.6 What worked and what did not work:	20
<b>6 Complete code printout:</b>	<b>21</b>

## 1 Theory Questions

### 1.1 Question 1:

*What are the strengths and weaknesses of Otsu's algorithm?* Otsu's algorithm has a few key strengths. First, it is quite a lightweight algorithm that runs quickly even without any GPU resources as would be

required for Neural Network implementations. This makes it optimal for realtime applications, where the low latency is necessary for efficient runtimes. Additionally, it is an unsupervised method for segmentation which means that it does not require any labeled training data.

One of the weaknesses of Otsu's algorithm is that it only works based off of the separate RGB channels. This reduces the accuracy of the segmentation, and also makes it so that users needed to pre-determine which thresholded class they want to keep depending on whether or not the foreground object is bright (keep class 1 which is closer to white pixels), or the foreground object is dark (keep class 0 which is closer to black pixels). Additionally, Otsu's algorithm is very sensitive to noise between the foreground and background objects. In this way, since it works on the brightness of the RGB channels, if the whole image is muddied/dark, Otsu's algorithm would likely fail to segment the foreground object. Lastly, finding a global threshold for the whole image could introduce performance losses.

## 1.2 Question 2:

*What are the strengths and weaknesses of the Watershed algorithm?* Similar to Otsu's algorithm, one of the strengths of the Watershed algorithm is that it is an unsupervised method for segmentation and therefore does not require a labeled dataset. Additionally, unlike Otsu's algorithm, the Watershed algorithm performs very well for complex objects, even when boundaries between objects overlap. It also makes no assumption on the pixel value distribution which makes it much more versatile than Otsu. Some of the weaknesses of the Watershed algorithm is that it is still much more sensitive to noise compared to Neural Network based approaches. Additionally, it tends to oversegment images with multiple segmented “fragments” in the image instead of just extracting the foreground. Lastly, the Watershed algorithm is much more computationally expensive compared to other methods while also having much lower performance than simple NN approaches.

## 2 Otsu's Algorithm:

Otsu's algorithm works by creating a histogram of all the pixel values in the image. It then finds a threshold value, that separates the histogram into two different classes (foreground and background) by maximizing the between class variance at a certain threshold. The following equations are used to calculate the between class variance:

$$p(i) = \frac{n_i}{N}$$

where:

- $n_i$  is the number of points with pixel value = i
- N is the total number of points in the image

We then calculate the probabilities of picking a point from each class:

$$\begin{aligned} \omega_0(t) &= \sum_{i=0}^t p(i) \\ \omega_1(t) &= \sum_{i=t}^N p(i) \end{aligned}$$

where:

- t = the current test value for the threshold

We can then calculate the class means as follows:

$$\begin{aligned} \mu_0 &= \frac{\sum_{i=0}^t i \times p(i)}{\omega_0} \\ \mu_1 &= \frac{\sum_{i=t}^N i \times p(i)}{\omega_1} \end{aligned}$$

Lastly, we calculate the between class variance:

$$\omega_B = \omega_0 \omega_1 (\mu_1 - \mu_0)^2$$

After calculating the optimal threshold as the value of t that maximizes  $\omega_B$ , we threshold the image as follows: For all (i,j), we fill in an image canvas using this rule:

- If pixel value as  $(i,j) > \text{threshold} \rightarrow 1$
- If pixel value as  $(i,j) \leq \text{threshold} \rightarrow 0$

This calculation is performed if we assume that the foreground pixel is a dark object. On the other hand, if the foreground object is brighter than the rest of the image, we need to threshold in the opposite direction by keeping class 1.

This algorithm was run separately for each of the RGB channels of the image. Lastly, I ran this in an iterative manner, where I would re-run Otsu's algorithm on only the image pixels that remain after applying the segmentation mask on the original image. It is important to note that the re-runs should be run on the remaining pixels, and not on the thresholded image as that would increase the number of 0 valued pixels which would affect the histogram pixel distribution and therefore the between class variance calculation.

## 2.1 Implementation:

### 2.1.1 Otsu's Algorithm:

```

1 # RGB Segmentation
2 def get_Otsus_segmentation(img_channel, img_name, iter):
3     num_bins = int(np.max(img_channel) - np.min(img_channel) + 1)
4     hist, bin_edges = np.histogram(img_channel, bins=num_bins)
5
6     # Get histogram probabilities
7     p_i = hist / sum(hist)
8
9     # List to save between class variances
10    variances = np.zeros((len(bin_edges)))
11
12    # Go through all threshold ts and find the one that maximizes the fisher
13    # discriminant
14    # I only work with the pixels left from the previous iteration of Otsu's algorithm.
15    # So, my loop for the t values can only get larger than the previous one
16    for t in range(len(bin_edges)):
17        # Probability of picking points from each class = sum(p(i))
18        w0 = np.sum(p_i[:t])
19        w1 = np.sum(p_i[t:])
20
21        if w0 == 0 or w1 == 0:
22            continue
23
24        # Mean levels for each class = sum(i * p(i))
25        i = np.arange(len(p_i))
26        mu0 = np.sum(i[:t] * p_i[:t])/w0
27        mu1 = np.sum(i[t:] * p_i[t:])/w1
28
29        # Between class varianced
30        omegaB = w0 * w1 * (mu1 - mu0)**2
31
32        # Append fisher's discriminant to the variances list
33        variances[t] = omegaB
34
35        # Save a graph of the histogram with the variances at each t plotted on top
36        plt.figure(figsize=(10, 6))
37        plt.bar(bin_edges[:-1], hist, width=np.diff(bin_edges), edgecolor="black", align="edge",
38                label="Pixel Value Histogram")
39        plt.plot(variances, color='r', label="Between class variances at each t")
40        plt.legend(loc='upper right')
41        plt.savefig(histogram_paths + img_name + "_hist_" + str(iter) + ".jpg", format="jpg")
42
43        # Return the t that maximized fisher's ratio
44        return np.argmax(variances)

```

### 2.1.2 Calling coode:

```

1 # Otsu's Algorithm for RGB images:
2 # Images are in the following order: Climb, Dog, Flower, RVL
3 iters = [[3,3,4],
4           [3,3,3],
5           [1,1,1],
6           [2,2,2]]
7 invert = [False, True, False, False]
8 for i, img_path in enumerate(img_paths):
9     img_name = img_path.split("/")[-1][-4]
10    run_Otsus_algorithm_iterative(iters=iters[i], save_path=rgb_path, \
11        invert=invert[i], img_path=img_path, img=None, img_name=img_name, \
12        histogram_folder="RGB/")

```

### 3 Texture Based Otsu's Algorithm:

#### 3.1 Explanation of the algorithm:

On the other hand, for the texture based implementation I run a windowing algorithm with a window size of N.

To do so, I first pad the image with 0s using a pad width of N. I can then use numpy's sliding window view method to return an array for each window of shape  $(2N+1, 2N+1)$ .

Then, I subtract the mean window pixel from all the window pixels and calculate the within window variance, normalized from a 0 to 255. I run this algorithm for three values of N, where lower values produced more fine-grain results compared larger values of N. By concatenating these texture maps together into the 3 channels of an image canvas, I can then run my Otsu's algorithm on the textured image data and produce much better contouring results for complex images. The only place where this did not work was for the dog image where the grass textures were taken as contours of foreground objects which produced a lot of noise in my output.

#### 3.2 Implementation:

```

1 # Texture Segmentation
2 def get_texture_otsu(grey_img, N):
3     # Create canvas for output image
4     output_channel = np.zeros_like(grey_img)
5
6     # Pad the image with extra 0s as necessary
7     padded_img = np.pad(grey_img, pad_width=N, mode="constant", constant_values=0)
8
9     # Create an N by N window of the nearby pixels with 0 padding
10    # Instead of using a double for loop which is slow, I do this through numpy
11    # vectorized functions
12    # This will return a numpy array of shape: (H, W, 2N + 1, 2N + 1) (list of windows)
13    # We do 2N + 1 since we want N on each direction from the input pixel
14    window_shape = (2*N + 1, 2*N + 1)
15    windows = np.lib.stride_tricks.sliding_window_view(padded_img, window_shape)
16
17    # Subtract the window's mean value
18    window_means = np.mean(windows, axis=(-2, -1))
19    windows = windows - window_means[:, :, None, None]
20
21    # Compute the variance within each window
22    window_variances = np.var(windows, axis=(-2, -1))
23
24    # Normalize the variance to 0-1 range and scale to 255 for BW image
25    output_channel = (window_variances - window_variances.min()) / (window_variances.max()
26        - window_variances.min()) * 255
27
28    # Return a new array with the within pixel variance as the center pixel value
29    # These values are normalized to 0-255
30    return output_channel.astype(np.uint8)

```

## 4 Contouring:

While the Texture based Otsu algorithm directly creates contour maps, we need to do some post-processing on the RGB Ostu results to extract the contour information. To do so, I create 3 by 3 windows for all points of the image. I then check for two different masks:

1. Is the center pixel value equal to 1?
2. Are there any 0s in the window?

Performing a logical and computation on these two masks allows for an extremely fast implementation of a contouring algorithm to return border pixels.

### 4.1 Contouring Implementation:

```
1 def get_contours(img):
2     window_shape = (3,3)
3     windows = np.lib.stride_tricks.sliding_window_view(img, window_shape)
4
5     # We only want to look at windows where the center pixel is 1:
6     center_pixels = windows[:, :, 1, 1]
7     center_pixel_mask = center_pixels == 1
8
9     # Create a mask that will look in each window and check for a 0 pixel
10    # If there is a 0, it will return a 1, otherwise return a 0
11    zero_mask = np.any(windows == 0, axis=(2, 3)).astype(int)
12
13    final_mask = np.logical_and(center_pixel_mask, zero_mask)
14
15    # Convert to grey scale (0->255)
16    return final_mask * 255
```

However, since the point clouds returned by my contouring algorithm were pretty sparse, I also employed a “closing” algorithm to clean up the lines. This involved 3 iterations of dilation which grows the pixel clouds and 1 iteration of erosion.

### 4.2 Closing Implementation:

```
1 def run_closing(contour, kernel_size):
2     # Closing is dilation followed by erosion
3     kernel = np.ones((kernel_size, kernel_size), dtype=np.uint8)
4
5     contour = cv2.dilate(contour.astype(np.float32), kernel, iterations=3)
6     contour = cv2.erode(contour.astype(np.float32), kernel, iterations=1)
7     return contour
```

## 5 Results:

### 5.1 Original Images:

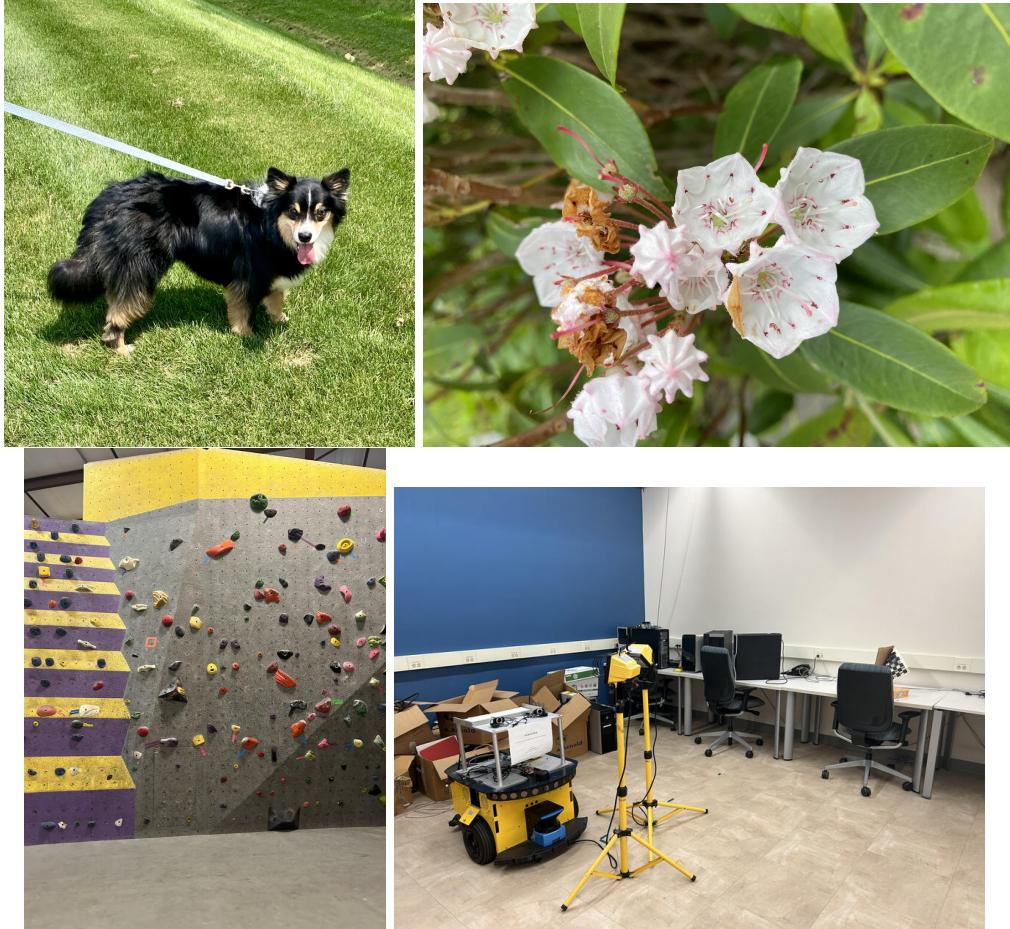


Figure 1: The original images I am performing the segmentation on.

For the climbing image, my goal was to have the colored climbing holds be the foreground and the background be the climbing wall. On the other hand, for the office image, my goal was to have the desk spaces and yellow light be the foreground and the background be the floor and walls of the room.

### 5.2 RGB Otsu's Results:

For the RGB images, I required an iterative approach to get optimal results. These were the hyperparameters that produced the best results:

Image Title	Green Iters	Blue Iters	Red Iters	Which class is kept
Dog	3	3	4	Class 1
Climb	3	3	3	Class 0
Flower	1	1	1	Class 0
Office	2	2	2	Class 0

Table 1: Optimal hyperparameters for the RGB Otsu segmentation.

#### 5.2.1 Histograms:

Included below are the pixel value histograms, along with the variance at every threshold value. I only include examples for the dog image, and my climbing wall image for this section to demonstrate results when thresholding on class 0, and class 1.

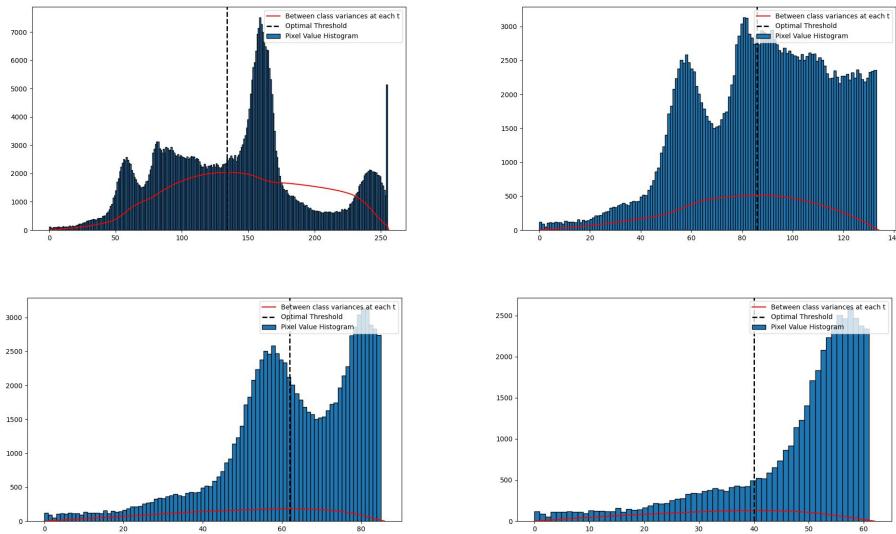


Figure 2: The histograms and variances at each threshold at iteration 0, 1, 2 and 3 for the climbing wall image.

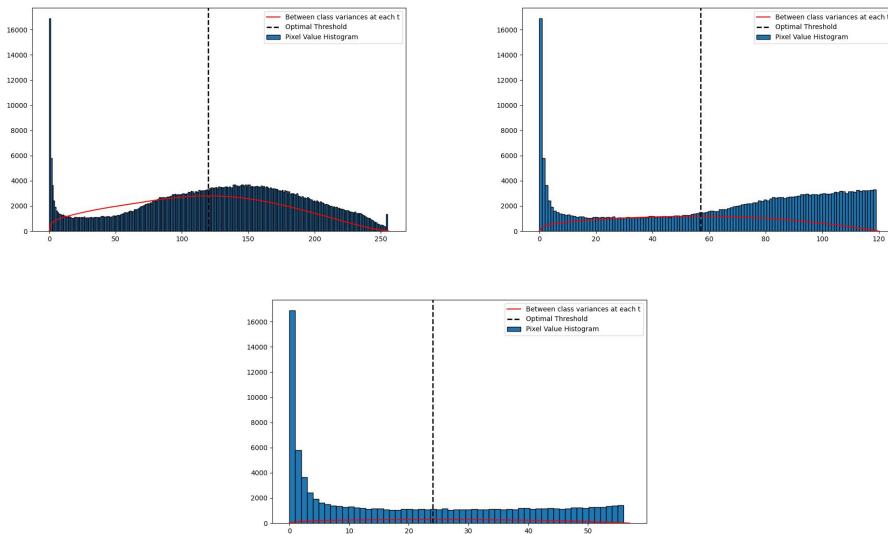


Figure 3: The histograms and variances at each threshold at iteration 0, 1, and 2 for the dog wall image.



### 5.2.2 Iteration Masks:

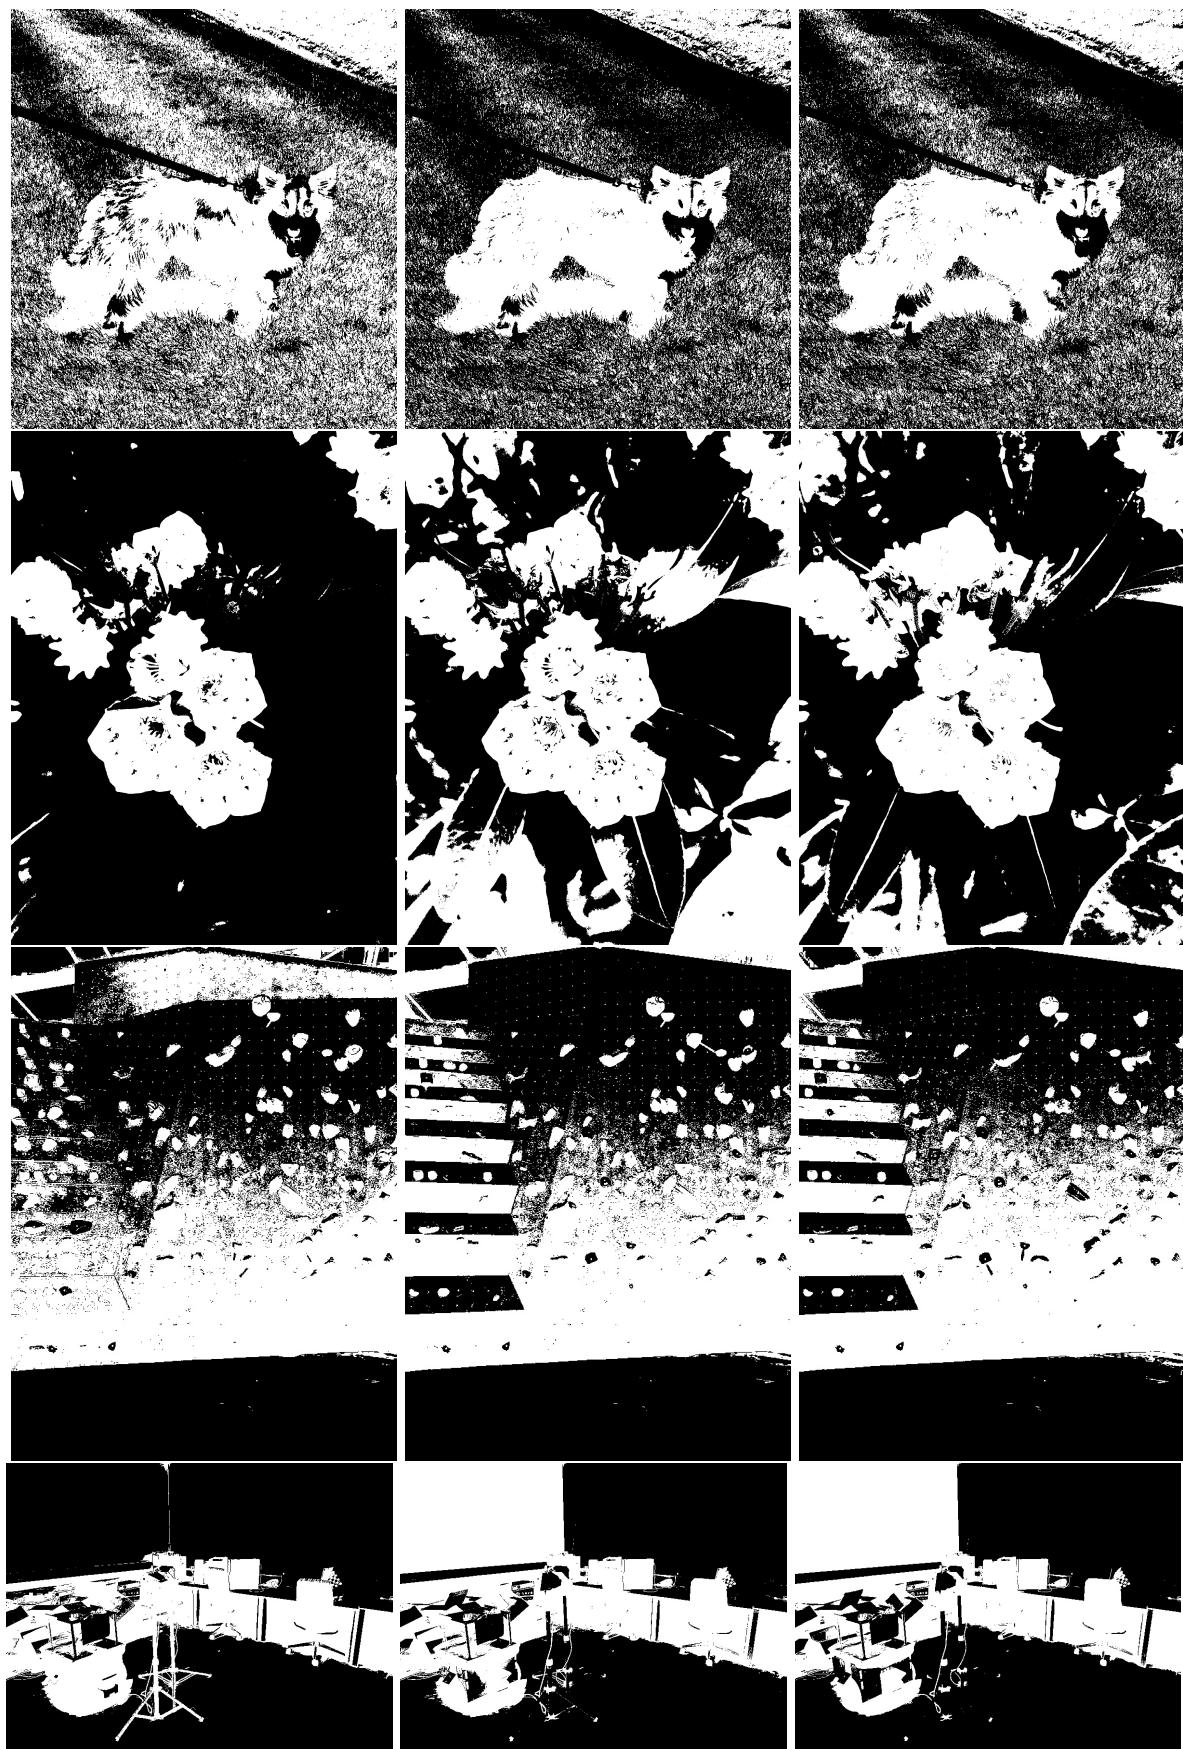


Figure 4: The masks for the first iteration for all images. The images are included in BGR ordering.

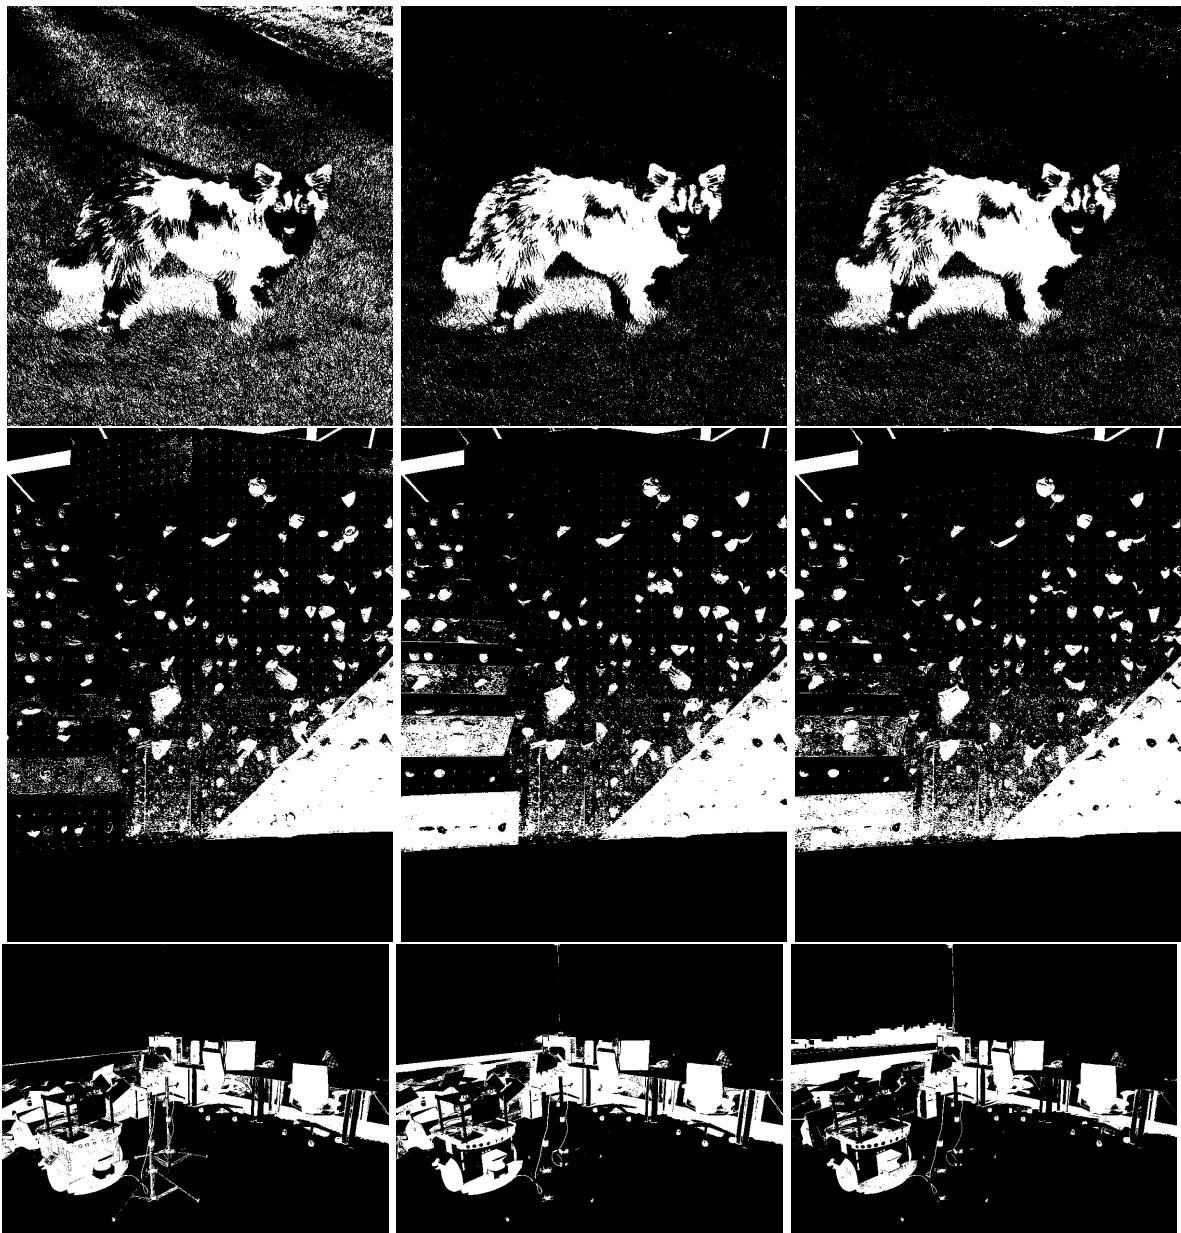


Figure 5: The masks for the second iteration for the images. The images are included in BGR ordering.

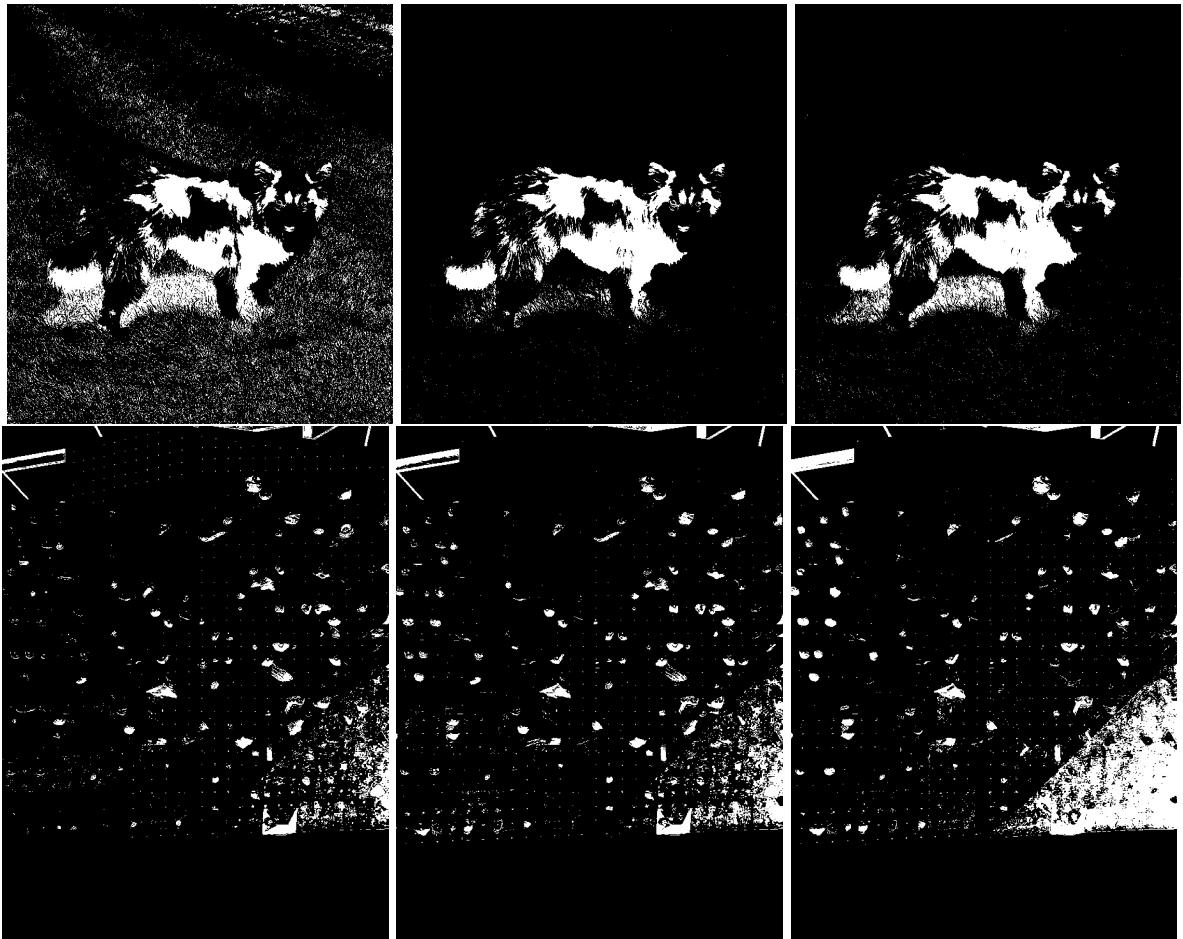


Figure 6: The masks for the third iteration for the images. The images are included in BGR ordering.

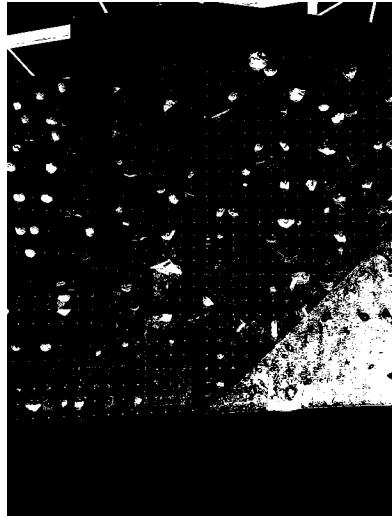


Figure 7: The masks for the third iteration for the climbing image's red channel.

### 5.2.3 BGR Final Masks:

For the following image, I plot the final graph but instead of just created a logical and of the black and white masks, I take the final mask for each channel and re-plot it in BGR format. Therefore, a white pixel in channel 1's output would result in a blue pixel in this image, while a white pixel in channel 2's output mask would result in a green pixel etc. This lets us understand which components of the image

are segmented from which channel through the thresholding.

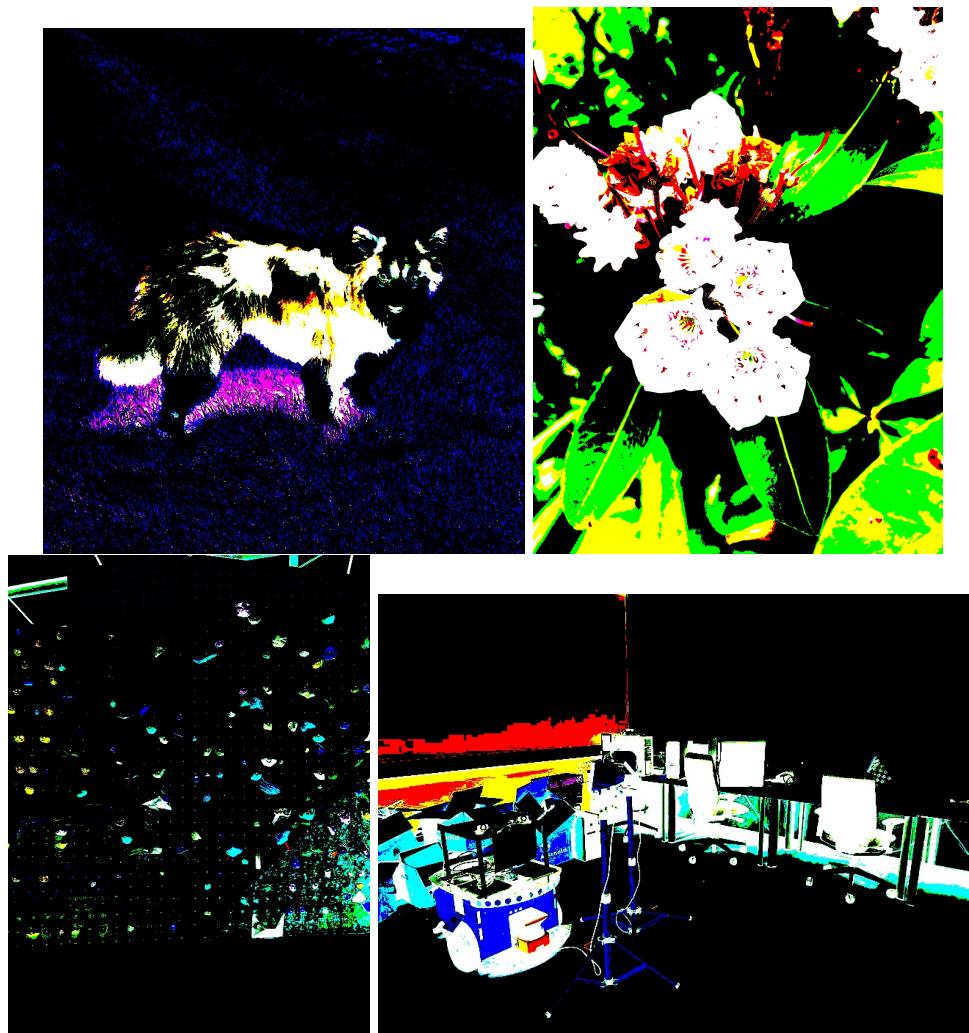


Figure 8: The final BGR channel masks for each image.

#### 5.2.4 Segmentation Result:

Lastly, I include below the final Segmentation results for the RGB based Otsu segmentation.

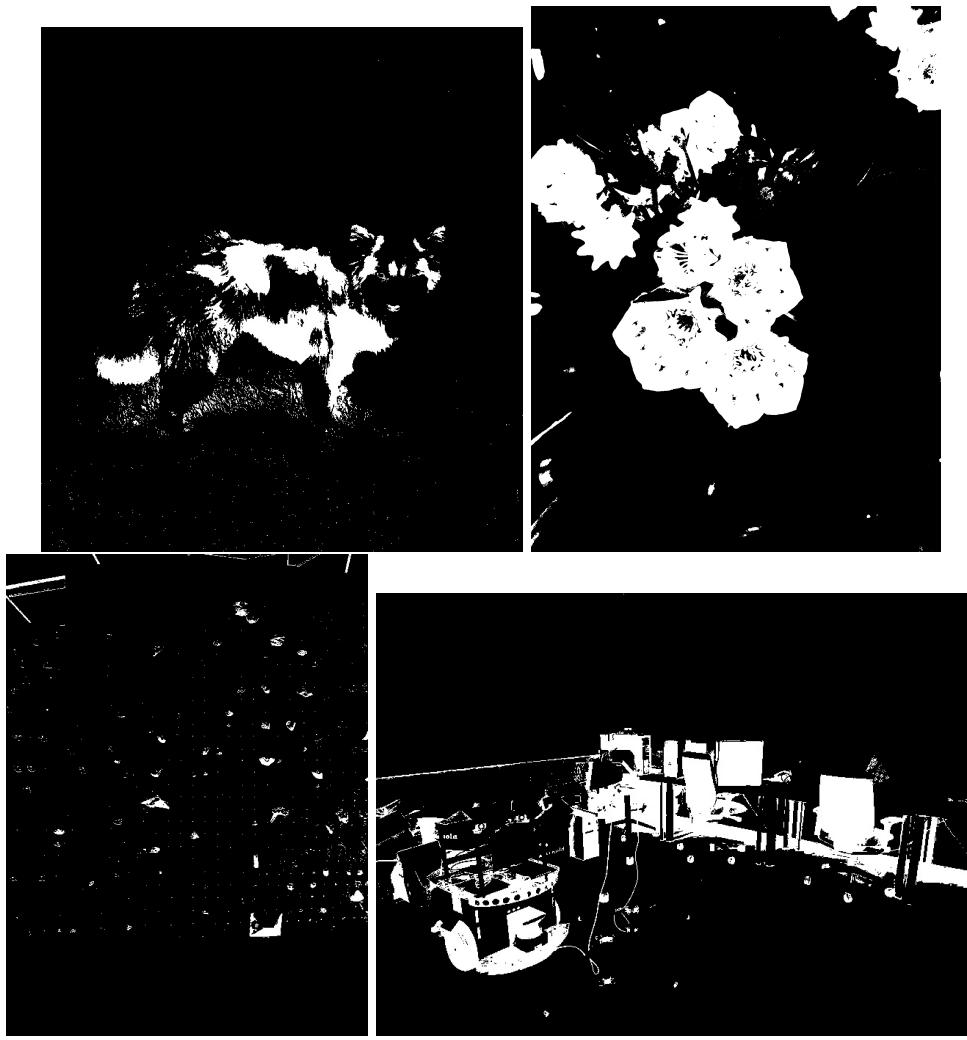


Figure 9: The final segmentation masks for each image with RGB Otsu.

#### 5.2.5 Contouring Results:

The following are the results of the contouring pipeline on the RGB Otsu's algorithm. These take the segmentation masks and return a point cloud that represents the border pixels of the image. The first set of results do not include any post-processing.

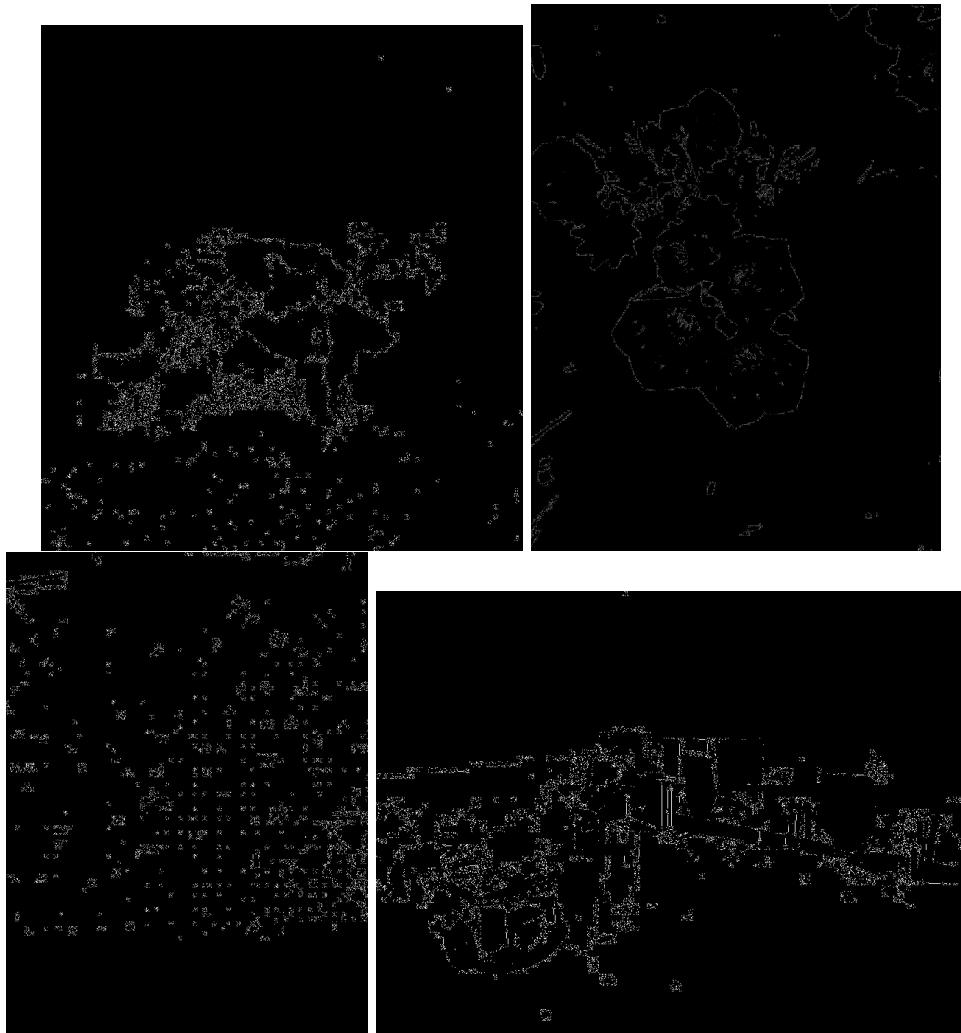


Figure 10: The image contours from the RGB Otsu's segmenation mask without any post-processing. The images are included in the following order: Dog, Flower, Climbing wall and Office space.

However, since these points clouds are pretty sparse, I decided to run erosion and dilation on the images to improve the results into continuous line segments. In this way, I ran a “closing” algorithm on the generated contours that involved 3 iterations of dilation, followed by 1 iteration of erosion which produced the best results. However, while this worked well for the flower image, you can see that for fine-grain details as seen in the climbing wall the dilation procedure overlapped with one another and created within-contour artifacts by filling in the space.

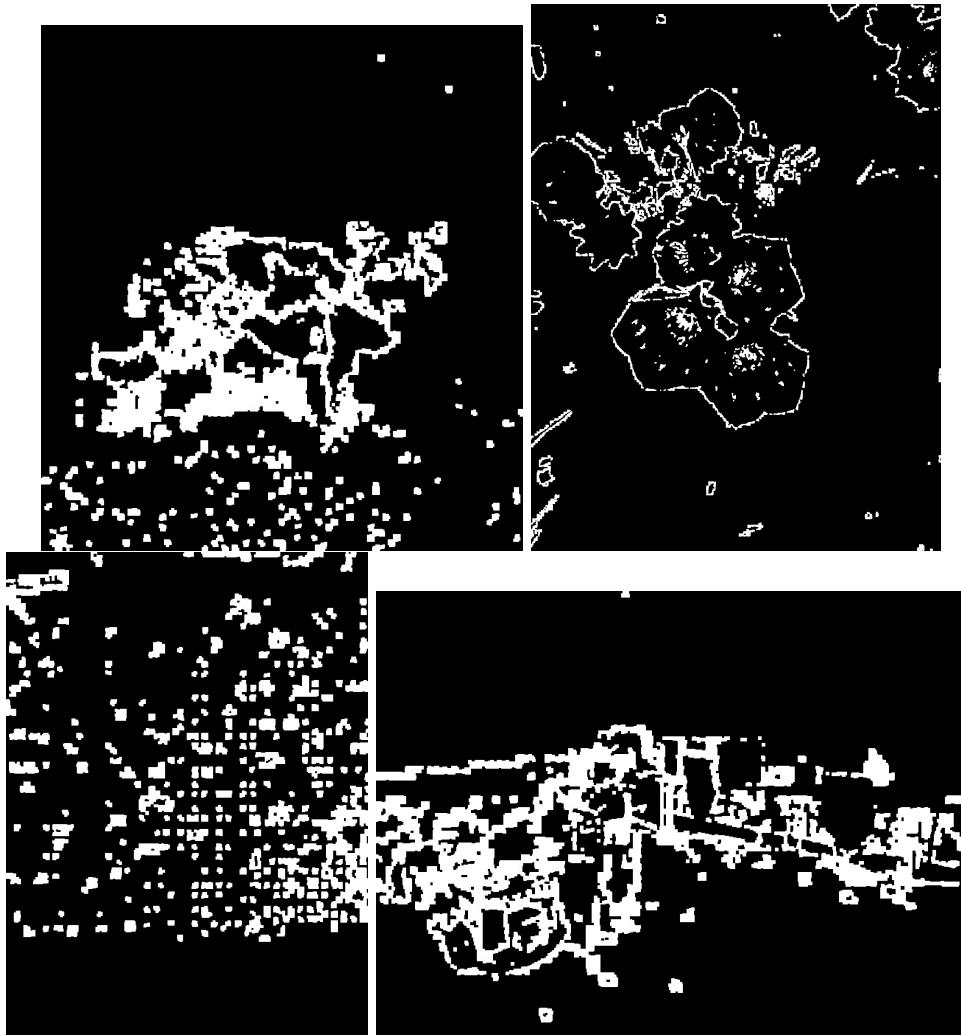


Figure 11: The image contours from the RGB Otsu's segmenation mask without any post-processing. The images are included in the following order: Dog, Flower, Climbing wall and Office space.

### 5.2.6 What worked and what did not work:

The best result was on the flower image where there was a very large color difference between the white flower and the green background leafs. On the other hand, the dog was especially hard to segment due to the large color differences within its coat: I could get the black portions well segemented, but that meant that I would miss the lighter brown areas. This is due to the fact that we are creating a global filter based on the intensity of the RGB values. I think it would be interesting to work on HSV or other color spaces instead of RGB and see if they algorithm performs there too. Working in that way could help to segment more holds on the climbing wall that were not pure red/green/blue. Working based on color saturation would have especially improved this image's result due to the bright color of the holds against the grey background wall. On the other hand, the office space segmentation worked very well but it oversegmented the task. In this way, it also captured the area beloow the desks instead of just the main objects.

As was previously stated, the contouring results are very promising, especially for the flower shape. The climbing wall contour suffered from lots of artifact point clouds from the bolt holes which greatly decreased the performance of the “closing” operation. This issue was also noticed to a lesser degree by the grass textures in the dog image.

### 5.3 Texture Based Otsu's Results:

#### 5.3.1 Optimal hyperparameters:

Image Title	N values	Channel 1 ( $C_1$ ) Iters	$C_2$ Iters	$C_3$ Iters	Class kept
Dog	[5, 4, 3]	1	1	1	Class 1
Climb	[3, 5, 7]	1	1	1	Class 1
Flower	[3, 5, 7]	2	2	2	Class 1
Office	[1, 2, 3]	3	2	2	Class 1

Table 2: Optimal hyperparameters for the RGB Otsu segmentation.

#### 5.3.2 Histograms:

Included below are the histograms for all images along with the chosen threshold.

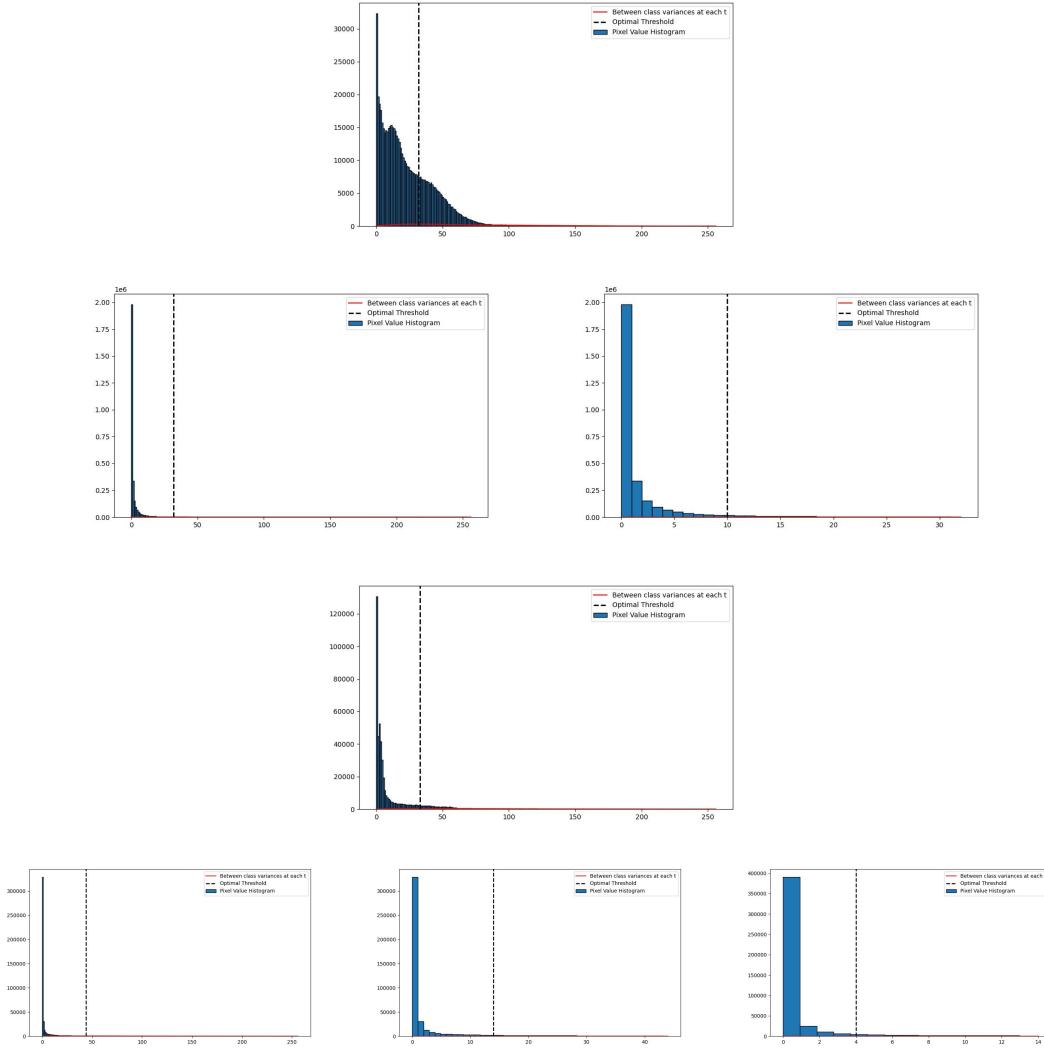


Figure 12: The histograms for all images, at all iterations. Histograms in the same row belong to the same image but at different iterations of Otsu's Algorithm. The columns are in the following order: Dog, Flower, Climbing wall and Office space.

#### 5.3.3 Iteration Masks:

Included below are the masks for each channel, at each iteration step of Otsu's Algorithm after running the texture windowing.

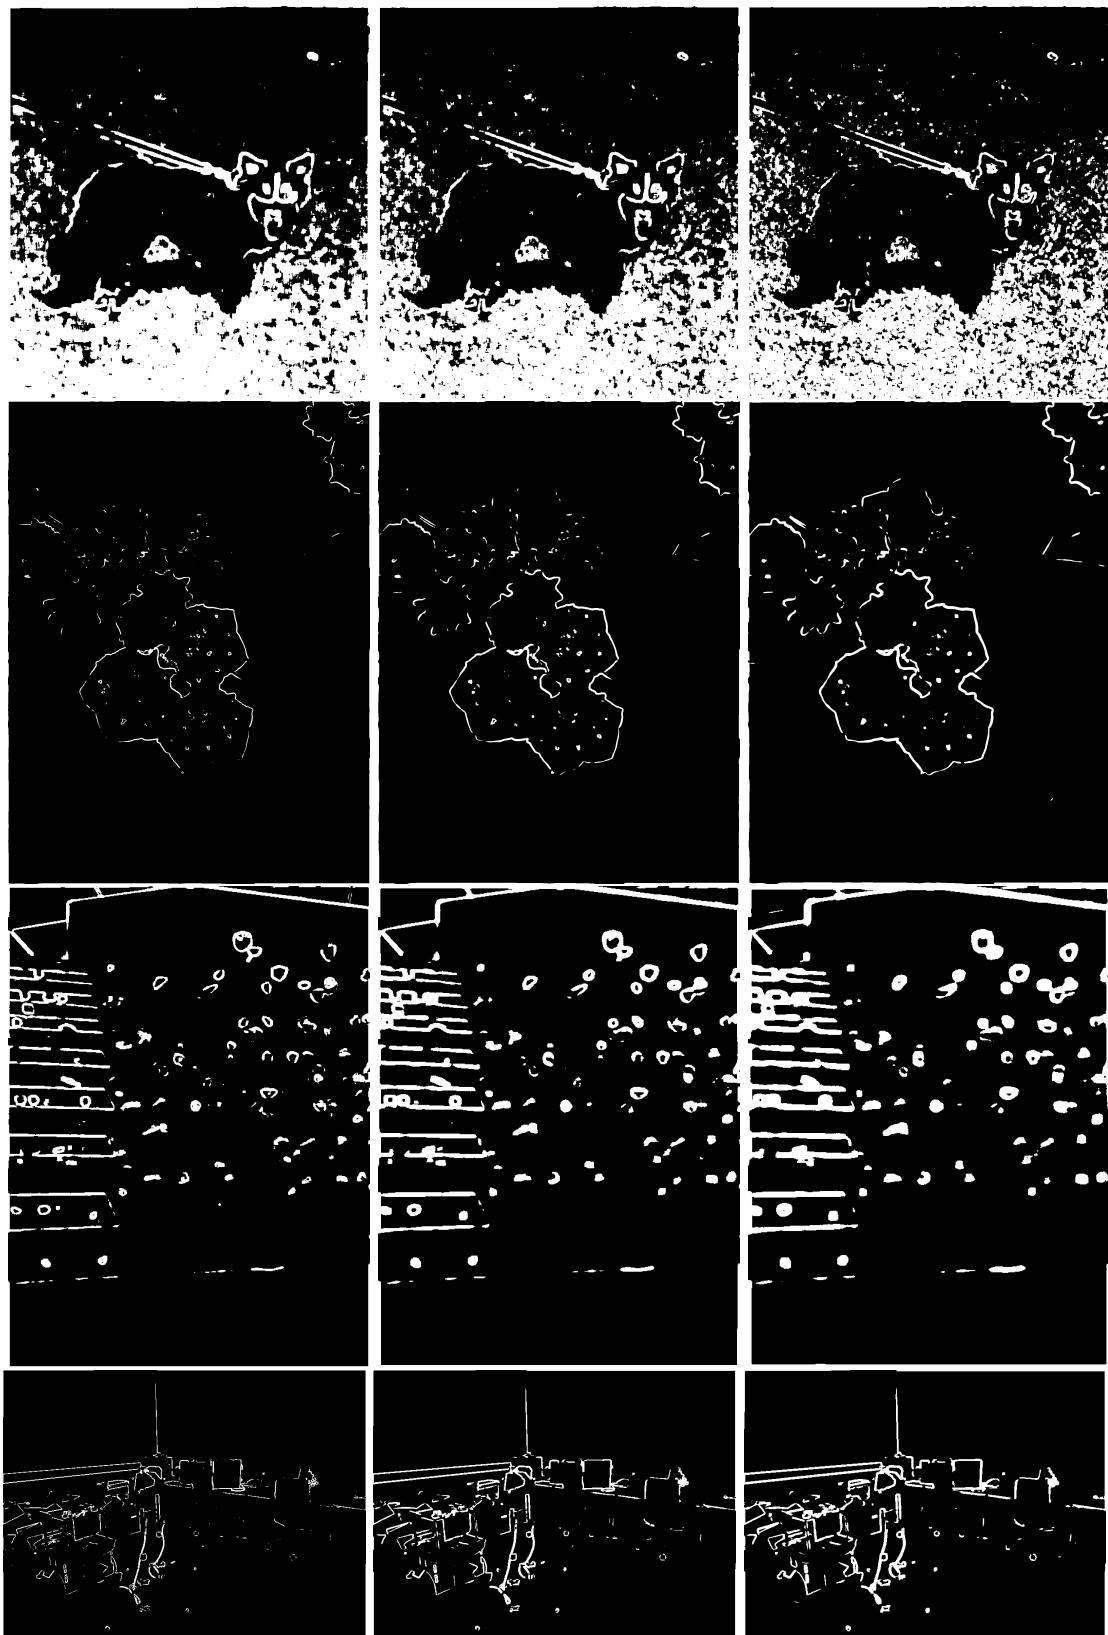


Figure 13: The results from the first iteration of Otsu’s algorithm on the outputs of the Texture Windowing method. The images are in the following order: each column denotes Channel 1, Channel 2, Channel 3 as defined by the N parameter list, and the rows are Dog, Flower, Climbing wall and lastly the office space.

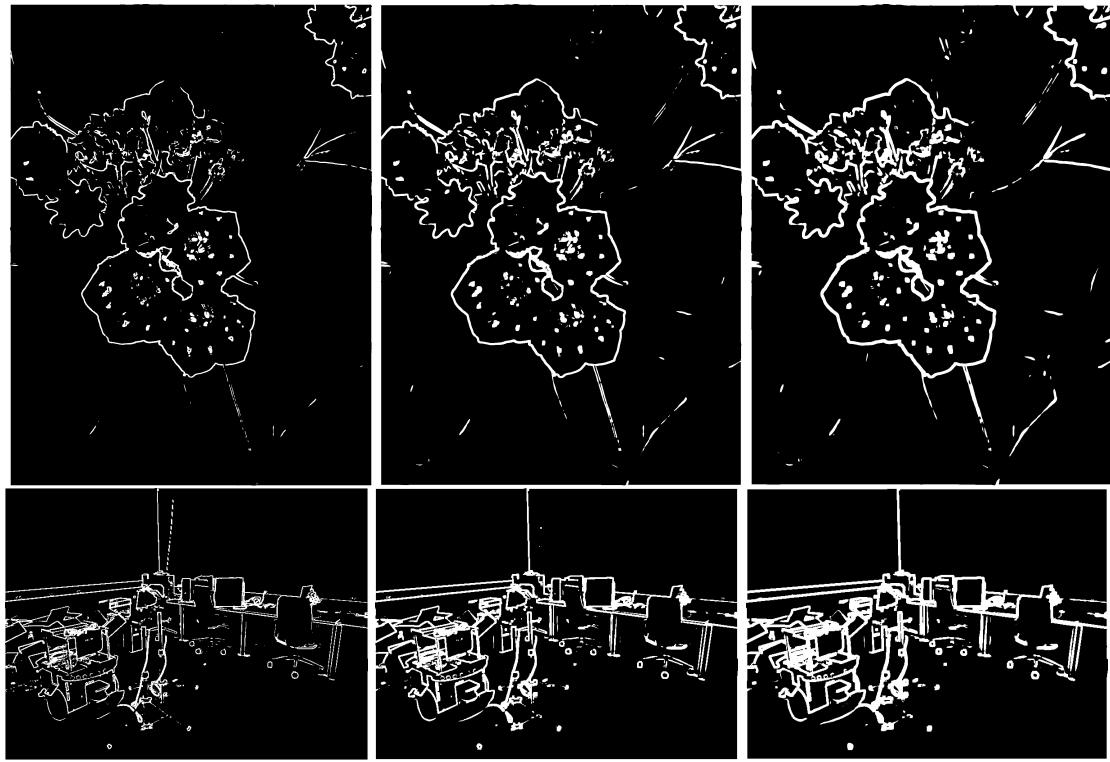


Figure 14: The results from the second iteration of Otsu’s algorithm on the outputs of the Texture Windowing method. The images are in the following order: each column denotes Channel 1, Channel 2, Channel 3 as defined by the N parameter list, and the rows are flower and the office space.

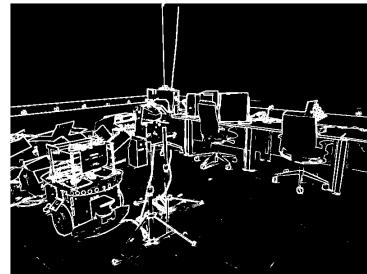


Figure 15: The results from the third iteration of Otsu’s algorithm on the outputs of the Texture Windowing method. This image is from the first channel of the office space image.

#### 5.3.4 BGR Final Masks:

For the following graphs, I take each channel’s mask output and graph them as RGB values. Therefore, a white pixel from the mask in channel 1 would show up as pure blue, a white pixel from channel 2 would be pure green etc. This allows for some intuition on which channel is responsible for extracting which portion of the information.

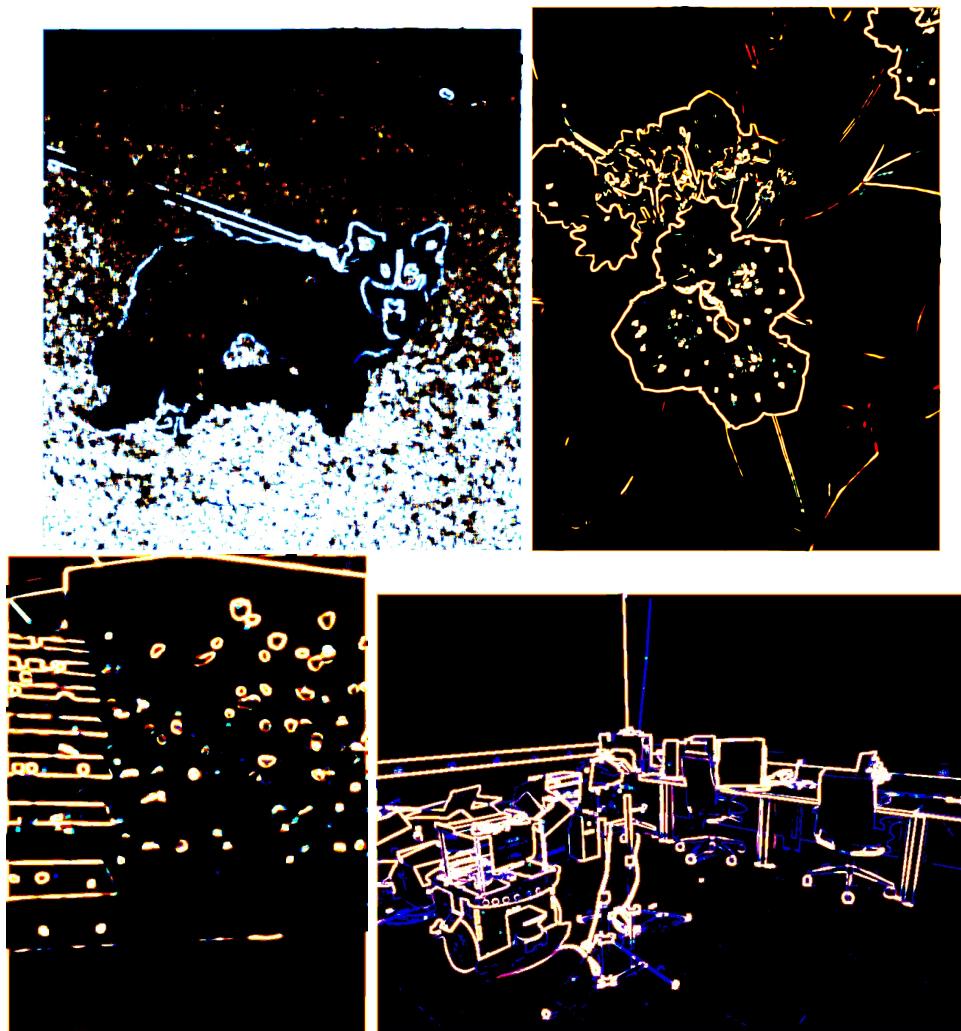


Figure 16: The BGR space outputs for the Texture based contouring. The results are in the following order: Dog, Flower, Climbing wall and Office space.

### 5.3.5 Final Contouring Result:

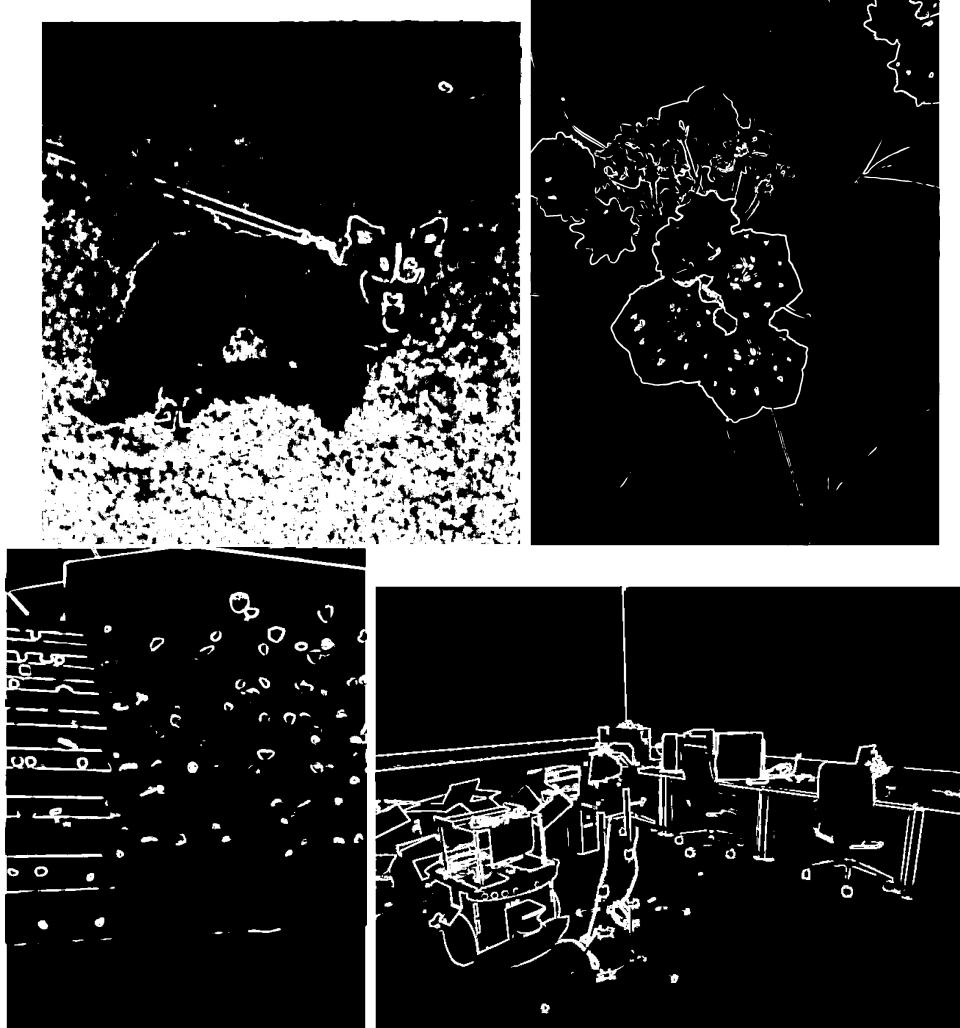


Figure 17: The final outputs for the Texture based contouring. The results are in the following order: Dog, Flower, Climbing wall and Office space.

### 5.3.6 What worked and what did not work:

To summarize the results for the texture based algorithm, I noticed a significant improvement in the contour output from this implementation compared to the previous implementation based on the RGB values following by contour extraction. This improvement is especially noticeable for the contour around the office space image which is very clear with this algorithm, but did not work on the RGB based method. I believe that the windowing effect based on variance was much more effective to determine boundary points compared to my previous contour implementation that naively looks for points on boundaries through the inclusion of both 1 and 0 pixels. The climbing wall also had much better performance, though there are larger artifacts in the horizontal edges on the left side of the wall that were not present in the RGB image. These lines are between very large shifts in color; however, they are not part of the actual climbing holds so an optimal algorithm would not segment them out. Lastly, the dog image was much more difficult to get accurate contour maps using this algorithm due to the fine detail in the grass textures. In this way, the algorithm would extract all of the grass textures before finding the contour of the dog due to the fine-grain differences in pixel values from one blade of grass to another. It may be necessary to somehow mask out these known green background first before running this algorithm and see how the performance improves.

## 6 Complete code printout:

```
1 # Imports
2 import cv2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import os
6
7 contour_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/Output_Images/
    Contour_Extraction/"
8 rgb_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/Output_Images/RGB_Otsus/"
9 texture_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/Output_Images/
    Texture_Segmentation/"
10 histogram_paths = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/Histograms/"
11
12 # Split images into RGB:
13 def get_bgr_split_images(img_path):
14     img = cv2.imread(img_path)
15     return cv2.split(img)
16
17 # Display image Object:
18 def display_img_object(img, isOpenCV=True):
19     if isOpenCV:
20         img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
21     plt.imshow(img)
22     plt.show()
23
24 # RGB Segmentation
25 def get_Otsus_segmentation(img_channel, img_name, iter, subfolder):
26     num_bins = int(np.max(img_channel) - np.min(img_channel) + 1)
27     hist, bin_edges = np.histogram(img_channel, bins=num_bins)
28
29     # Get histogram probabilities
30     p_i = hist / sum(hist)
31
32     # List to save between class variances
33     variances = np.zeros((len(bin_edges)))
34
35     # Go through all threshold ts and find the one that maximizes the fisher
36     # discriminant
37     # I only work with the pixels left from the previous iteration of Otsu's algorithm.
38     # So, my loop for the t values can only get larger than the previous one
39     for t in range(len(bin_edges)):
40         # Probability of picking points from each class = sum(p(i))
41         w0 = np.sum(p_i[:t])
42         w1 = np.sum(p_i[t:])
43
44         if w0 == 0 or w1 == 0:
45             continue
46
47         # Mean levels for each class = sum(i * p(i))
48         i = np.arange(len(p_i))
49         mu0 = np.sum(i[:t] * p_i[:t])/w0
50         mu1 = np.sum(i[t:] * p_i[t:])/w1
51
52         # Between class varianced
53         omegaB = w0 * w1 * (mu1 - mu0)**2
54
55         # Append fisher's discriminant to the variances list
56         variances[t] = omegaB
57
58     optimal_threshold = np.argmax(variances)
59
60     # Save a graph of the histogram with the variances at each t plotted on top
61     plt.figure(figsize=(10, 6))
62     plt.bar(bin_edges[:-1], hist, width=np.diff(bin_edges), edgecolor="black", align="edge",
63             label="Pixel Value Histogram")
64     plt.plot(variances, color='r', label="Between class variances at each t")
65     plt.axvline(x=optimal_threshold, color="k", linestyle="--", linewidth=2, label="Optimal Threshold")
66     plt.legend(loc='upper right')
67     plt.savefig(histogram_paths + subfolder + img_name + "_hist_" + str(iter) + ".jpg",
68                 format="jpg")
```

```

66     plt.close()
67
68     # Return the t that maximized fisher's ratio
69     return optimal_threshold
70
71
72 # Texture Segmentation
73 def get_texture_otsu(grey_img, N):
74     # Create canvas for output image
75     output_channel = np.zeros_like(grey_img)
76
77     # Pad the image with extra 0s as necessary
78     padded_img = np.pad(grey_img, pad_width=N, mode="constant", constant_values=0)
79
80     # Create an N by N window of the nearby pixels with 0 padding
81     # Instead of using a double for loop which is slow, I do this through numpy
82     # vectorized functions
83     # This will return a numpy array of shape: (H, W, 2N + 1, 2N + 1) (list of windows)
84     # We do 2N + 1 since we want N on each direction from the input pixel
85     window_shape = (2*N + 1, 2*N + 1)
86     windows = np.lib.stride_tricks.sliding_window_view(padded_img, window_shape)
87
88     # Subtract the window's mean value
89     window_means = np.mean(windows, axis=(-2, -1))
90     windows = windows - window_means[:, :, None, None]
91
92     # Compute the variance within each window
93     window_variances = np.var(windows, axis=(-2, -1))
94
95     # Normalize the variance to 0-1 range and scale to 255 for BW image
96     output_channel = (window_variances - window_variances.min()) / (window_variances.max()
97     () - window_variances.min()) * 255
98
99     # Return a new array with the within pixel variance as the center pixel value
100    # These values are normalized to 0-255
101    return output_channel.astype(np.uint8)
102
103
104 def get_contours(img):
105     window_shape = (3,3)
106     windows = np.lib.stride_tricks.sliding_window_view(img, window_shape)
107
108     # We only want to look at windows where the center pixel is 1:
109     center_pixels = windows[:, :, 1, 1]
110     center_pixel_mask = center_pixels == 1
111
112     # Create a mask that will look in each window and check for a 0 pixel
113     # If there is a 0, it will return a 1, otherwise return a 0
114     zero_mask = np.any(windows == 0, axis=(2, 3)).astype(int)
115
116     final_mask = np.logical_and(center_pixel_mask, zero_mask)
117
118     # Convert to grey scale (0->255)
119     return final_mask * 255
120
121
122 def run_closing(contour, kernel_size):
123     # Closing is dillation followed by erosion
124     kernel = np.ones((kernel_size, kernel_size), dtype=np.uint8)
125
126     contour = cv2.dilate(contour.astype(np.float32), kernel, iterations=3)
127     contour = cv2.erode(contour.astype(np.float32), kernel, iterations=1)
128
129     return contour
130
131
132 def run_Otsus_algorithm_iterative(iters, save_path, invert, img_path=None, img=None,
133     img_name="", histogram_folder=""):
134     """
135         Runs an iterative version of Otsus algorithm. Will print the following results:
136             - Separate masks for each channel dimension
137             - A joined mask highlighting the RGB extractions for each mask
138             - The histogram with variance values at each tested threshold for Otsu's
139             algorithm.
140
141     Args:

```

```

135     iters (int): number of iters to run Otsus
136     save_path (string): path pointing to where you want the images to be saved
137     img_path (string, optional): Path to the image. Defaults to None.
138     img (OpenCV Image Object, optional): OpenCV BGR Image. Defaults to None.
139     img_name (str, optional): NameID for image printouts. Defaults to "".
140
141 """
142     print("Working on the following image:", img_name)
143     if img_path != None:
144         channels = get_bgr_split_images(img_path)
145     else:
146         channels = cv2.split(img)
147
148     final_rgb_img = np.zeros((channels[0].shape[0], channels[0].shape[1], 3), dtype=np.
149     uint8)
150     final_mask = np.ones_like(channels[0], dtype=np.uint8)
151
152     for i, channel in enumerate(channels):
153         flattened_channel = channel.flatten()
154         # Channel I am working on:
155         channel_name = ["_blue", "_green", "_red"][i]
156         print("Working on the following channel: ", channel_name)
157
158         for iter in range(iters[i]):
159             # As I run Otsu's algorithm, more of my pixels become 0 due to the masking.
160             optimal_threshold = get_Otsus_segmentation(flattened_channel, img_name, iter,
161             histogram_folder)
162
163             # Print the optimal threshold:
164             print("Optimal Threshold: ", optimal_threshold)
165
166             # If the values in the channel are less than the threshold -> replace it with
167             # 0 in the bw_img or do the opposite.
168             # This choice can be made depending on whether the foreground is a bright
169             # object, or a dark one.
170             # If foreground is closer to white, you want to keep Class 1 (brighter
171             # pixels), else keep class 0 (darker pixels)
172             if not invert:
173                 bw_img = np.zeros(channel.shape, dtype=np.uint8)
174                 bw_img[channel < optimal_threshold] = 255
175             else:
176                 bw_img = np.ones_like(channel) * 255
177                 bw_img[channel < optimal_threshold] = 0
178
179             # Fill in the channels for the final image:
180             final_rgb_img[:, :, i] = bw_img
181
182             # Save the masked image to a file
183             cv2.imwrite(save_path + "RGB_Seperate/Iter" + str(iter) + "/" + img_name +
184             channel_name + ".jpg", bw_img)
185
186             # Update channel to new image for next iteration:
187             # Use the bw_img as a mask on top of the old channel
188             # channel = channel * (bw_img // 255)
189             flattened_channel = flattened_channel[flattened_channel < optimal_threshold]
190
191
192             # Final mask is a logical and of the masks from all of the final masks after all
193             # iterations of thresholding
194             final_mask = np.logical_and(bw_img // 255, final_mask).astype(np.uint8)
195
196             # Save the image with the combined masks:
197             final_mask_img = final_mask * 255
198             cv2.imwrite(save_path + "Final_Images/" + img_name + ".jpg", final_mask_img)
199
200             # Save the final image in RGB with all channels masked separately:
201             cv2.imwrite(save_path + "RGB_Final_Masks/" + img_name + ".jpg", final_rgb_img)
202
203
204 # Calling code:
205 input_folder = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/pics/"
206 img_paths = [input_folder+img_name for img_name in os.listdir(input_folder)]
207
208 # Otsu's Algorithm for RGB images:

```

```

201 # Images are in the following order: Climb, Flower, dog_small, RVL
202 iters = [[3,3,4],
203           [1,1,1],
204           [3,3,3],
205           [2,2,2]]
206 invert = [False, True, False, False]
207 for i, img_path in enumerate(img_paths):
208     img_name = img_path.split("/")[-1][-4]
209     run_Otsus_algorithm_iterative(iters=iters[i], save_path=rgb_path, invert=invert[i],
210                                     img_path=img_path, img=None, img_name=img_name, histogram_folder="RGB/")
211
212 # Calling code for Texture Contouring
213 # Images are in the following order: Climb, Flower, dog_small, RVL
214 iters = [[1,1,1],
215           [2,2,2],
216           [1,1,1],
217           [3,2,2]]
218 invert = [True, True, True, True]
219 window_sizes = [[3, 5, 7],
220                  [3, 5, 7],
221                  [5, 4, 3],
222                  [1, 2, 3]]
223
224 # Otsu's algorithm based on texture:
225 for i, img_path in enumerate(img_paths):
226     img_name = img_path.split("/")[-1][-4]
227     img = cv2.imread(img_path)
228     img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
229
230     final_mask = np.ones_like(img, dtype=np.uint8)
231
232     # Initialize image to fill in with texture data
233     new_img = np.zeros((img.shape[0], img.shape[1], 3))
234
235     for channel_idx ,N in enumerate(window_sizes[i]):
236         # Run Texture Otsu on the image:
237         channel = get_texture_ostu(img, N)
238
239         # Input the texture Otsu outputs as channels for BGR Otsu
240         new_img[:, :, channel_idx] = channel
241     # I now have a new image that I need to segment using BGR Otsu's Algorithm
242     run_Otsus_algorithm_iterative(iters=iters[i], save_path=texture_path, invert=invert[i],
243                                   img_path=None, img=new_img, img_name=img_name, histogram_folder="Texture/")
244
245 # Calling code for contouring based on the images with closing.
246 final_image_dirs = ["/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/Output_Images/
247                      Texture_Segmentation/Final_Images/",
248                      "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW6/Output_Images/RGB_Otsus/
249                      Final_Images/"]
250 types = ["texture", "rgb"]
251 for type, final_img_dir in zip(types, final_image_dirs):
252     for img_name in os.listdir(final_img_dir):
253         # Run the contouring on that image
254
255         # Open the image:
256         img = cv2.imread(final_img_dir + img_name)
257         img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
258
259         # Get the image contour:
260         contour = get_contours(img)
261
262         # Save the images:
263         cv2.imwrite(contour_path+img_name[:-4] + "_" + type + ".jpg", contour)
264
265         # Run Closing on the images:
266         closed_contour = run_closing(contour, 3)
267         cv2.imwrite(contour_path+"Post_Closing/" + img_name[:-4] + "_" + type + ".jpg",
268                     closed_contour)

```