

# ECE 66100 Homework #9

by

Adrien Dubois (dubois6@purdue.edu)

December 3, 2024

## Contents

<b>1 Theory:</b>	<b>2</b>
1.1 Task 1: Image Rectification . . . . .	2
1.2 Task 2: Loop and Zhang's Algorithm . . . . .	4
1.3 Task 3: Dense Stereo Matching using the Census Transform . . . . .	4
1.4 Task 4: Dense Correspondences using Depth Maps . . . . .	5
<b>2 Results: Task 1</b>	<b>5</b>
2.1 Manually Selected Points: . . . . .	5
2.2 Improvement in F caused by LM: . . . . .	6
2.3 Stereo images before and after rectification:	7
2.3.1 Before Rectification: . . . . .	7
2.3.2 Post Rectification: . . . . .	7
2.4 Point Matches . . . . .	8
2.4.1 Rectified point matches using a small angle . . . . .	8
2.4.2 Rectified point matches using a large angle . . . . .	8
2.5 Correspondences using the Canny operator on rectified images: . . . . .	9
2.5.1 Output of the canny operator: . . . . .	9
2.5.2 Removal of lines on the output of the Canny operator using HoughLinesP: . . . . .	9
2.5.3 Point matches using SSD on the Canny edges: . . . . .	11
2.5.4 Point matches from the warped scene mapped back to the original images . . . . .	11
2.6 3D reconstructed scene: . . . . .	12
2.6.1 2 views of the 3D scene: . . . . .	12
2.6.2 Both left and right image points with their corresponding world points: . . . . .	13
<b>3 Results: Task 2</b>	<b>14</b>
3.1 Before Rectification: . . . . .	14
3.2 After Rectification: . . . . .	15
3.3 Commentary on difference in results from my implementation and Loop & Zhang . . . . .	15
<b>4 Results: Task 3</b>	<b>15</b>
4.1 Estimated disparity maps for different window sizes: . . . . .	16
4.2 Accuracy and error masks: . . . . .	17
4.3 Observations on the output quality based on window sizes: . . . . .	17
<b>5 Results: Task 4</b>	<b>18</b>
5.1 Image and depth map for each image pair . . . . .	18
5.2 Depth check process for each image pair . . . . .	18
5.3 3D world coordinates for each pair . . . . .	19
<b>6 Code Printout:</b>	<b>19</b>

# 1 Theory:

## 1.1 Task 1: Image Rectification

Image rectification involves a left and right images that are facing the same object, and point correspondences of that object. We denote points on the left imaging plane as  $\{x_1, x_2, \dots, x_n\}$  and pixel coordinates on the right image as  $\{x'_1, x'_2, \dots, x'_n\}$ .

These pixel coordinates are related to one another through a  $3 \times 3$  fundamental matrix  $F$ :

$$x'Fx = 0$$

where:

$$F = \begin{bmatrix} f_{11} & f_{12} & f_{13} \\ f_{21} & f_{22} & f_{23} \\ f_{31} & f_{32} & f_{33} \end{bmatrix}$$

By expanding the previous equation, initial least-squares estimate of this fundamental matrix can therefore be calculated as follows:

$$\begin{bmatrix} x_1x'_1 & x'_1y_1 & x'_1 & x_1y'_1 & y_1y'_1 & y'_1 & x_1 & y_1 & 1 \\ x_2x'_2 & x'_2y_2 & x'_2 & x_2y'_2 & y_2y'_2 & y'_2 & x_2 & y_2 & 1 \\ \dots & \dots \\ x_nx'_n & x'_ny_n & x'_n & x_ny'_n & y_ny'_n & y'_n & x_n & y_n & 1 \end{bmatrix} \begin{bmatrix} f_{11} \\ f_{12} \\ f_{13} \\ f_{21} \\ f_{22} \\ f_{23} \\ f_{31} \\ f_{32} \\ f_{33} \end{bmatrix} = 0$$

We would therefore need at least 8  $x$  and  $x'$  correspondences to estimate an initial solution for  $F$ , though using closer to 20 proved more successful for my implementation.

The next step in image rectification requires us to convert the cameras to canonical form. This involves centering the left camera on the 0,0 point, and making the imaging planes parallel to one another vis-a-vis the imaged object. To do this, we will first need to calculate the projective matrices  $P$  and  $P'$  for the left and right images respectively. By definition in canonical form, we know that the parameters for the left camera are:

$$P = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \text{ and } e_p = [0 \ 0 \ 0]$$

For the right camera, we rely on the fundamental matrix. We know that  $F^T e'_p = 0$ , so the right epipole is a null vector of the fundamental matrix. Then, we can calculate the right projection matrix as follows:

$$P' = [[e']]_X F \quad |e'|$$

where  $[e']$  is denoted as the skew-symmetric or cross product representation of  $e'$ :

$$[e']_X = \begin{bmatrix} 0 & -e'_3 & e'_2 \\ e'_3 & 0 & -e'_1 \\ -e'_2 & e'_1 & 0 \end{bmatrix}$$

One issue with this estimation of  $P'$  is that it is based on a linear least-squares estimate of the fundamental matrix. Instead we can use the Levenberg-Marquadt algorithm to better refine  $P'$ . To do so, we must first calculate the world points created by the intersection of the rays passing through the centers of projections of each camera and the image pixel correspondences. We do so through the following equation, where  $P_i$  will denote the i-th row of the  $P$  matrix.

$$\begin{bmatrix} x_i P_3^T - P_1^T \\ y_i P_3^T - P_2^T \\ x'_i P_3'^T - P_1'^T \\ y'_i P_3'^T - P_1'^T \end{bmatrix} X_i = 0$$

This gives us a 4 dimensional homogenous coordinate which we can convert to physical coordinate by dividing by the last element. After computing this initial linear least-squares estimate for the world coordinates, we can pass the following parameters to the L.M. optimizer:

$$\text{learnable parameters: } [P' \ X_1 \ X_2 \ \dots \ X_N]$$

And we minimize the following geometric cost function:

$$\epsilon_{geo} = \sum_{i=1}^N (||x_i - \hat{x}_i||^2 + ||x'_i - \hat{x}'_i||^2)$$

where  $\hat{x}$  and  $\hat{x}'$  are the back-projected world coordinates to the respective image planes:

$$\hat{x}_i = P X_i$$

$$\hat{x}'_i = P' X_i$$

After acquiring a non-linear estimate for  $P'$ , also know the right epipole to be its fourth column due to the canonical form setup. Additionally, we can refine our estimate for  $F$  using this improved  $P'$ :

$$F = [e']_X P' P +$$

Next, we now calculate the homographies that would rectify the left and right images. This would make it so that corresponding points in the left and right images can be found on the same rows. We first calculate the right homography:

$$1. \theta = \arctan 2(-\left(e'_2 - \frac{h}{2}\right), -\left(e'_1 - \frac{w}{2}\right))$$

$$2. R = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$3. \tilde{e} = R\tilde{e}' = [\tilde{e}_1 \ \tilde{e}_2 \ \tilde{e}_3]$$

$$4. f = \frac{\tilde{e}_1}{\tilde{e}_3}$$

$$5. G = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ \frac{-1}{f} & 0 & 1 \end{bmatrix}$$

$$6. T_2 = \begin{bmatrix} 1 & 0 & \frac{w}{2} \\ 0 & 1 & \frac{h}{2} \\ 0 & 0 & 1 \end{bmatrix}$$

$$7. T = \begin{bmatrix} 1 & 0 & -\frac{w}{2} \\ 0 & 1 & -\frac{h}{2} \\ 0 & 0 & 1 \end{bmatrix}$$

$$8. H = T_2 G R T$$

On the other hand, to compute the left rectification homography, we follow these steps:

$$1. M = P' P +$$

$$2. H_0 = H' M$$

$$3. H_A = \begin{bmatrix} a & b & c \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

4. The parameters  $a$ ,  $b$ , and  $c$  are calculated through the following manner:

$$(a) z_i = H_0 x_i = [\xi_i \ \nu_i \ 1]$$

$$(b) z'_i = H' x'_i = [\xi'_i \ \nu'_i \ 1]$$

(c) Solve the following linear least-square equation:

$$\begin{bmatrix} \xi_1 & \nu_1 & 1 \\ \dots & \dots & \dots \\ \xi_N & \nu_N & 1 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \end{bmatrix} = \begin{bmatrix} \xi'_1 \\ \nu'_1 \\ 1 \end{bmatrix}$$

5.  $H = H_A H_0$

6. Normalize the homographies by their last entry ( $h_{3,3}$ )

It is important to note that this code only produced solid results when all homogenous matrices were normalized. I did this by dividing all definitions of  $e'$ ,  $P$ ,  $P'$ ,  $H$ ,  $H'$ ,  $z_i$ ,  $z'_i$  by their scaling coordinate (the last entry in each vector/matrix).

Additionally, I used a canny operator to create an edge map from the rectified image. It is important to note that there was significant noise in extracting matching points from this edge map due to the canny operator extracting the border of the images. Therefore, I used a HoughLinesP transform to extract the straight lines from these edge maps, and inpainted the lines in the edge map with black pixels to erase them. This helped to reduce the amount of matching points found on the image borders when using SSD.

Additionally, since I know that I am working on rectified images, I prune the points to have a low absolute difference in width ( $\text{diff}(|x_1 - x_2|) < 4$ ) and absolute difference in height ( $\text{diff}(|y_1 - y_2|) < 4$ ) value. It is clear that setting an upper bound height difference between the coordinate matches would work since both images are rectified. However, I found anecdotally that my images were also located in around the same location in the camera plane so I set a much looser boundary on those points to also speed up the computation of SSD.

Lastly, we need to create the 3D reconstruction from the point matches on the rectified edge map. First, I remapped the point matches to the original, unrectified images by applying the corresponding inverse homography. Then, I re-ran the steps described above up until the LM optimization. Then, I do one last world point reconstruction using linear least-squares with the final  $P$  and  $P'$  matrices and plot those points in 3 dimensions.

## 1.2 Task 2: Loop and Zhang's Algorithm

The Loop and Zhang's algorithm works by decomposing the rectifying homographies calculated just above into three separate components:

$$H = H_{sh} H_{sim} H_p \text{ and } H' = H'_{sh} H'_{sim} H'_p$$

where  $H_{sh}$  is a shearing homography,  $H_{sim}$  is a similarity transformation and  $H_p$  is the purely projective portion of the rectifying homography. The role of  $H_p$  is therefore to send the corresponding epipole to infinity in each image plane. On the other hand, Next, the similarity transformations will rotate the epipoles to coincide with the world x-axis as required by canonical form. Lastly, the shear homographies are there to reduce, as much as possible, the projective distortion.

## 1.3 Task 3: Dense Stereo Matching using the Census Transform

Dense stereo matching involves finding per-pixel correspondences in a given stereo-rectified image pair. In this way, we calculate a disparity map by searching through all pixels in the right image for the one that matches the pixel in the left. We then define the disparity for that pixel as the difference in the x-coordinates between the pair.

More formally, we define the problem as such:

1. For each pixel in the left image  $(x_i, y_i)$ , we define a candidate pixel in the same row by:  $(x'_i - d, y'_i)$
2. For each candidate pixel, as scanned along the row indexed by  $y_i$ , we consider a  $M \times M$  window centered on  $(x'_i - d, y'_i)$
3. We then create two bitvectors of shape  $M \times M$  (one for each images). The values in these bitvectors are defined as follows:
  - 1 if the pixel in the window is strictly larger than the center pixel

- 0 otherwise
- We then XOR these two bitvectors together and the sum of the digits of the resulting bitvector (the number of ones) is our data cost
  - By finding the x-coordinate that minimizes the data-cost, we can set the correct disparity value  $d$  and create the disparity map accordingly.

## 1.4 Task 4: Dense Correspondences using Depth Maps

*Explain in your own words the automatic extraction of dense correspondences using depth maps and why it is useful* A depth map is created by calculating the distance between each real world coordinate in the image, and the camera's principle plane. These are especially useful to real time navigation based computer vision such as person-tracking or autonomous vehicles. In this way, we can use these depth maps to check whether our extracted dense correspondences are accurate.

This is done by first calculating an inverse of the camera matrix  $K_A$ :

$$K^{-1} = \begin{bmatrix} \frac{1}{\alpha_x} & -\frac{s}{\alpha_x \alpha_y} & \frac{s y_0 - \alpha_y x_0}{\alpha_x \alpha_y} \\ 0 & \frac{1}{\alpha_y} & -\frac{y_0}{\alpha_y} \\ 0 & 0 & 1 \end{bmatrix}$$

We can then calculate the 3D coordinate of any point on the ray corresponding to the camera pixel  $x_i$ :

$$X_{cam} = K^{-1} x_i$$

We can then get this 3D coordinate in terms of the image coordinate frame by multiplying it by its depth value at that point. Afterwards, using the rotation and translation matrices we convert the 3D image coordinate to the world coordinate frame.

$$X_{worldhc} = T^{-1} X_{camhc}$$

Next, we estimate the depth value in the right image by calculating  $X_{cam}^B = [R_B \ | t_B] X_{world}$ . The Z value, is therefore this estimated depth value. We can also determine the true depth value by projecting the 3D world coordinate onto image B and getting the depth value at that point.

It is the estimated depth coordinate, and the true depth coordinate that we compare for accuracy. If the difference between the two points is less than a given threshold, we consider the two points correspondences.

## 2 Results: Task 1

### 2.1 Manually Selected Points:

It is important to note that my implementation of stereo rectification ended up being much more robust to large angular changes compared to the provided Loop and Zhang code. Therefore, the angle between the images looks very small so that Loop and Zhang works. For demonstration, I will also include an image with a much larger difference between the left and right images in the rectification section; however, since that image did not work for Loop and Zhang I do not report its results everywhere.



(a) Points chosen for the left image



(b) Points chosen for the right image

## 2.2 Improvement in F caused by LM:

To check whether my algorithm was working correctly, I calculated  $x_i'Fx_i$  for all points i and report the mean absolute value. Therefore, I could confirm that my LM was working not just through the decreased cost function, but also by this error value getting closer to 0.

Metric	Before LM	After LM
Cost	1.6670e+02	3.3692e+00
Error	0.086739	0.0753223

Table 1: Comparison of metrics before and after applying LM.

## 2.3 Stereo images before and after rectification:

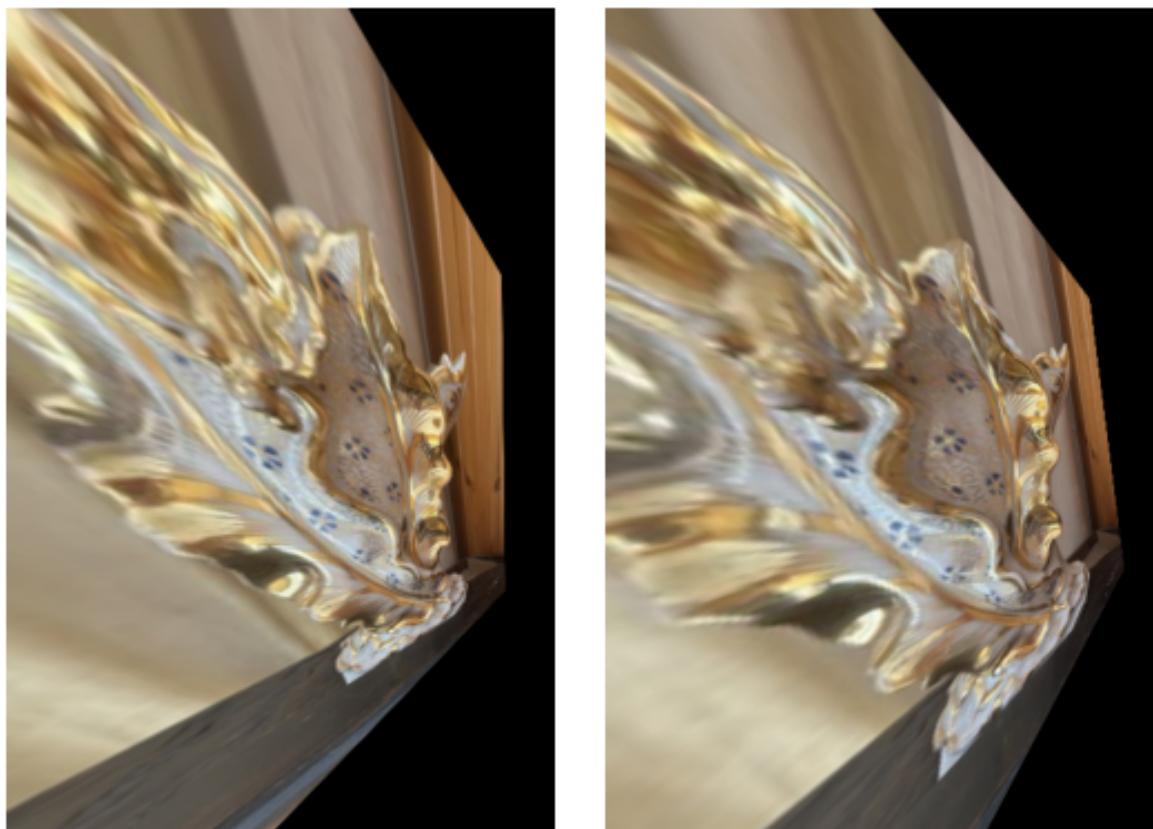
### 2.3.1 Before Rectification:



(a) Left image before rectification

(b) Right image before rectification

### 2.3.2 Post Rectification:



(a) Left image after rectification

(b) Right image after rectification

Please note that I applied a translation homography to these images for display; however, this was not applied for the next images to stay consistent with the instructions.

## 2.4 Point Matches

To support my point that my algorithm was more robust to larger angle changes than the Loop and Zhang implementation, I also report point matches with another abject, and greater rotation, which caused errors with the provided Loop and Zhang code.

### 2.4.1 Rectified point matches using a small angle

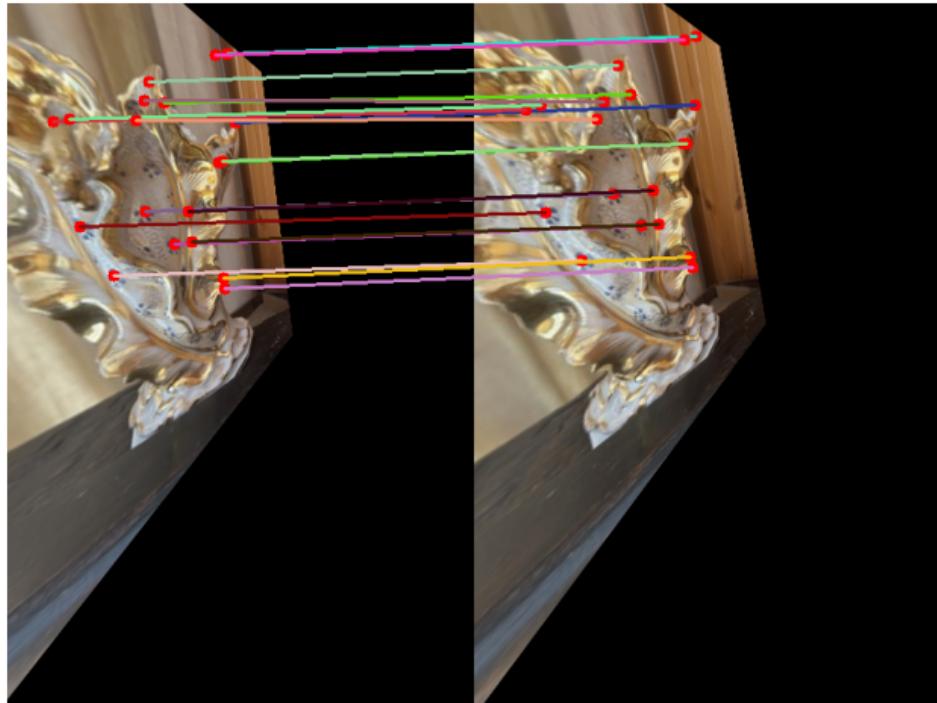


Figure 4: Point matches printed on the rectified images with a small angle between the left and right images.

### 2.4.2 Rectified point matches using a large angle

Original Images:



(a) Left image before rectification



(b) Right image before rectification

Rectified Images:

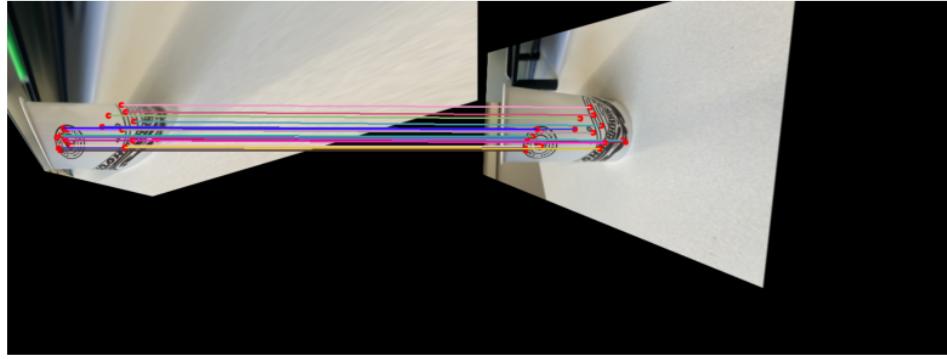


Figure 6: Point matches printed on the rectified images with a large angle between the left and right images.

## 2.5 Correspondences using the Canny operator on rectified images:

### 2.5.1 Output of the canny operator:

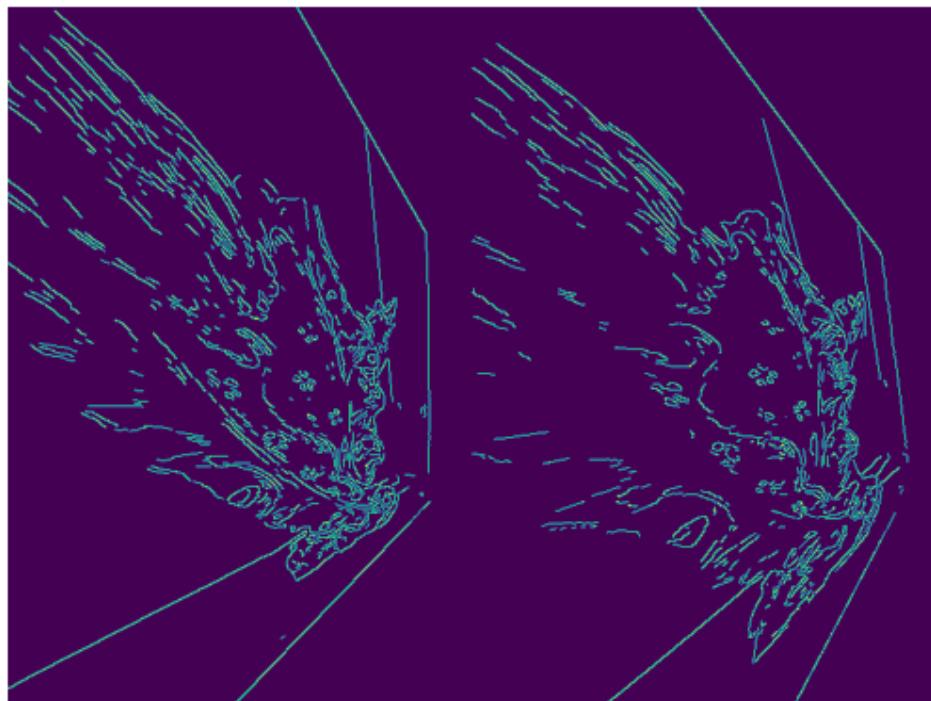


Figure 7: Output of the canny operator with thresholds 150, 200 applied on the rectified image pairs.

### 2.5.2 Removal of lines on the output of the Canny operator using HoughLinesP:

After applying the Canny operator to get the edges of the image, I used my previous implementation of the SSD point matching algorithm to collect 200 point matches between the two images. It is important to note that the canny operator had one main issue in this context which was that it extract lines at the borders of the images due to the perspective transform we applied. This means that all of the top 200 point matches actually occurred on the edges of the image instead of the object in question. This issue is demonstrated in the image below.

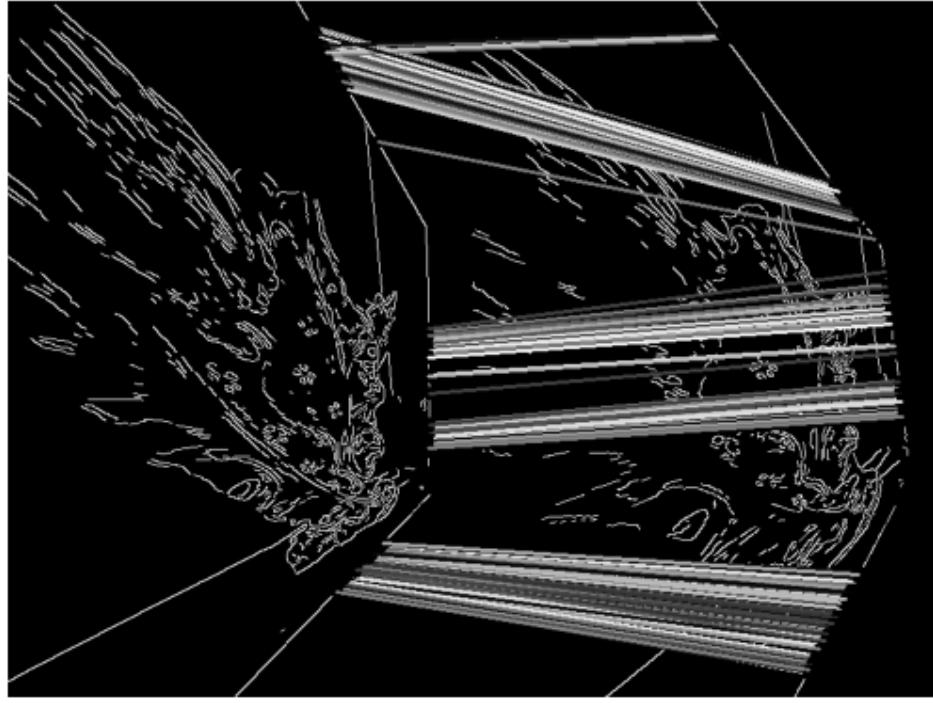


Figure 8: Top 200 point matches from the edged images.

However, to get a good 3D reconstruction of the object, we obviously want the point matches to be detected on the object in question and not on the borders of the image. Therefore, we use a HoughLine transform to detect the location of straight lines on the image, and remove these lines from the edge image using inpainting with black pixels. With this method, I purposefully over-remove lines from the image to ensure that as much of the image border edges are removed. In this way, not point matches will be detected on the borders, and there are still more than enough points left in the canny edge image for hundreds of solid point matches.

The parameters I used for HoughLinesP to remove the lines were:

Parameter	Value
<code>rho</code>	1 (pixel resolution)
<code>theta</code>	$\pi/180$
<code>threshold</code>	5
<code>minLineLength</code>	40 (minimum line length in pixels)
<code>maxLineGap</code>	10 (maximum gap between segments to connect as a single line)

Table 2: Parameters used in the `cv2.HoughLinesP` function call.

and the resulting edge map was:

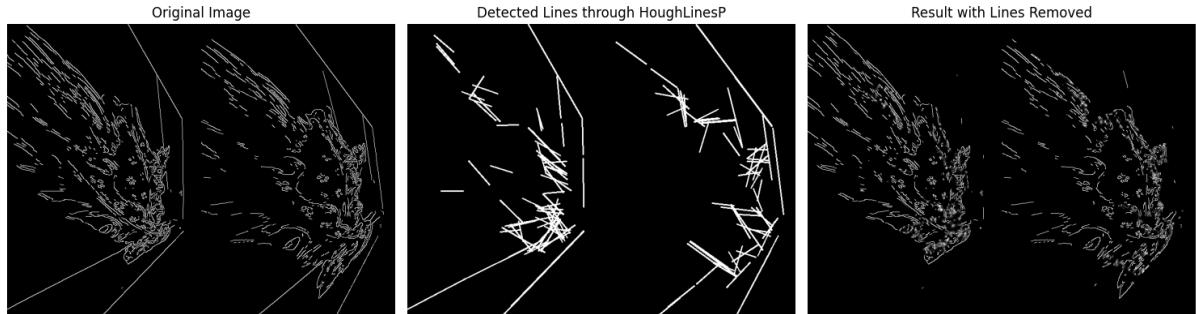


Figure 9: The original edges, the lines removed and the final edges with straight lines removed.

### 2.5.3 Point matches using SSD on the Canny edges:

Included below are the point matches found before and after removing the lines using HoughLinesP. For the first image, I report 200 matches while I report 1000 matches in the final image with lines removed.

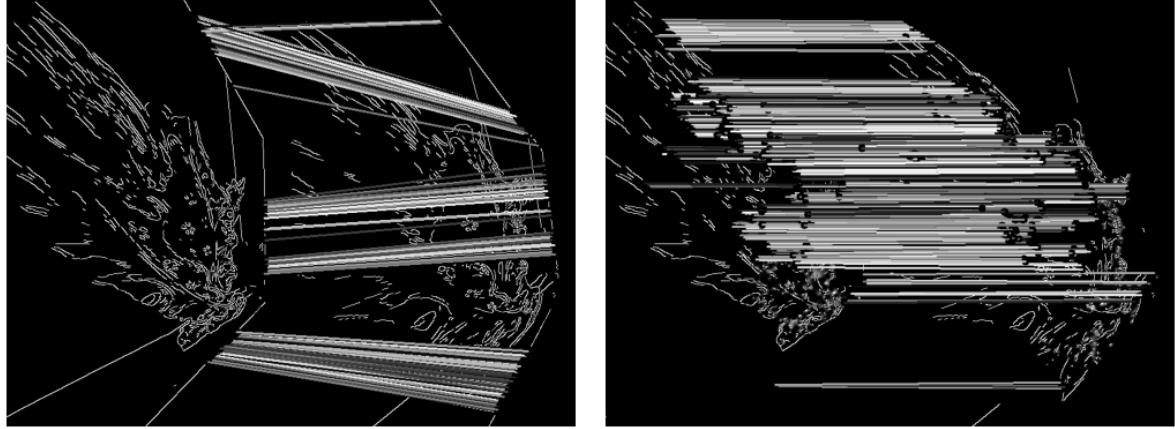


Figure 10: The the point matches before (left) and after (right) removing the image borders using houghlinesP.

### 2.5.4 Point matches from the warped scene mapped back to the original images

By applying the inverse homographies ( $H^{-1}$  and  $H'^{-1}$ ) to the point matches on the rectified images, we can find their corresponding locations in the original images. The result of this operation is shown below. This allows us to re-run the previous operations completed in task one but with almost 1000 matching points instead of around 20 manually selected matches.

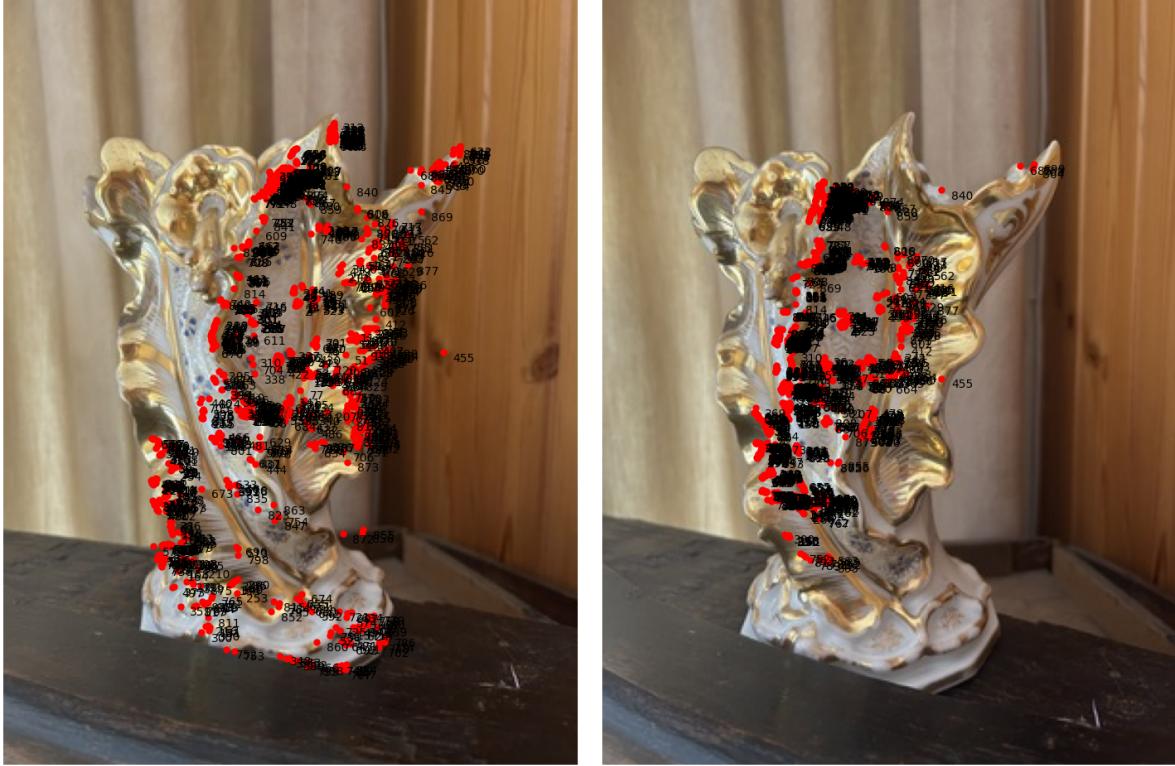


Figure 11: Point matches from removing the homographies on the canny point matches.

## 2.6 3D reconstructed scene:

Since the error  $x_i' F x_i$  was already reducing from 0.0753223 to 0.01464 through the use 1000 point matches, I only ran my LM algorithm on 100 random points samples from the point matches. I then use the improved P and  $P'$  to do a final linear least-squares estimate of all the world coordinates. While it would have been possible to run LM on all 1000 points, the time constraints were not feasible in this setting as you would need to optimize over 6012 parameters for each run of the cost function.

### 2.6.1 2 views of the 3D scene:

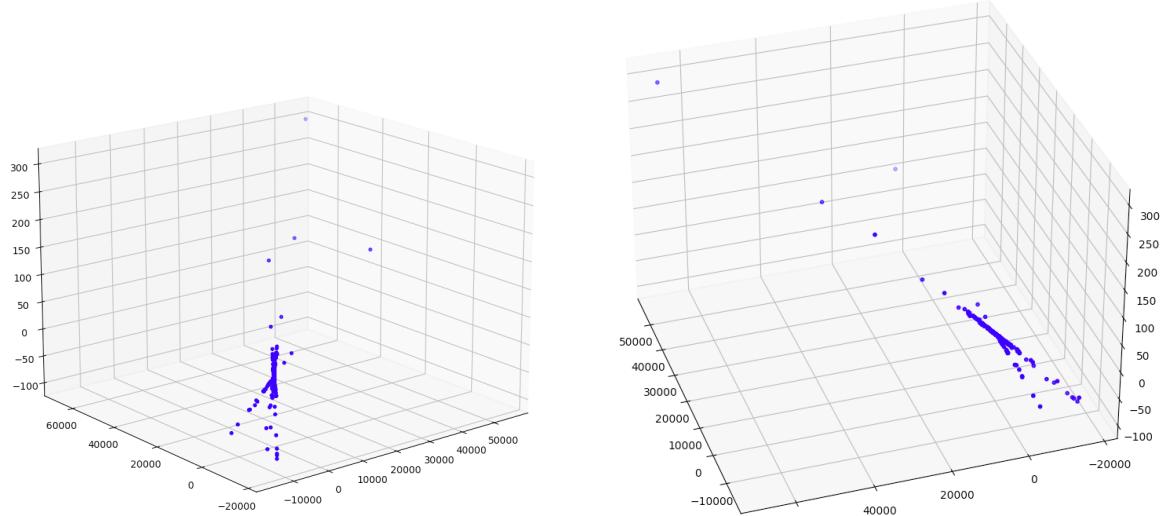


Figure 12

### 2.6.2 Both left and right image points with their corresponding world points:

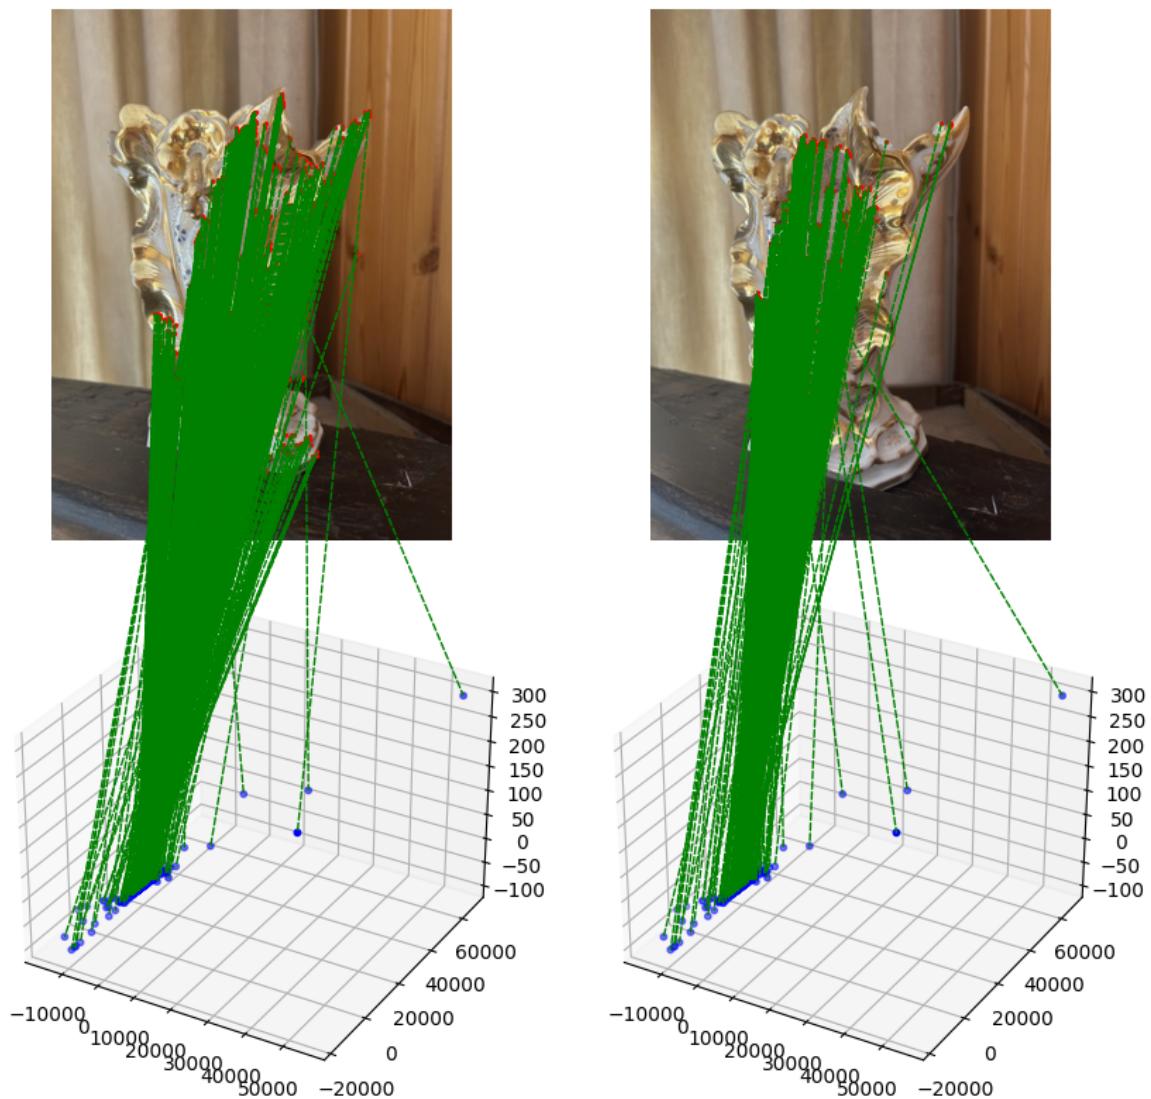
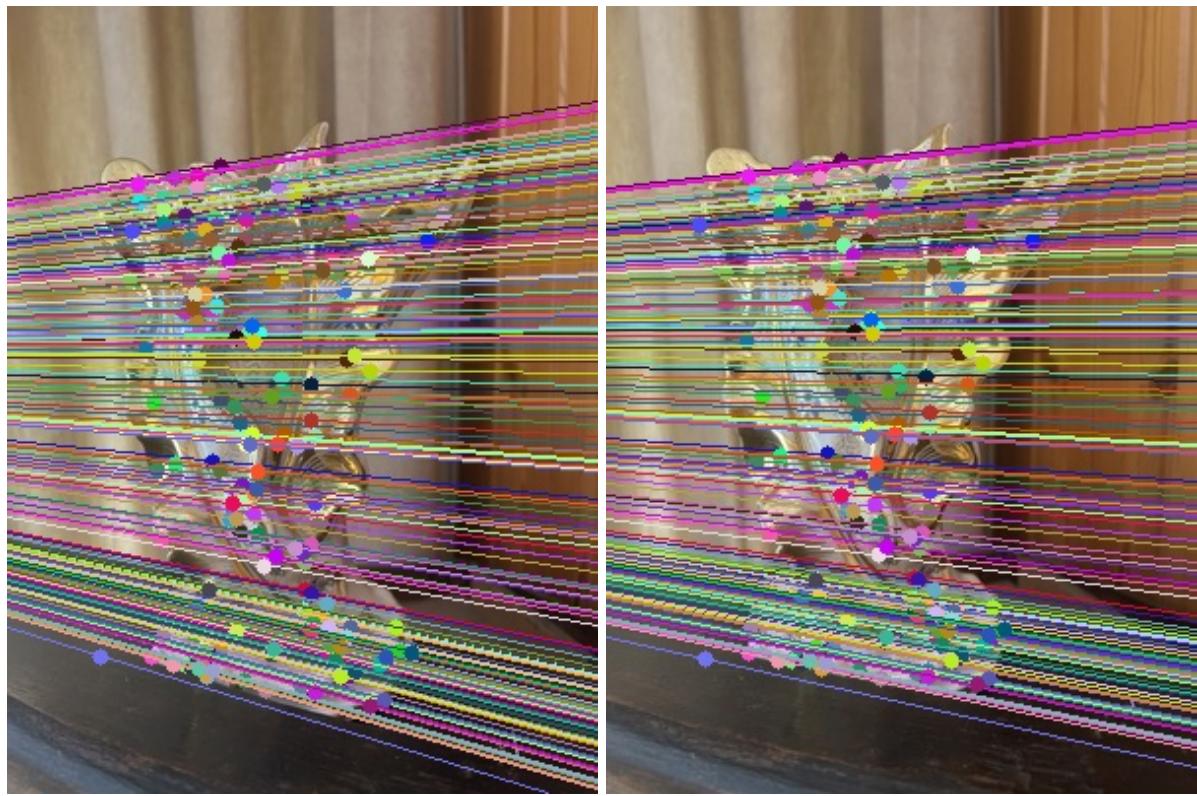


Figure 13

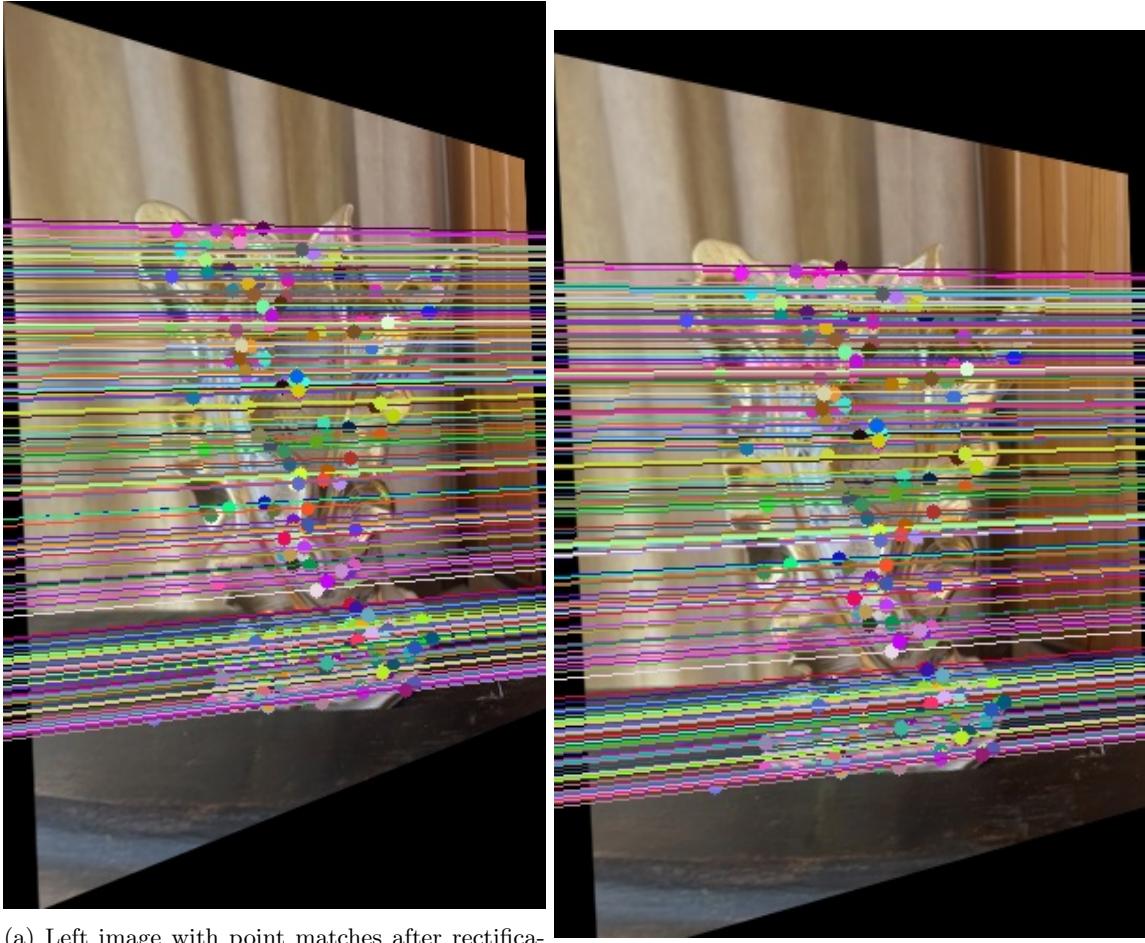
### 3 Results: Task 2

#### 3.1 Before Rectification:



(a) Left image with point matches before rectification (b) Right image with point matches before rectification

### 3.2 After Rectification:



(a) Left image with point matches after rectification

(b) Right image with point matches after rectification

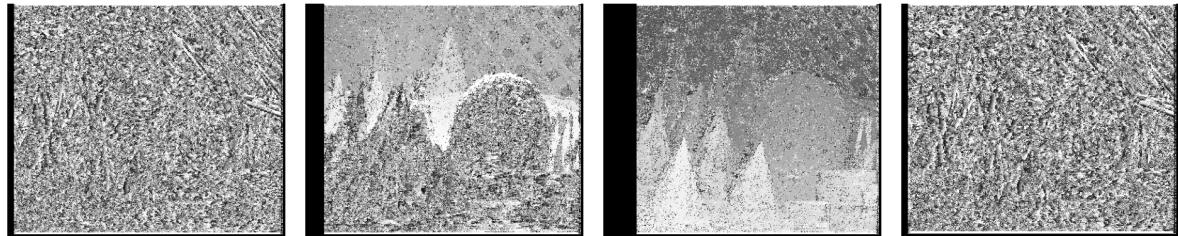
### 3.3 Commentary on difference in results from my implementation and Loop & Zhang

When the Loop and Zhang implementation worked, it found a lot more point matches and the lines were much more horizontal without requiring any pruning. On the other hand, my implementation ran into some errors when the point correspondences ended up outside of the valid pixel range which required manual pruning, and a translocation homography to move more of the image back into the frame. On the other hand, as shown in the cup image results section of task 1, Loop and Zhang did not work on some images that worked for my implementation. The error reported by the provided code pointed to a too larger distance between the left and right images which caused floating point errors in their equations. Therefore, the ideal option for 3D reconstruction would be lots of images with small angle differences, and rectification done using Loop and Zhang's algorithm. If only a few images of the scene were taken, and the angle between them was therefore larger I would recommend trying our implementation with hand picked correspondences for rectification.

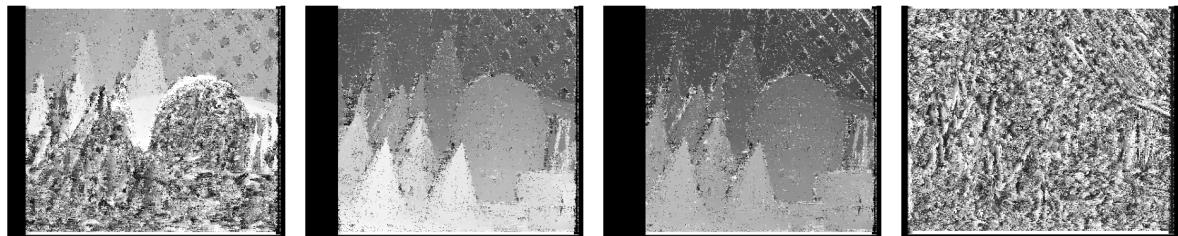
## 4 Results: Task 3

For this section, I report results from a range of window sizes and  $d_{max}$  values. Note that the maximum value found in the ground truth disparity map was 54; however, I wanted to test different ranges to see how this would affect the results.

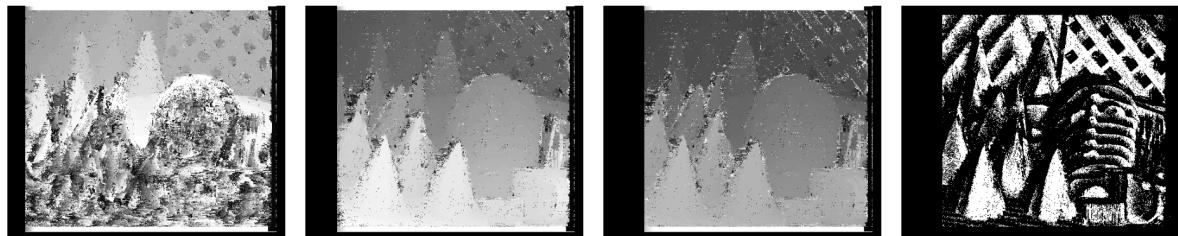
#### 4.1 Estimated disparity maps for different window sizes:



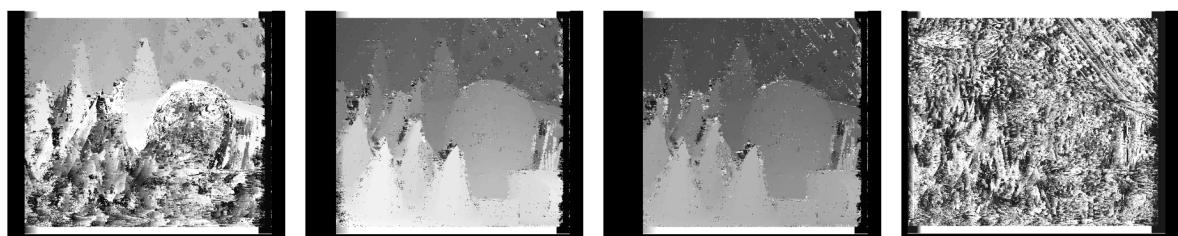
(a) Disparity maps with a window size of **5** and a dmax of [10, 30, 54, 70]



(b) Disparity maps with a window size of **9** and a dmax of [10, 30, 54, 70]

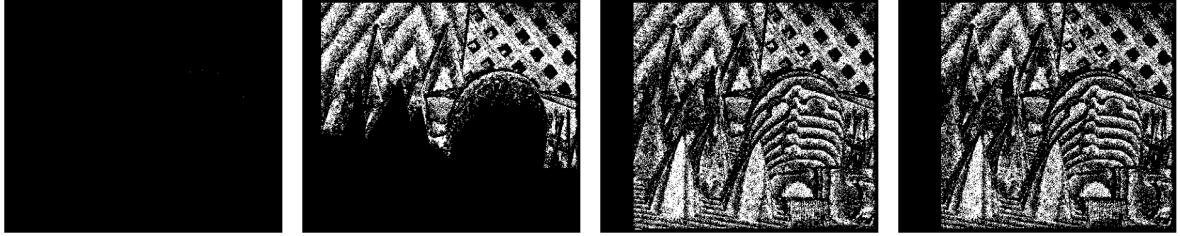


(c) Disparity maps with a window size of **16** and a dmax of [10, 30, 54, 70]

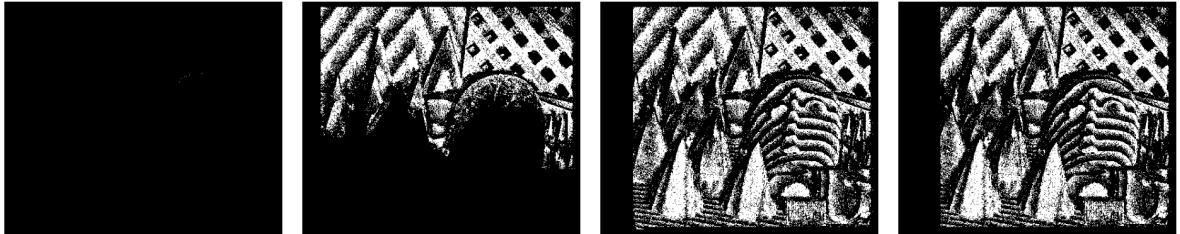


(d) Disparity maps with a window size of **25** and a dmax of [10, 30, 54, 70]

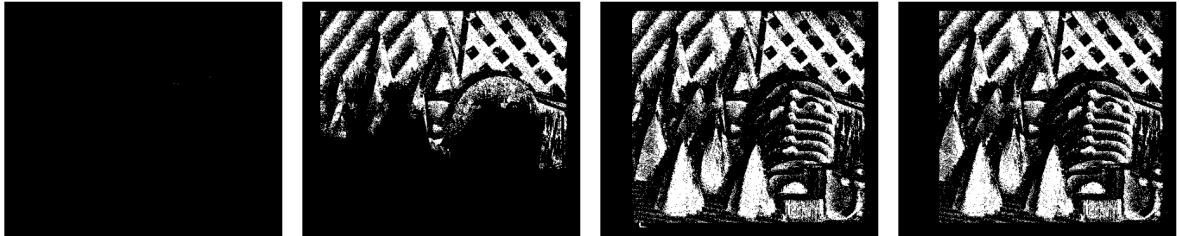
## 4.2 Accuracy and error masks:



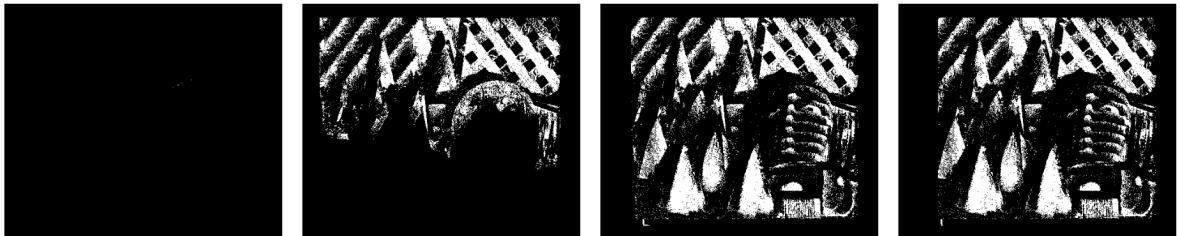
(a) Accuracy maps with a window size of **5** and a dmax of [10, 30, 54, 70]



(b) Accuracy maps with a window size of **9** and a dmax of [10, 30, 54, 70]



(c) Accuracy maps with a window size of **16** and a dmax of [10, 30, 54, 70]



(d) Accuracy maps with a window size of **25** and a dmax of [10, 30, 54, 70]

## 4.3 Observations on the output quality based on window sizes:

Included below are histograms of the impact on changing the window size and  $d_{max}$  values on the accuracy of the disparity map. As you can see, the best results were found with a  $d_{max}$  that matched the maximum in the ground truth since it creates an upper bound for the disparity error per point. Additionally, we can see that the **accuracy is highest around a window size of 9 for all values of  $d_{max}$** . Lastly, as the  $d_{max}$  value went down, the accuracy also went down sharply (much larger effect than increasing  $d_{max}$ ).

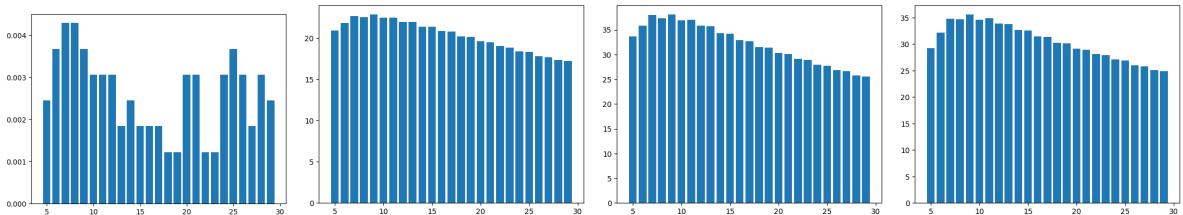


Figure 18: Histogram of the accuracy for using  $d_{max}$  values of [10, 30, 54, 70] (top to bottom)

## 5 Results: Task 4

For plotting these graphs, I include 100 probe points for the point matching process, and include 2500 for the 3D reconstruction so that the object can better be seen. Some images show only a few matching or no matching points since the estimated and real depths of the detected point matches were too different. This demonstrates the power of using depth checking for correspondence points to remove noise in the output of an image.

### 5.1 Image and depth map for each image pair

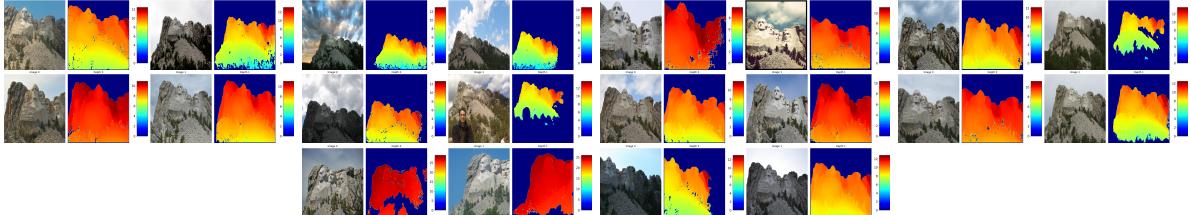


Figure 19: Depth checked images with corresponding matching points.

### 5.2 Depth check process for each image pair

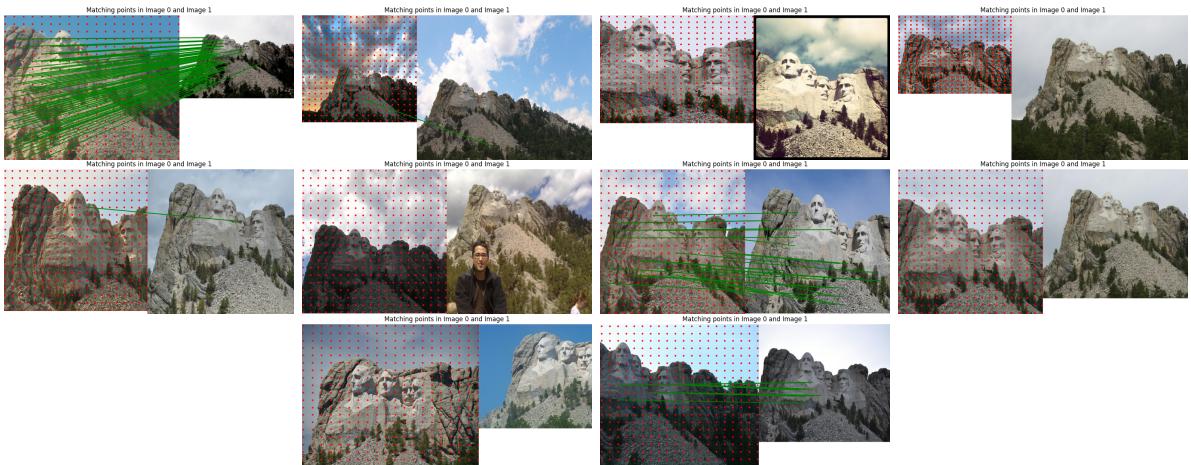


Figure 20: Depth checked images with corresponding matching points.

### 5.3 3D world coordinates for each pair

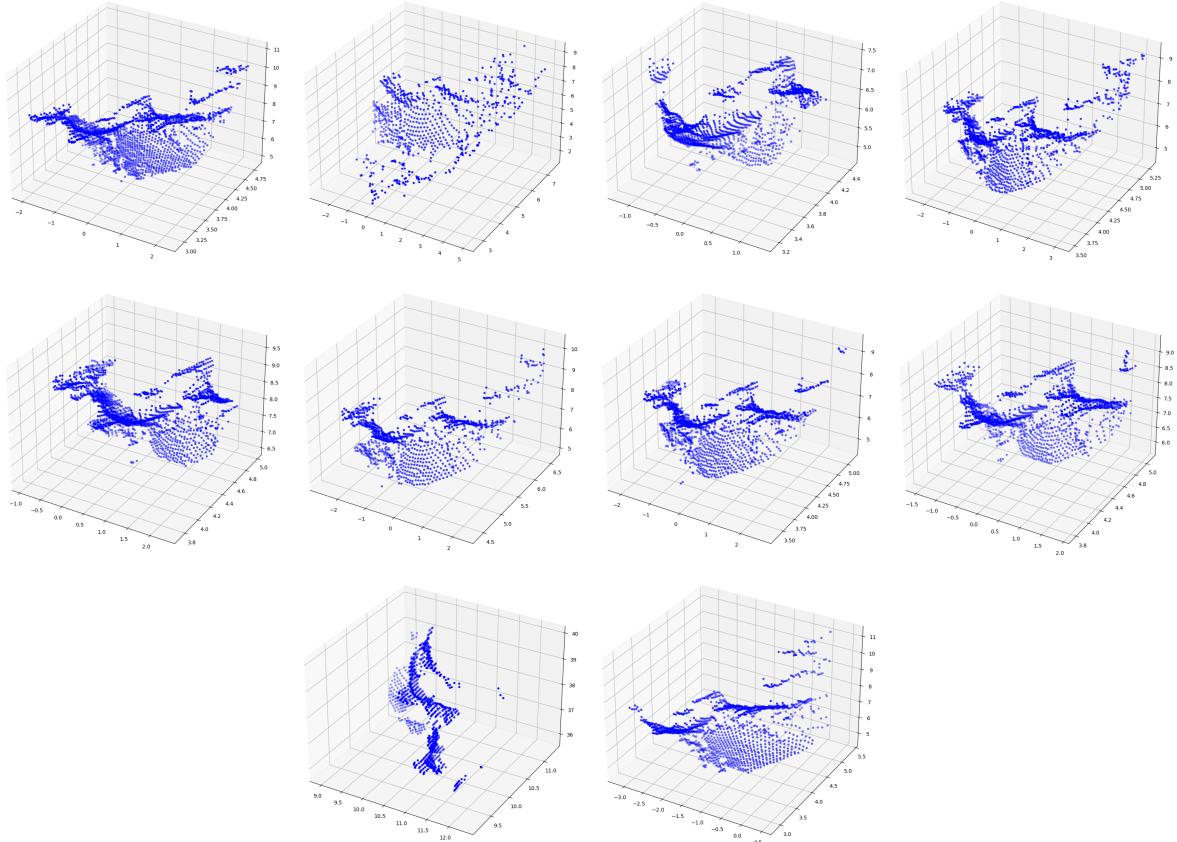


Figure 21: 3D point reconstructions for all image pairs.

## 6 Code Printout:

```

1 # %%
2 import cv2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from scipy.optimize import least_squares
6 import pickle
7 import random
8 from tqdm import tqdm
9
10 # %%
11 img_dir = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW9/Images/"
12
13 # %%
14 class Point():
15     def __init__(self, x, y):
16         """Defines a point using its physical space coordinates"""
17         self.x = x
18         self.y = y
19         self.hc = self.get_hc()
20
21     @classmethod
22     def from_hc(cls, hc):
23         """Defines a point from its representation in homogeneous coordinates"""
24         if np.isclose(hc[2], 0):
25             x = hc[0]
26             y = hc[1]
27         else:
28             x = hc[0] / hc[2]
29             y = hc[1] / hc[2]
30         return cls(x, y)

```

```

30     def get_hc(self):
31         """Returns the point in homogeneous coordinates"""
32         return np.array([self.x, self.y, 1])
33     def apply_homography(self, H):
34         self.hc = H @ self.hc
35         return self.hc
36     def apply_x_y_offset(self, x_offset, y_offset):
37         self.__init__(self.x + x_offset, self.y + y_offset)
38         return self.hc
39     def __repr__(self):
40         """To string method for debugging"""
41         return f"Point(x={self.x}, y={self.y}, hc={self.hc})"
42
43 # %%
44 def open_image_in_grayscale(img_path):
45     img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
46     return img
47
48 # %%
49 # For view 3 and 4:
50 # Left Image:
51 img1_coords = [(135.39086538156292, 142.9686706914469), (153.67407839960097,
52     150.35743901737024), (97.53385960303707, 153.29330863104485), (150.01673332962673,
53     242.3620485533101), (154.906077176412, 213.23856563984987), (166.59798637524642,
54     217.06500865037748), (168.20133062334986, 236.29312439383074), (174.10846178282134,
55     228.46617560753106), (158.30688593123517, 251.50398712946978), (140.20318683016325,
56     286.359314173019), (137.86601754532887, 291.3675340690927), (123.99692097166292,
57     154.54893583465073), (117.26992782683912, 120.09360509287032), (141.2245863425531,
58     121.40618912112862), (110.44752003348668, 122.43507091165702), (144.4775692921619,
59     133.73249617162207), (156.46386194602724, 141.5855844620856), (110.71053591493266,
60     256.13154910371065), (125.36770173155205, 261.28376496652237), (136.11628999707295,
61     264.9258485936823)]
62
62 # # Right Image:
63 img2_coords = [(103.94194116801071, 146.10462964705587), (125.13056380524532,
64     155.512106943920315), (76.58935558176239, 150.3423541745028), (131.6160480507922,
65     243.6500281999699), (134.20041793306098, 216.25570744792097), (151.7741331324886,
66     223.31965179278893), (158.1723424180141, 242.3256019935444), (170.37685143709615,
67     236.66553983976723), (143.2063567947179, 253.9351568492766), (125.23883285132553,
68     280.4958444177697), (123.20772144902898, 284.4018278837246), (95.91480113148509,
69     155.57379406148593), (86.17062836890693, 121.6411453823902), (107.37853379334175,
70     125.30953983418433), (81.40673972671068, 123.36838476930289), (112.33528369563939,
71     137.254669816577), (126.6729958282682, 146.34430177937818), (97.56755810749733,
72     244.93423422613654), (108.28612115429344, 253.4884720423296), (118.69549488243196,
73     259.4661322030428)]
74
74 # img1_points = [Point(x, y) for (x, y) in img1_coords]
75 # img2_points = [Point(x, y) for (x, y) in img2_coords]
76
76 # %%
77 # Left image:
78 img1_coords = [(95.36836337466062, 168.76952141057927), (133.9695606660342,
79     162.22694559848208), (102.56519676796754, 202.13665805227512), (147.05471229022865,
80     189.0515064280807), (210.41955903039025, 120.35446040105984), (213.69084693643885,
81     120.35446040105984), (162.0044980208708, 191.66853675291958), (165.2757859269194,
82     162.8812031796918), (164.8668749386633, 69.32236906670147), (237.48946645294242,
83     79.79049036605704), (168.79242042592162, 316.6317347639765), (155.05301122051745,
84     308.7806437894599), (71.39112403685812, 80.72231890381454), (104.32980104565014,
85     112.07505961218324), (111.64950704760392, 108.04922131110865), (112.3708226634825,
86     230.9239916255029), (124.80177630933298, 259.05706761752094), (100.56290877527834,
87     314.84556210945397), (85.67466735946616, 303.6793810475948), (78.58502859003177,
88     311.3007427247368), (184.86684848498908, 243.86891309260727), (185.6544659729186,
89     232.5469117036203), (159.46618449926174, 265.7252983826517)]
90
90 # Right iamge
91 img2_coords = [(89.12020347410777, 171.38655173541815), (124.45011285943275,
92     165.49823350453067), (95.66277928620497, 205.40794595832375), (138.84377964604664,
93     193.63130949654874), (202.83017108835747, 121.66297556347928), (208.0642317380353,
94     119.70020281985012), (153.76085249762832, 194.93982465896818), (157.68639798488664,
95     167.46100624815983), (151.16835683208478, 73.24791455395984), (236.22184238934864,
96     75.210687297589), (163.59925087506952, 324.4828257384932), (150.51409925087506,
97     315.32321960155707), (66.09429952536783, 85.5322304671512), (93.9072102215332,
98     115.42739116195997), (100.89762093126461, 112.15528402123464), (103.59034888191881,
99     115.42739116195997)]

```

```

233.7957134749712), (115.61328702986482, 262.897965627476), (97.57875165958416,
317.83723897614044), (87.65648352360236, 303.7371737302716), (79.56200162319618,
312.7455487484656), (178.4377679783561, 252.25864453774142), (179.23394652593706,
239.5197877764464), (151.6662643159471, 271.6654966350267)]
65
66 # img2_coords = [(95, 168), (290, 390)]
67
68 img1_points = [Point(x, y) for (x, y) in img1_coords]
69 img2_points = [Point(x, y) for (x, y) in img2_coords]
70
71 # %%
72 plt.figure(figsize=(16, 12))
73 plt.imshow(cv2.cvtColor(cv2.imread(img_dir+"view1.jpg"), cv2.COLOR_BGR2RGB))
74 plt.axis("off")
75 for i, (x, y) in enumerate(img1_coords):
76     plt.plot(x, y, 'ro', markersize=5) # Plot red circle
77     plt.text(x + 5, y + 5, str(i + 1), color="black", fontsize=10)
78
79 plt.show()
80
81 # %%
82 class Image_Rectifier():
83     def __init__(self, points1, points2, img1_path, img2_path, img1=None, img2=None):
84         self.F = None
85         self.points1 = np.array([point.hc for point in points1])
86         self.points2 = np.array([point.hc for point in points2])
87         if img1_path != None:
88             self.img1 = open_image_in_grayscale(img1_path)
89             self.img2 = open_image_in_grayscale(img2_path)
90         else:
91             self.img1 = img1
92             self.img2 = img2
93
94     def calculate_fundamental_matrix(self):
95         points_mat = []
96         for (point1, point2) in zip(self.points1, self.points2):
97             x, y, _ = point1 # x (left image)
98             xp, yp, _ = point2 # x' (right image)
99
100         point_mat_entry = np.array([x*xp, xp*y, xp, x*yp, y*yp, yp, x, y, 1])
101         points_mat.append(point_mat_entry)
102
103     points_mat = np.vstack(points_mat)
104
105     # Estimate fundamental matrix as the Null space
106     _, _, Vt = np.linalg.svd(points_mat)
107     f = Vt[-1]
108
109     F = f.reshape((3,3))
110     F = F / F[-1, -1]
111
112     # Ensure rank 2:
113     U, D, Vt = np.linalg.svd(F)
114     D_rank2 = np.array([[D[0], 0, 0],
115                         [0, D[1], 0],
116                         [0, 0, 0]])
117     self.F = U @ D_rank2 @ Vt
118     self.F = self.F / self.F[-1, -1]
119
120     @staticmethod
121     def get_cross_form(ep_vec):
122         ep1, ep2, ep3 = ep_vec
123         ep_cross = np.array([[0, -ep3, ep2],
124                             [ep3, 0, -ep1],
125                             [-ep2, ep1, 0]])
126
127         return ep_cross
128
129     def get_cononical_form(self):
130         # Left projection matrix in canonical form:
131         self.P = np.array([[1, 0, 0, 0],
132                           [0, 1, 0, 0],
133                           [0, 0, 1, 0]])
134
135         self.e = np.array([0, 0, 0]) # Unsure if this needs to be in 3D coords or HC

```

```

134     # Now to calculate the right epipole and right projection matrix
135     # For ep, we know that: F^T ep = 0, so we can get the Null vector to solve this:
136     _, _, Vt = np.linalg.svd(self.F.T)
137     self.ep = Vt[-1]
138     # Normalize e':
139     self.ep = self.ep / self.ep[-1]
140
141     # Make it into cross product form:
142     ep_cross = Image_Rectifier.get_cross_form(self.ep)
143
144     self.Pp = np.hstack((ep_cross @ self.F, self.ep.reshape(-1, 1)))
145     self.Pp = self.Pp / self.Pp[-1, -1]
146
147     @staticmethod
148     def cost_func(learnable_params, x_points, xp_points, triangulation=False):
149         # Unpack the learnable parameters:
150         P = np.array([[1, 0, 0, 0],
151                      [0, 1, 0, 0],
152                      [0, 0, 1, 0]])
153         Pp = learnable_params[:12].reshape((3,4))
154
155         # Every three points past is a 3 element vector for the world points;
156         world_points = np.array(learnable_params[12:]).reshape(-1, 3)
157         # Add w parameter back into world_points:
158         world_points = np.hstack((world_points, np.ones((world_points.shape[0], 1))))
159
160         # Backproject the world points in image hc coordinates
161         x_hat_h = (P @ world_points.T).T
162         xp_hat_h = (Pp @ world_points.T).T
163
164         # Get the normalized homogenous coordinates
165         x_hat = x_hat_h / (x_hat_h[:, 2, np.newaxis] + 1e-8)
166         xp_hat = xp_hat_h / (xp_hat_h[:, 2, np.newaxis] + 1e-8)
167
168         # Only compare real pixel coordinates, not homogenous coords since we want it to
169         # be invariant to a scalar factor of k
170         diff_x = x_points[:, :2] - x_hat[:, :2]
171         diff_xp = xp_points[:, :2] - xp_hat[:, :2]
172
173         residuals = np.hstack((diff_x[:, 0], diff_x[:, 1], diff_xp[:, 0], diff_xp[:, 1]))
174         .flatten()
175         return residuals
176
177     @staticmethod
178     # As we refine pp, we need to re-estimate the world coordinates:
179     def estimate_world_coordinates(P, Pp, points1_hc, points2_hc):
180         # Create the matrix to find the world points. A_mat @ X = 0:
181         world_points = []
182         for i, (point1, point2) in enumerate(zip(points1_hc, points2_hc)):
183             x1, y1, _ = point1
184             x2, y2, _ = point2
185
186             P1, P2, P3 = P[0, :, :], P[1, :, :], P[2, :, :]
187             Pp1, Pp2, Pp3 = Pp[0, :, :], Pp[1, :, :], Pp[2, :, :]
188
189             A_mat = np.array([x1*P3 - P1,
190                             y1*P3 - P2,
191                             x2*Pp3 - Pp1,
192                             y2*Pp3 - Pp2])
193             _, _, Vt = np.linalg.svd(A_mat)
194
195             # Get the homogenous coordinates:
196             # Each point is 1 column, each point has 4 coords: x, y, z, w
197             X_h = Vt[-1]
198
199             # Convert to inhomogeneous world points:
200             # Each point is 1 column, each point has 4 coords: x, y, z, 1
201             X = X_h[:-1] / X_h[-1]
202             world_points.append(X)
203             world_points = np.array(world_points)
204             return world_points
205
206     def refine_proj_prime_matrix_with_nonlinear_leastsquares(self):

```

```

205     learnable_params = self.Pp.flatten()
206     idx = np.random.choice(self.points1.shape[0], size=75, replace=False)
207     points1 = self.points1[idx]
208     points2 = self.points2[idx]
209     print(points1.shape, points2.shape)
210     world_points = Image_Rectifier.estimate_world_coordinates(self.P, self.Pp,
211     points1, points2)
212     learnable_params = np.hstack((learnable_params, world_points.flatten()))
213
214     # Apply non-linear least squares optimization to the P' matrix:
215     optim_l = least_squares(Image_Rectifier.cost_func, learnable_params, args=(
216     points1, points2, False), method='lm', verbose=2)
217
218     learnable_params = optim_l.x
219     self.Pp = learnable_params[:12].reshape((3, 4))
220     self.Pp = self.Pp / self.Pp[-1, -1]
221
222     # # Every three points past is a 3 element vector for the world points;
223     world_points = np.array(learnable_params[12:]).reshape(-1, 3)
224     # Add w parameter back into world_points:
225     self.world_points = np.hstack((world_points, np.ones((world_points.shape[0], 1)))
226   )
227
228   def get_refined_fund_matrix_from_Pp(self):
229     # The epipole of P' is its 4th column
230     self.ep = self.Pp[:, -1]
231     # Normalize e':
232     self.ep = self.ep / self.ep[-1]
233
234     # We can then calculate a refined value for F:
235     ep_cross = Image_Rectifier.get_cross_form(self.ep)
236
237     # Refined version of F: F = [e']x P' P+ (P+ = pseudo-inv of P)
238     self.F = ep_cross @ self.Pp @ np.linalg.pinv(self.P)
239     self.F = self.F / self.F[-1, -1]
240
241   def get_homographies_from_refined_F(self):
242     h, w = self.img1.shape
243
244     # Extract the columns of e' to calculate the angle between the components of e'
245     ep1, ep2, _ = self.ep
246
247     # Calculate the rotation matrix for the right camera:
248     theta = np.arctan2(-(ep2 - h/2), -(ep1 - w/2))
249     self.R = np.array([[np.cos(theta), -np.sin(theta), 0],
250                       [np.sin(theta), np.cos(theta), 0],
251                       [0, 0, 1]])
252
253     # After applying the rotation matrix, the epipole should align with the x_axis
254     # Ideally, we would get: [f, 0, 1] so: f = er1 / er3 (normalize by third coord)
255     # and er2=0 ideally
256     ep_rot = self.R @ self.ep
257     print(ep_rot)
258     focal_len = ep_rot[0] / ep_rot[2]
259
260     # RIGHT HOMOGRAPHY #####
261     # We can now calculate the right homography: H' = G R T
262     # Calculate G: the geometric transformation based on the focal length
263     G = np.array([[1, 0, 0],
264                  [0, 1, 0],
265                  [-1/focal_len, 0, 1]])
266     h, w = self.img1.shape
267     # Calculate T: the translation matrix to shift the center of the image
268     T = np.array([[1, 0, -w/2],
269                  [0, 1, -h/2],
270                  [0, 0, 1]])
271     T2 = np.array([[1, 0, w/2],
272                  [0, 1, h/2],
273                  [0, 0, 1]])
274     self.Hp = T2 @ G @ self.R @ T
275     self.Hp = self.Hp / self.Hp[-1, -1]
276
277     # LEFT HOMOGRAPHY #####

```

```

274     M = self.Pp @ np.linalg.pinv(self.P)
275     M = M / M[-1, -1]
276     H_O = self.Hp @ M
277     H_O = H_O / H_O[-1, -1]
278
279     # We then compute z_i, z_i'
280     z, zp = [], []
281     for x, xp in zip(self.points1, self.points2):
282         z.append(H_O @ x)
283         zp.append(self.Hp @ xp)
284     z = np.array(z)
285     zp = np.array(zp)
286     # Normalize z, zp coordinates:
287     z = z / z[:, -1][:, np.newaxis]
288     zp = zp / zp[:, -1][:, np.newaxis]
289
290     # Using these values, we can build a linear system to solve for the left
291     # homography:
292     # A * x = b
293     x = np.linalg.pinv(z) @ zp[:, 0]
294     a, b, c = x
295
296     H_A = np.array([[a, b, c],
297                     [0, 1, 0],
298                     [0, 0, 1]])
299     # Calculate the value of H and normalize it by its last value
300     self.H = H_A @ H_O
301     self.H = self.H / self.H[-1, -1]
302
303 # %% [markdown]
304 # ### Initial estimation of the fundamental matrix, P and P' in canonical form
305
306 # %%
307 #
308 ##### 3.1: Image Rectification
309 img_rectifier = Image_Rectifier(img1_points, img2_points, img_dir+"view1.jpg", img_dir+
310     "view2.jpg")
311 # Get the Fundamentaal Matrix
312 img_rectifier.calculate_fundamental_matrix()
313 print(img_rectifier.F)
314
315 # You can check if your fundamental is correct by checking x'Fx = 0?
316 x = img_rectifier.points1
317 xp = img_rectifier.points2
318 F = img_rectifier.F
319
320 errors = []
321 for x_point, xp_point in zip(x, xp):
322     errors.append(np.abs(xp_point @ F @ x_point))
323 print("Average error: ", np.mean(errors))
324
325 # %% [markdown]
326 # ### Perform non-linear least squares to refine P':
327
328 # %%
329 # Calculate the canonical form of the matrix:
330 img_rectifier.get_cononical_form()
331
332 img_rectifier.refine_proj_prime_matrix_with_nonlinear_leastsquares()
333 img_rectifier.get_refined_fund_matrix_from_Pp()
334
335
336 # %%
337 # You can check if your fundamental is correct by checking x'Fx = 0?
338 x = img_rectifier.points1
339 xp = img_rectifier.points2
340 F = img_rectifier.F
341

```

```

342 errors = []
343 for x_point, xp_point in zip(x, xp):
344     errors.append(np.abs(xp_point @ F @ x_point))
345
346 print("Average error: ", np.mean(errors))
347
348 # %%
349 # After finishing LM estimation of F:
350 img_rectifier.get_homographies_from_refined_F()
351 H, Hp = img_rectifier.H, img_rectifier.Hp
352
353 # %%
354 H_translation = np.array([[1, 0, 150],
355                           [0, 1, 150],
356                           [0, 0, 1]])
357 # H_translation = np.array([[1, 0, 0],
358 #                           [0, 1, 0],
359 #                           [0, 0, 1]])
360
361 # %%
362 img1 = cv2.imread(img_dir+"view1.jpg")
363 img1 = cv2.cvtColor(img1, cv2.COLOR_BGR2RGB)
364 warped_img1 = cv2.warpPerspective(img1, H_translation@H, (400, 600))
365 plt.imshow(warped_img1)
366 plt.axis("off")
367 plt.show()
368
369 # %%
370 img2 = cv2.imread(img_dir+"view2.jpg")
371 img2 = cv2.cvtColor(img2, cv2.COLOR_BGR2RGB)
372 warped_img2 = cv2.warpPerspective(img2, H_translation@Hp, (400, 600))
373 plt.imshow(warped_img2)
374 plt.axis("off")
375 plt.show()
376
377 # %%
378 # Print Point Correspondences on rectified images:
379 # h, w, _ = warped_img1.shape
380 img1_pointshc = [H @ point.hc for point in img1_points]
381 img2_pointshc = [Hp @ point.hc for point in img2_points]
382
383 # Get the combined warped image:
384 combined_image = np.hstack((warped_img1, warped_img2))
385
386 for i, (point1, point2) in enumerate(zip(img1_pointshc, img2_pointshc)):
387     if i>1: #i != 12 and i != 15 and i != 16 and i != 22:
388         x1, y1 = (point1[:2] / point1[2]).astype(np.uint8)
389         x2, y2 = (point2[:2] / point2[2]).astype(np.uint8)
390         # x1, y1 = x1+150, y1+150
391         # y2+= 150
392         # x2 += 150
393         x2 = x2 + warped_img1.shape[1]
394
395         cv2.circle(combined_image, (x1, y1), radius=5, color=(0, 0, 255), thickness=-1)
396         cv2.circle(combined_image, (x2, y2), radius=5, color=(0, 0, 255), thickness=-1)
397
398         line_color = (random.randint(0, 255), random.randint(0, 255), random.randint(0, 255))
399         cv2.line(combined_image, (x1, y1), (x2, y2), color=line_color, thickness=2)
400
401 combined_image = cv2.cvtColor(combined_image, cv2.COLOR_BGR2RGB)
402 plt.figure(figsize=(10, 6))
403 plt.imshow(combined_image)
404 plt.axis("off")
405 plt.show()
406
407 # %% [markdown]
408 # ### 3.2 Interest Point Detection
409
410 # %%
411 #####

```

```

412 # ##### 3.2: Interest Point Detection
413 ######
414 # Open the images:
415 img1_bw = np.expand_dims(open_image_in_grayscale(img_dir+"view1.jpg"), axis=-1)
416 img2_bw = np.expand_dims(open_image_in_grayscale(img_dir+"view2.jpg"), axis=-1)
417 warped_img1_bw = cv2.warpPerspective(img1_bw, H, (400, 600))
418 warped_img2_bw = cv2.warpPerspective(img2_bw, Hp, (400, 600))
419
420 # Run canny operator for corner detection on the image
421 edges = cv2.Canny(image=warped_img1_bw, threshold1=150, threshold2=200, apertureSize=3)
422 edges2 = cv2.Canny(image=warped_img2_bw, threshold1=150, threshold2=200, apertureSize=3)
423 combined_image = np.hstack((edges, edges2))
424 plt.imshow(combined_image)
425 plt.axis("off")
426 plt.show()
427
428 # %%
429 combined_image = np.hstack((edges, edges2))
430 # Use Hough Line Transform to detect straight lines
431 lines = cv2.HoughLinesP(combined_image, 1, np.pi/180, threshold=5, minLineLength=40,
432 maxLineGap=10)
433
434 # Create a mask to draw the detected lines
435 line_mask = np.zeros_like(combined_image)
436
437 # Draw the detected lines on the mask
438 if lines is not None:
439     for line in lines:
440         x1, y1, x2, y2 = line[0]
441         cv2.line(line_mask, (x1, y1), (x2, y2), 255, thickness=2)
442
443 # Inpaint the image to remove the lines
444 img_wlines_removed = cv2.inpaint(combined_image, line_mask, inpaintRadius=3, flags=cv2.
445 INPAINT_TELEA)
446
447 # Show the original, mask, and result images
448 plt.figure(figsize=(15, 5))
449 plt.subplot(1, 3, 1)
450 plt.title("Original Image")
451 plt.imshow(combined_image, cmap="gray")
452 plt.axis("off")
453
454 plt.subplot(1, 3, 2)
455 plt.title("Detected Lines through HoughLinesP")
456 plt.imshow(line_mask, cmap="gray")
457 plt.axis("off")
458
459 plt.subplot(1, 3, 3)
460 plt.title("Result with Lines Removed")
461 plt.imshow(img_wlines_removed, cmap="gray")
462 plt.axis("off")
463
464 plt.tight_layout()
465 plt.show()
466
467 # %%
468 edge_to_remove = edges.copy()
469 # Use Hough Line Transform to detect straight lines
470 lines = cv2.HoughLinesP(edge_to_remove, 1, np.pi/180, threshold=5, minLineLength=40,
471 maxLineGap=10)
472
473 # Create a mask to draw the detected lines
474 line_mask = np.zeros_like(edge_to_remove)
475
476 # Draw the detected lines on the mask
477 if lines is not None:
478     for line in lines:
479         x1, y1, x2, y2 = line[0]
480         cv2.line(line_mask, (x1, y1), (x2, y2), 255, thickness=2)
481
482 # Inpaint the image to remove the lines
483 left_edge = cv2.inpaint(edge_to_remove, line_mask, inpaintRadius=3, flags=cv2.
484 INPAINT_TELEA)

```

```

480
481
482 # %%
483 edge_to_remove = edges2.copy()
484 # Use Hough Line Transform to detect straight lines
485 lines = cv2.HoughLinesP(edge_to_remove, 1, np.pi/180, threshold=5, minLineLength=40,
486 maxLineGap=10)
487
488 # Create a mask to draw the detected lines
489 line_mask = np.zeros_like(edge_to_remove)
490
491 # Draw the detected lines on the mask
492 if lines is not None:
493     for line in lines:
494         x1, y1, x2, y2 = line[0]
495         cv2.line(line_mask, (x1, y1), (x2, y2), 255, thickness=2)
496
497 # Inpaint the image to remove the lines
498 right_edge = cv2.inpaint(edge_to_remove, line_mask, inpaintRadius=3, flags=cv2.
499 INPAINT_TELEA)
500
501 # %%
502 # Extract edge points from the canny edge map:
503 kp1 = np.column_stack(np.where(left_edge > 254))
504 kp2 = np.column_stack(np.where(right_edge > 254))
505
506 # %%
507 # I now need to run NCC on the points in the canny edge detector:
508 # The following code is from my solution for hw4:
509 def get_masked_corners(corners, area_bound, img):
510     return corners[(corners[:, 0] >= area_bound) & (corners[:, 0] <= img.shape[0] -
511     area_bound) &
512                     (corners[:, 1] >= area_bound) & (corners[:, 1] <= img.
513     shape[1] - area_bound)]
514
515 def SSD(m, img1, corners1, img2, corners2, num_matches):
516     area_bound = m+1
517     # Remove boundary points that are outside of the neighborhood search area
518     # Make sure to check for the boundaries on all sides of the images
519     # Corners is a list of points: [(x1,y1), (x2,y2), (x3,y3), ...]
520     masked_corners1 = get_masked_corners(corners1, area_bound, img1)
521     masked_corners2 = get_masked_corners(corners2, area_bound, img2)
522
523     matching_points = []
524     for coord1 in masked_corners1:
525         # Points are (y,x)
526         # Get (m+1) by (m+1) area in the BW image around that corner point
527         neighborhood1 = img1[coord1[0]-(area_bound):coord1[0]+(area_bound), coord1[1]-(area_bound):coord1[1]+(area_bound)]
528         for coord2 in masked_corners2:
529             # Points are meant to be on a horizontal line
530             if np.abs(coord2[0] - coord1[0]) < 4 and np.abs(coord2[1] - coord1[1] < 10):
531                 neighborhood2 = img2[coord2[0]-(area_bound):coord2[0]+(area_bound),
532                                     coord2[1]-(area_bound):coord2[1]+(area_bound)]
533
534             # Calculate the sum of squared distances and add the result to a list to
535             # later get the topk
536             ssd = np.sum((neighborhood1 - neighborhood2)**2)
537             matching_points.append([ssd, [coord1, coord2]])
538
539     # Get the top num_matches corner matches (with the highest SSD values first)#
540     # Convert to array for easier manipulation
541     matching_points = np.array(matching_points, dtype=object)
542     print("Number of matches found: ", matching_points.shape)
543
544     if num_matches > len(matching_points):
545         num_matches = len(matching_points)
546     # Sort matching points by SSD (smallest first)
547     sorted_matches = sorted(matching_points, key=lambda x: x[0])
548     # To store the top correspondences without repeats
549     topk_matches = []
550     used_coords1 = set()
551     used_coords2 = set()

```

```

545     # Iterate over sorted matches and select the top num_matches with unique points
546     for match in sorted_matches:
547         coord1, coord2 = match[1]
548         # Ensure no coordinate from coord1 or coord2 is repeated
549         if tuple(coord1) not in used_coords1 and tuple(coord2) not in used_coords2:
550             topk_matches.append([coord1, coord2])
551             used_coords1.add(tuple(coord1))
552             used_coords2.add(tuple(coord2))
553
554         # # Stop when we have enough matches
555         if len(topk_matches) == num_matches:
556             break
557
558     return topk_matches
559
560 # %%
561 kp_matches = SSD(20, warped_img1, kp1, warped_img2, kp2, num_matches=1000)
562
563 # %%
564 kp_matches = np.array(kp_matches)
565
566 # %%
567 with open("/mnt/cloudNAS3/Adubois/Classes/ECE661/HW9/kp_matches_no_translate.pkl", "wb") as file:
568     pickle.dump(kp_matches, file)
569
570 # %%
571 with open("/mnt/cloudNAS3/Adubois/Classes/ECE661/HW9/kp_matches_no_translate.pkl", "rb") as file:
572     kp_matches = pickle.load(file)
573
574 # %%
575 combined_edges = np.hstack((left_edge, right_edge))
576 for kp_match in kp_matches:
577     y1, x1 = kp_match[0]
578     y2, x2 = kp_match[1]
579     x2 += right_edge.shape[1]
580     cv2.circle(combined_edges, (x1, y1), radius=5, color=(0, 0, 255), thickness=-1)
581     cv2.circle(combined_edges, (x2, y2), radius=5, color=(0, 0, 255), thickness=-1)
582
583     line_color = (random.randint(0, 255), random.randint(0, 255), random.randint(0, 255))
584     cv2.line(combined_edges, (x1, y1), (x2, y2), color=line_color, thickness=2)
585 plt.imshow(combined_edges, cmap="gray")
586 plt.axis("off")
587 plt.show()
588
589 # %% [markdown]
590 # ### 3D Reconstruction:
591
592 # %%
593 # Recreate the points from the rectified images onto the original image
594 # I also remove border points again which helps to remove the final errors. I do this by
595 # applying a mask around the border of the image
596
597 img_1 = cv2.imread(img_dir+"view1.jpg")
598 img_2 = cv2.imread(img_dir+"view2.jpg")
599
600 og_points1, og_points2 = [], []
601 for points in kp_matches:
602     y1, x1 = points[0]
603     y2, x2 = points[1]
604
605     hc1 = np.array([x1, y1, 1])
606     hc2 = np.array([x2, y2, 1])
607
608     og_coord1 = np.linalg.inv(H) @ hc1
609     og_coord2 = np.linalg.inv(Hp) @ hc2
610
611     og_point1 = Point.from_hc(og_coord1)
612     og_point2 = Point.from_hc(og_coord2)
613
614     x1, y1, _ = og_point1.hc

```

```

614     x2, y2, _ = og_point2.hc
615
616     if (y1 < img_1.shape[0] - 30 and x1 < img_1.shape[1] - 30) and (y2 < img_2.shape[0]
617 - 30 and x2 < img_2.shape[1] - 30):
618         og_points1.append(og_point1)
619         og_points2.append(og_point2)
620
621 # Remove the homography on the points so that they apply to the original images:
622 # kp1_points = [Point.from_hc(Point(x, y).apply_homography(np.linalg.inv(H))) for (x, y)
623 #                 in kp1_coords]
624 # kp2_points = [Point.from_hc(Point(x, y).apply_homography(np.linalg.inv(Hp))) for (x, y)
625 #                 in kp2_coords]
626
627 # %%
628 plt.figure(figsize=(16, 12))
629 img_2 = cv2.cvtColor(cv2.imread(img_dir+"view2.jpg"), cv2.COLOR_BGR2RGB)
630 plt.imshow(img_2)
631 plt.axis("off")
632 for i, points in enumerate(og_points2):
633     x, y, _ = points.hc
634     plt.plot(x, y, 'ro', markersize=5) # Plot red circle
635     plt.text(x + 5, y + 5, str(i + 1), color="black", fontsize=10)
636
637 plt.show()
638
639 # Repeat steps 1-4:
640 img_rectifier_final = Image_Rectifier(og_points1, og_points2, img_dir+"view1.jpg",
641                                     img_dir+"view2.jpg")#, warped_img1, warped_img2)
642 # Get the Fundamentaal Matrix
643 img_rectifier_final.calculate_fundamental_matrix()
644 print(img_rectifier_final.F)
645
646 # %%
647 # Calculate the canonical form of the matrix:
648 img_rectifier_final.get_cononical_form()
649
650 img_rectifier_final.refine_proj_prime_matrix_with_nonlinear_leastsquares()
651 img_rectifier_final.get_refined_fund_matrix_from_Pp()
652
653 # %%
654 # You can check if your fundamental is correct by checking  $x'Fx = 0$ ?
655 x = img_rectifier_final.points1
656 xp = img_rectifier_final.points2
657 F = img_rectifier_final.F
658
659 errors = []
660 for x_point, xp_point in zip(x, xp):
661     errors.append(np.abs(xp_point @ F @ x_point))
662
663 print("Average error: ", np.mean(errors))
664
665 # %%
666 # Reconstruct the points:
667 P = img_rectifier_final.P
668 Pp = img_rectifier_final.Pp
669 x = [point.hc for point in og_points1]
670 xp = [point.hc for point in og_points2]
671 world_points = Image_Rectifier.estimate_world_coordinates(P, Pp, x, xp)
672
673 with open("./world_points.pkl", "wb") as file:
674     pickle.dump(world_points, file)
675
676 # %%
677 from matplotlib.patches import ConnectionPatch
678 from mpl_toolkits.mplot3d import proj3d
679
680 # Create a figure
681 fig = plt.figure(figsize=(10, 8))
682 img_points = np.array([point.hc[:2] for point in og_points2])

```

```

683
684 # Plot 2D image with points
685 ax1 = fig.add_subplot(2, 1, 1) # Top subplot
686 ax1.imshow(img_2)
687 ax1.scatter(img_points[:, 0], img_points[:, 1], c='red', s=2, label='2D Points') # Note
688     x, y swap
689 ax1.axis("off")
690
691 # Plot 3D points
692 ax2 = fig.add_subplot(2, 1, 2, projection='3d') # Bottom subplot
693
694 # Plot 3D points in the world frame
695 ax2.scatter(
696     world_points[:, 0],
697     world_points[:, 1],
698     world_points[:, 2],
699     c='b', s=10, label='3D points (world frame')
700 )
701
702 # Draw lines connecting points from 2D to 3D across subplots
703 # This part of the code was based on a Piazza post question 91 on how to plot lines
704     between both figure axes.
705 for i in range(len(img_points)):
706     # Get 2D point coordinates
707     x2d, y2d = img_points[i, 0], img_points[i, 1]
708
709     # Get 3D point in screen coordinates
710     x3d, y3d, _ = proj3d.proj_transform(world_points[i, 0], world_points[i, 1],
711     world_points[i, 2], ax2.get_proj())
712
713     # Create a ConnectionPatch from the 2D point to the 3D screen coordinates
714     con = ConnectionPatch(
715         xyA=(x2d, y2d),
716         coordsA=ax1.transData,
717         xyB=(x3d, y3d),
718         coordsB=ax2.transData,
719         color="green", linestyle="--", linewidth=1
720     )
721     fig.add_artist(con)
722
723 plt.tight_layout()
724 plt.show()
725
726 # %%

```

```

1 # %%
2 import cv2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from tqdm import tqdm
6 import pickle
7
8 # %%
9 def open_image_in_grayscale(img_path):
10     img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
11     return img
12
13 # %%
14 task3_img_dir = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW9/Task3Images/"
15
16 left_img_path = task3_img_dir+"im2.png"
17 right_img_path = task3_img_dir + "im6.png"
18
19 accuracy_maps = []
20 for M in [5, 9, 16, 25]:
21     accuracy_map = []
22     for d_max in [10, 30, 54, 70]:
23         # Maximum disparity from original image is 54
24         window_shape = (M,M)
25
26         left_img = open_image_in_grayscale(left_img_path)
27         right_img = open_image_in_grayscale(right_img_path)

```

```

28
29     # Pad the left and right images such that the windows are created for all image
30     # pixels
31     left_img = np.pad(left_img, M)
32     right_img = np.pad(right_img, M)
33
34     # Create windows from the images. This outputs a shape of (H, W, M, M) for an
35     # MxM window centered on each pixel
36     left_windows = np.lib.stride_tricks.sliding_window_view(left_img, window_shape)
37     right_windows = np.lib.stride_tricks.sliding_window_view(right_img, window_shape)
38 )
39
40     # Get the center pixels for each window:
41     left_center_pixels = left_windows[:, :, M//2][..., np.newaxis, np.newaxis]
42     right_center_pixels = right_windows[:, :, M//2][..., np.newaxis, np.
43 newaxis]
44
45     # Create new windows with only the center pixel values:
46     left_center_windows = np.ones_like(left_windows) * left_center_pixels
47     right_center_windows = np.ones_like(right_windows) * right_center_pixels
48
49     # For each entry in both windows, check if the entry in the img window is
50     # strictly larger than the center pixel value
51     # If strictly greater than, put in a 1, otherwise put in a 0 (right now these
52     # are True and False values)
53     # I do this for entries from (x-d,y) to (x,y)
54     disparity_map = np.zeros_like(left_img)
55
56     for row_idx in range(left_windows.shape[0]):
57         for col_idx in range(left_windows.shape[1]):
58             if col_idx > d_max:
59                 left_window = left_windows[row_idx, col_idx, ...]
60                 left_center_pixel_window = left_center_pixels[row_idx, col_idx, ...]
61                 bitvec_sums = []
62                 for d in range(d_max, 0, -1):
63                     right_window = right_windows[row_idx, col_idx-d]
64                     right_center_pixel_window = right_center_pixels[row_idx, col_idx-
65                     -d, ...]
66
67                     # Now I can compare the two windows
68                     left_bitvec_comparator = left_window > left_center_pixel_window
69                     right_bitvec_comparator = right_window >
70                     right_center_pixel_window
71
72                     # XOR the two comparators. This is treated as a bitwise XOR
73                     # since the bits are represented as logical True/False values
74                     # This results in a list of windows: (M, M)
75                     xor_bitvec = np.logical_xor(left_bitvec_comparator,
76                     right_bitvec_comparator)
77
78                     # I can then get the sum of the bitvector. This is the data cost
79                     bitvec_sums.append(np.sum(xor_bitvec))
80
81
82                     # Calculate the disparity value:
83                     disparity = d_max - np.argmin(bitvec_sums)
84                     # Fill in that location in the disparity map:
85                     disparity_map[row_idx, col_idx] = disparity
86
87                     ground_truth_dmap = open_image_in_grayscale(task3_img_dir+"disp2.png")
88                     # Lower the resolution to match out images:
89                     ground_truth_dmap = ground_truth_dmap.astype(np.float32) / 4
90                     ground_truth_dmap = ground_truth_dmap.astype(np.uint8)
91                     # Pad the ground_truth_dmap as needed:
92                     ground_truth_dmap = np.pad(ground_truth_dmap, M)
93
94                     # Calculate the error (H, W), we only want to compute for non-black (padded)
95                     # pixels
96                     non_black_mask = ground_truth_dmap > 0
97                     diff = np.abs(ground_truth_dmap - disparity_map)
98                     # We also allow for an error of a few pixels between the ground truth and the
99                     # computed disparity map
100                    accuracy_value = np.count_nonzero(diff[non_black_mask] < 3) / np.count_nonzero(
101                    non_black_mask)
102                    accuracy_map.append(accuracy_value)

```

```

88     accuracy_printout = np.zeros_like(diff)
89     accuracy_printout[non_black_mask] = (diff[non_black_mask] < 3) * 255
90
91     plt.imshow(accuracy_printout, cmap="gray")
92     plt.axis("off")
93     plt.show()
94     plt.close()
95
96
97     print(f"M: {M}, Dmax: {d_max}, accuracy: {np.round(accuracy_value * 100, 4)}")
98     accuracy_maps.append(accuracy_map)
99
100 # %%
101 np_accuracy_maps = np.array(accuracy_maps)
102 with open("/mnt/cloudNAS3/Adubois/Classes/ECE661/HW9/accuracy_maps.pkl", "wb") as file:
103     pickle.dump(np_accuracy_maps, file)
104
105 # %%
106 plt.bar(range(5, 30), np_accuracy_maps[:, 3]*100)
107 plt.show()
108
109 # %%
110 plt.imshow(disparity_map, cmap="gray")
111 plt.axis("off")
112 plt.show()
113
114 # %%
115 ground_truth_dmap = open_image_in_grayscale(task3_img_dir+"disp2.png")
116 # Lower the resolution to match out images:
117 ground_truth_dmap = ground_truth_dmap.astype(np.float32) / 4
118 ground_truth_dmap = ground_truth_dmap.astype(np.uint8)
119 # Pad the ground_truth_dmap as needed:
120 ground_truth_dmap = np.pad(ground_truth_dmap, M)
121
122 # Calculate the error, we only want to compute for non-black (padded) pixels
123 non_black_mask = ground_truth_dmap > 0
124 diff = np.abs(ground_truth_dmap - disparity_map)
125 # We also allow for an error of a few pixels between the ground truth and the computed
# disparity map
126 accuracy_value = np.count_nonzero(diff[non_black_mask] < 3) / np.count_nonzero(
    non_black_mask)
127
128 print(np.round(accuracy_value * 100, 4))

```

```

1 import numpy as np
2 import pickle as pkl
3 import matplotlib.pyplot as plt
4 import h5py # for reading depth maps
5
6 """
7 A few notes on the scene_info dictionary:
8 - depth maps are stored as h5 files. Depth is the distance of the object from the camera
# (ie Z coordinate in camera coordinates). The depth map can contain invalid points (
# depth = 0) which correspond to points where the depth could not be estimated.
9 - The intrinsics are stored as a 3x3 matrix.
10 - The poses [R,t] are stored as a 4x4 matrix to allow for easy transformation of points
# from one camera to the other. The resulting transformation matrix is a 4x4 matrix is
# of the form:
11     T = [[R, t]
12           [0, 1]] where R is a 3x3 rotation matrix and t is a 3x1 translation vector.
13 """
14
15 DEPTH_THR = 0.1
16
17 def plot_image_and_depth(img0, depth0, img1, depth1, plot_name):
18     # Enable constrained layout for uniform subplot sizes
19     fig, ax = plt.subplots(1, 4, figsize=(20, 5), constrained_layout=True)
20
21     # Image 0
22     ax[0].imshow(img0, aspect='auto')
23     ax[0].set_title('Image 0')
24     ax[0].axis('off')
25

```

```

26     # Depth 0
27     im1 = ax[1].imshow(depth0, cmap='jet', aspect='auto')
28     ax[1].set_title('Depth 0')
29     ax[1].axis('off')
30     cbar1 = fig.colorbar(im1, ax=ax[1], shrink=0.8, aspect=20)
31     cbar1.ax.yaxis.set_ticks_position('left')
32     cbar1.ax.yaxis.set_label_position('left')
33     cbar1.ax.tick_params(labelsize=15)
34
35     # Image 1
36     ax[2].imshow(img1, aspect='auto')
37     ax[2].set_title('Image 1')
38     ax[2].axis('off')
39
40     # Depth 1
41     im2 = ax[3].imshow(depth1, cmap='jet', aspect='auto')
42     ax[3].set_title('Depth 1')
43     ax[3].axis('off')
44     cbar2 = fig.colorbar(im2, ax=ax[3], shrink=0.8, aspect=20)
45     cbar2.ax.yaxis.set_ticks_position('left')
46     cbar2.ax.yaxis.set_label_position('left')
47     cbar2.ax.tick_params(labelsize=15)
48
49     plt.savefig(plot_name, bbox_inches='tight', pad_inches=0)
50     plt.close()
51
52 if __name__ == "__main__":
53     scene_info = pkl.load(open('./data/scene_info/1589_subset.pkl', 'rb'))
54     for i_pair in range(len(scene_info)):
55         # print(scene_info[i_pair].keys())
56         # ['image0','image1','depth0', 'depth1', 'K0', 'K1', 'T0', 'T1', 'overlap_score']
57         # print(scene_info[i_pair]['image0']) # path to image0
58         # print(scene_info[i_pair]['image1']) # path to image1
59         # print(scene_info[i_pair]['depth0']) # path to depth0
60         # print(scene_info[i_pair]['depth1']) # path to depth1
61         # print(scene_info[i_pair]['K0']) # intrinsic matrix of camera 0 [3,3]
62         # print(scene_info[i_pair]['K1']) # intrinsic matrix of camera 1 [3,3]
63         # print(scene_info[i_pair]['T0']) # pose matrix of camera 0 [4,4]
64         # print(scene_info[i_pair]['T1']) # pose matrix of camera 1 [4,4]
65         # print('-----')
66
67         # read images
68         img0 = plt.imread(scene_info[i_pair]['image0'])
69         img1 = plt.imread(scene_info[i_pair]['image1'])
70         print("Hi, images read")
71         # read depth
72         with h5py.File(scene_info[i_pair]['depth0'], 'r') as f:
73             depth0 = f['depth'][:]
74         with h5py.File(scene_info[i_pair]['depth1'], 'r') as f:
75             depth1 = f['depth'][:]
76
77         # check shapes
78         h0, w0 = img0.shape[:-1]
79         h1, w1 = img1.shape[:-1]
80         assert img0.shape[:-1] == depth0.shape, f"depth and image shapes do not match: {img0}, {depth0}"
81         assert img1.shape[:-1] == depth1.shape, f"depth and image shapes do not match: {img1}, {depth1}"
82
83         # plot image and depth
84         plot_name = f'./pics/image_and_depth_pair_{i_pair}.png'
85         # plot_image_and_depth(img0, depth0, img1, depth1, plot_name)
86
87         print("Images are plotted")
88         #(1) make meshgrid of points in image 0
89         x = np.linspace(10, img0.shape[1]-10, 20) # ignore a border of 10 pixels
90         y = np.linspace(10, img0.shape[0]-10, 20)
91         """
92         <Student code>
93         # meshgrid of x and y coordinates
94         xx, yy = ...
95         """

```

```

96     xx, yy = np.meshgrid(x, y)
97     # make homogeneous coordinates for points0 #[3, N]
98     """
99     <Student code>
100    points0 = ...
101    """
102    points0 = np.vstack((xx.flatten(), yy.flatten(), np.ones_like(xx.flatten()))).T
103
104    #(2) get depth values at points0
105    """
106    <Student code>
107    depth_values0 = ...
108    """
109    x_coords = points0[:, 0].astype(int)
110    y_coords = points0[:, 1].astype(int)
111    depth_values0 = depth0[y_coords, x_coords]
112
113    # remove points with depth 0 (invalid points)
114    """
115    <Student code>
116    valid_points = ...
117    # mask points0 and depth_values0
118    """
119    valid_mask = depth_values0 > 0
120    points0 = points0[valid_mask, :].T
121    depth_values0 = depth_values0[valid_mask]
122
123    # (3) Find the 3D coordinates of these points in camera 0 frame
124    K0 = scene_info[i_pair]['K0'] # [3,3]
125    T0 = scene_info[i_pair]['T0'] # [4,4]
126    """
127    <Student code>
128    # inverse of K0
129    K0_inv = ...
130    # convert points0 to camera coordinates
131    xyz_cam0 = ...
132    # normalize xyz_cam0 to set z = 1 (sanity check)
133    xyz_cam0 = ...
134    # get the point at depth
135    xyz_cam0 = ...
136    # make homogeneous coordinates [4,N]
137    xyz_cam0_hc = ...
138    # convert to world frame [4,N]
139    xyz_world_hc = ...
140    """
141    K0_inv = np.linalg.inv(K0)
142
143    # Get the camera coords
144    xyz_cam0 = (K0_inv @ points0).T
145    xyz_cam0 = xyz_cam0 / xyz_cam0[:, 2][:, np.newaxis]
146    # Get the points at depth:
147    xyz_cam0 = xyz_cam0 * depth_values0[:, np.newaxis]
148
149    # Make homogenous coordinates (and [4,N]):
150    xyz_cam0_hc = np.hstack((xyz_cam0, np.ones((xyz_cam0.shape[0], 1)))).T
151    # Conver to wold frame:
152    xyz_world_hc = T0 @ xyz_cam0_hc
153
154    # (4) Transform these points to camera 1 frame
155    T1 = scene_info[i_pair]['T1']
156    """
157    <Student code>
158    # transform points to camera 1 frame
159    xyz_cam1_hc = ...
160    # convert to camera 1 coordinates
161    xyz_cam1 = ...
162    # get z coordinates for depth check
163    estimated_depth_values1 = ...
164    """
165    # Cam frame 1 and camera coords:
166    xyz_cam1_hc = np.linalg.inv(T1) @ xyz_world_hc
167    xyz_cam1 = xyz_cam1_hc[:3, :] / xyz_cam1_hc[3, :]

```

```

168     estimated_depth_values1 = xyz_cam1[2, :]
169
170     # project to image 1
171     """
172     <Student code>
173     points1 = ... # [3, N]
174     # normalize by dividing by last row
175     points1 = ... # [3, N]
176     # check if points1 are within image bounds
177     ...
178     # get the depth values at these points using the depth map
179     true_depth_values1 = ...
180     """
181     K1 = scene_info[i_pair]["K1"]
182     T1 = scene_info[i_pair]["T1"]
183     points1 = K1 @ xyz_cam1
184     #Normalize by last row:
185     points1 = points1[:, 2, :] / points1[2, :]
186
187     # Create mask for valid points similar to hw3:
188     valid_x = (points1[0, :] >= 0) & (points1[0, :] < img1.shape[1])
189     valid_y = (points1[1, :] >= 0) & (points1[1, :] < img1.shape[0])
190     valid_mask = valid_x & valid_y
191
192     # Get the depth values at the valid points
193     points1 = points1[:, valid_mask]
194     estimated_depth_values1 = estimated_depth_values1[valid_mask]
195     true_depth_values1 = depth1[points1[1, :].astype(int), points1[0, :].astype(int)]
196
197     # (5) plot matching points in image 0 and image 1 with depth check such that the
198     # depth values match
199     fig, ax = plt.subplots(1, 1, figsize=(10, 5))
200
201     # Horizontally stack the images
202     combined_img = np.ones((max(img0.shape[0], img1.shape[0]), img0.shape[1] + img1.
203     shape[1], 3), dtype=np.uint8) * 255
204     combined_img[:img0.shape[0], :img0.shape[1]] = img0
205     combined_img[:img1.shape[0], img0.shape[1]:] = img1
206
207     ax.imshow(combined_img, aspect='auto')
208     ax.scatter(xx, yy, c='r', s=5)
209     ax.set_title('Matching points in Image 0 and Image 1')
210     ax.axis('off')
211
212     # draw lines between matching points
213     for i in range(points1.shape[1]):
214         # if depth values match
215         if np.abs(estimated_depth_values1[i] - true_depth_values1[i]) < DEPTH_THR
216         and true_depth_values1[i] != 0:
217             ax.plot([points0[0,i], points1[0,i] + img0.shape[1]],
218                     [points0[1,i], points1[1,i]], 'g')
219
220     plt.savefig(f'../pics/depth_check_pair_{i_pair}.png', bbox_inches='tight',
221     pad_inches=0)
222     plt.close()
223     print(f"Done with pair {i_pair}")
224
225     # (6) Plot all 3D points for the pair
226     """
227     <Student code>
228     ...
229     """
230     fig = plt.figure(figsize=(10, 10))
231     ax = fig.add_subplot(111, projection='3d')
232
233     # Plot 3D points in the world frame
234     ax.scatter(
235         xyz_world_hc[0, :],
236         xyz_world_hc[1, :],
237         xyz_world_hc[2, :],
238         c='b', s=10, label='3D points (world frame)')
239
240     # Save the plot

```

```
236     plt.savefig(f'../pics/3D_points_pair_{i_pair}.png', bbox_inches='tight',
237         pad_inches=0)
            plt.close()
```