

# ECE 66100 Homework #10

by

Adrien Dubois (dubois6@purdue.edu)

December 6, 2024

## Contents

<b>1 Theory Questions:</b>	<b>1</b>
1.1 Question 1: Overfitting to training data . . . . .	1
1.2 Question 2: Reparametrization Trick . . . . .	2
<b>2 PCA:</b>	<b>2</b>
<b>3 LDA:</b>	<b>3</b>
<b>4 Autoencoders:</b>	<b>4</b>
<b>5 Cascading AdaBoost Classifiers:</b>	<b>4</b>
5.1 Data Preprocessing . . . . .	4
5.2 Weak Classifier Class: . . . . .	4
5.3 Strong Classifier: . . . . .	5
5.4 Strong classifier Cascade: . . . . .	6
5.5 Adaboost Testing: . . . . .	6
<b>6 Results:</b>	<b>6</b>
6.1 Classification Accuracy as a function of p: . . . . .	6
6.2 UMAP plots for PCA, LDA and Autoencoder: . . . . .	7
6.3 Confusion Matrix for PCA, LDA and Autoencoder: . . . . .	9
6.4 Comparison: . . . . .	9
6.4.1 PCA vs LDA: . . . . .	9
6.4.2 Autoencoders vs PCA & LDA: . . . . .	10
6.5 AdaBoost results: . . . . .	10
6.5.1 Training Results: . . . . .	10
6.5.2 Testing Results: . . . . .	10
<b>7 Source Code Listing:</b>	<b>11</b>

## 1 Theory Questions:

### 1.1 Question 1: Overfitting to training data

*What is overfitting to training data and why is it important to control.* My understanding of overfitting is when the model starts to memorize the training data compared to learning the overall distribution. This can often occur because the training dataset is too small, or not representative of the overall distribution of the data. In this way, it is very important to consistently measure both the training loss and evaluation loss. Comparing these two can give a good indication of when to stop training the model so that it does not overfit to the training data, since the true accuracy measurement should always be performed on a significant amount (at least 10% of the total dataset) of unseen testing samples.

## 1.2 Question 2: Reparametrization Trick

*How do you understand the reparametrization trick used in variational autoencoders.* The reparametrization trick is required when using variational autoencoders since we cannot backpropagate through the distribution parametrized by the  $(\mu, \sigma)$  parameters. Instead, we sample  $\mathbf{z}$  from a standard normal distribution with  $(\mu = 0, \sigma = 1)$  and compute the latent space variable as:

$$\mathbf{z} = \mu_{\text{learned}} + \sigma_{\text{learned}} \times \mathbf{z}$$

This separates the randomness in the latent space from the learned parameters, so that we can backpropagate solely through the deterministic parameters, instead of not being able to backpropagate through a randomly sampled distribution. This is required because the sampling operation is non-differentiable. So, with this reparametrization trick, backpropagation of the gradients is possible and we can use the reconstruction and KL divergence losses to learn the best parameters.

## 2 PCA:

The process for principle component analysis is as follows:

- Vectorize each black and white image such that it goes from a shape of  $(H, W, 1)$  array to  $(H \times W)$ .
- We then normalize the vector by dividing each image vector by its L2 Norm.

$$\vec{x}_i = \frac{x_i}{\|x_i\|}$$

- We can then compute the global mean from each image and subtract it from the image vector:

$$\vec{m} = \frac{1}{N} \sum_{i=1}^N \vec{x}_i$$

$$X = [x_0 - \vec{m} \| x_1 - \vec{m} \| \dots \| x_N - \vec{m} \|]$$

- Next, we know that the covariance matrix is  $C = X^T X$ . However, for large image vectors, with matrix would be much too large when calculating the eigen decomposition of  $C$ . Therefore, we estimate this matrix from a smaller covariance matrix:  $C = X X^T$ .
- Decompose the  $C$  matrix into its eigenvalues  $\lambda$ , and eigenvectors  $\mathbf{v}$ .
- Retain the top  $p$  eigenvectors according to the largest eigenvalues. This allows us to create our projection matrix. Note that since `np.linalg.eig` already sorts the eigenvectors for you, I return the top  $p$  eigenvectors through indexing in the  $\mathbf{v}$  array from the end.

$$W_p = X^T \times [\mathbf{v}_{-p} \| \mathbf{v}_{-p+1} \| \dots \| \mathbf{v}_{-1}]$$

- Normalize the projection matrix:

$$\hat{W}_p = \frac{W_p}{\|W_p\|}$$

- Lastly, our feature vectors are calculated as follows:

$$y_i = X \times \hat{W}_p$$

Then we train a classifier using the PCA embeddings for each image. I did this by storing each embedding vector and its corresponding class into an array. I can then search for the closest embedding to a probe test image PCA embedding and assign it the class of that nearby train embedding. This technique was utilized for all such testing in PCA, LDA and autoencoders for this report.

### 3 LDA:

The idea behind linear discriminant analysis is to find the directions in the underlying latent space that maximally discriminate between the classes. This is done through two metrics: the within and between class class scatters. We can calculate the  $p$  dimensional LDA embedding as follows:

- Compute the overall mean of all image vectors:

$$\vec{m} = \frac{1}{N} \sum_{i=1}^N x_i$$

- Compute the per-class means:

$$\vec{m}_c = \frac{1}{N_c} \sum_{i=1}^{N_c} x_i^c$$

- The between class scatter is defined as:

$$S_B = \frac{1}{N} \sum_{i=1}^N \{(\mathbf{m}_c - \mathbf{m})(\mathbf{m}_c - \mathbf{m})^T\}$$

- On the other hand, the within class scatter is defined as:

$$S_W = \frac{1}{N} \sum_{c=1}^N \frac{1}{|N_c|} \sum_{k=1}^{|N_c|} \{(\mathbf{x}_k^c - \mathbf{m}_c)(\mathbf{x}_k^c - \mathbf{m}_c)^T\}$$

- It is important to note for debugging purposes that both the  $S_W$  and  $S_B$  matrices are of shape: (feature dimension, feature dimension)
- Our goal with LDA is to maximize the Fischer Discriminant function:

$$J(\mathbf{w}) = \frac{\mathbf{w}^T S_B \mathbf{w}}{\mathbf{w}^T S_W \mathbf{w}}$$

- Instead of directly optimizing over this cost, it is equivalent to instead find the few largest eigenvalues of the matrix  $S_W^{-1} S_B$ . However, there is a high probability that the  $S_W$  matrix is singular and therefore non-invertible. We therefore need to use Yu and Yang's approach for LDA calculation to get around this problem.
- We first calculate the eigenvalues  $\lambda$ , and eigenvectors  $\mathbf{v}$  of the  $S_B$  matrix.
- We then remove all eigenvalues that are near 0, and their corresponding eigenvectors. The eigenvectors are also normalized columnwise.
- Next, we create a diagonal matrix of the eigenvalues  $D_B$  and calculate the low dimensional projection matrix of  $S_B$ :

$$Z = [\hat{v}_0 \| \hat{v}_1 \| \dots \| \hat{v}_n] D_B^{-0.5}$$

- We then require the following eigen decomposition:

$$\text{eig-values, } U = \text{eig-decomp} (Z^T S_W Z)$$

- Retain the top  $p$  columns of  $U$  and normalize these vectors.
- The projection matrix  $W_p$  is then equal to:

$$W_p = (\hat{U}^T Z^T)$$

- and the final feature vector is equal to:

$$y_i = (\hat{x}_i - \vec{m}) \times W_p$$

## 4 Autoencoders:

An autoencoder works through a pair of models: the encoder and decoders. The encoder's role is to project the initial image vector into a lower dimensional space, while the decoder will reproject the latent space vector to the original image dimensionality. You can then train both models together to recreate the original image values in the decoder output. This will also push the encoder to learn the optimal latent space representation that maintains the most information of the original image so that it can be reconstructed by the decoder.

For this assignment, I have tested and demonstrate results using a latent space dimension of 3, 8 and 16. I also trained autoencoders with latent space dimensions from 1 to 16 and will be reporting accuracy for each such model, but will not include the UMAP representation, nor the confusion matrices. The accuracy for each model was also calculated using the same approach described in the PCA section.

## 5 Cascading AdaBoost Classifiers:

Instead of including a step by step theory based approach for explaining the Adaboost classifiers, I will use a functional/class based approach that is more closely related to the actual programming task. Please refer to other previous solutions such as 2020, solution 1 by Brian Helfrecht for a more theoretical step by step approach without any bias to a certain programming style.

### 5.1 Data Preprocessing

The purpose of the data preprocessing is to take input images of shape (128, 128) into low level feature vectors. I did this as follows:

1. Convert every image to  $(64 \times 64)$ , grayscale pytorch tensors or numpy arrays.
2. We then need to apply Haar filters of sizes 0, 2, 4, 6, 8, ..., 64. For each "Haar size", I am referring to the dimension of the convolutional kernel applied when performing haar filtering. For example, a Haar size of 6 would have two kernels:

$$\text{haar\_dx} = \begin{bmatrix} -1 & -1 & -1 & 1 & 1 \end{bmatrix}, \quad \text{haar\_dy} = \text{haar\_dx}^T$$

We can then horizontally stack the features extracted each horizontal and vertical filter to construct a high dimensional feature representation of the image. It is important to note that by a Haar size of 0, I also include the actual image in the feature vector.

This pre-processing procedure is applied to the training and testing images. It is also important to note that I create a very large vertically stacked array of all the positive samples and negative samples in each data set partition. I also randomize these arrays and their corresponding labels (1 for positive sample, -1 for negative samples) accordingly to ensure realistic training and testing.

### 5.2 Weak Classifier Class:

- For each img in the array of image feature vectors, we associate a weight. This weight is initialize as a uniform distribution across all samples.
- Next, for each feature vector (a column in the (num-samples, feature-dim) img array), we do the following:
  - Sort the feature column. Also sort the weight matrix, and labels according to the ascending sorting procedure on the features.
  - We then define four metrics used to qualify the classification error for both polarities. I remind you that this weak classifier works by assuming all features below a certain threshold are class 1, and the rest are class 2. So the polarity will switch whether the first class refers to the positive or negative samples. The metrics used are:
    - \*  $S^+$  The cumulative sum of all weights for positive images whose value is less than the current threshold.
    - \*  $S^-$  The cumulative sum of all weights for negative images whose value is less than the current threshold.

- \*  $T^+$  The sum of the weights for positive images
- \*  $T^-$  The sum of the weights for negative images
- We can then calculate the errors for each polarity:

$$e^{+1} = S^+ + T^- - S^-$$

$$e^{-1} = S^- + T^+ - S^+$$

- The four metrics above combine into the overall prediction error:  $\epsilon = \min(e^{+1}, e^{-1})$
- So, depending on which error was the minimum, we associate the polarity accordingly. We also store the feature index and feature value that created the minimum error.
- This procedure is repeated for every feature vector (column) in the image matrix. The optimal values for the feature, threshold, polarity and error are what define a weak classifier.

### 5.3 Strong Classifier:

A strong classifier is defined by one cascade of weak classifiers.

- We first initialize one weak classifier.
- We then calculate the following parameters:

$$\beta = \frac{\epsilon}{1 - \epsilon}$$

$$\alpha = \ln\left(\frac{1}{\beta}\right)$$

- We then have to calculate our predicted labels through our classifier by thresholding the features based on what you calculated for the weak classifier. Make sure to take into account the polarity when assigning labels.
- We can then update each weight in the weight matrix as follows, where  $\delta$  equals 1 if the prediction was incorrect, and 0 if the prediction was correct. Make sure to also normalize the updated weights so that they still sum up to 1.

$$w_i = \text{old}(w_i) \times \beta^{1-\delta}$$

- Next, we need to calculate the accuracy of the weak classifiers so far. In this way, we loop through each weak classifier and use the alpha parameter as a scalar that informs us how much each classifier impacts the final result. Therefore:

$$\text{label}(x) = \sum_{t=0}^{T-1} \alpha_t \times \text{pred}_t$$

where  $T$  is the total number of weak classifiers, and  $\text{pred}_t$  is a one if  $\text{feature} * \text{polarity} \geq \text{threshold} * \text{polarity}$  else it is a -1.

- We can then get the sign of each output to get the final predicted label.
- Next we need to compare the predicted labels to the ground truth labels to calculate the accuracy.
- If the accuracy is 100%, we can stop adding to the cascade. Otherwise, we keep adding new weak classifiers up to a user specified limit.

## 5.4 Strong classifier Cascade:

When one strong classifier is not enough to reduce the false positive rate to below 1%, we need to keep adding new cascades. This creates a cascade of strong classifiers.

- In this way, we first define a strong classifier.
- We then get the predicted labels in the same way that was performed for a strong classifier.
- Now, instead of calculating accuracy, we calculate the false positive rate. If this rate is less than 1% we stop our strong classifier cascade.
- If the false positive rate is not low enough, we need to update our training data. In this way, we look for correct nagtive predictions and remove them from the image and label arrays. This will help us hone in on reducing the false positive rate with a mix of positive samples, and incorrectly predicted negative samples.
- Lastly, regardless of the false positive rate, if we have removed all the incorrectly predicted negative images from the image array, or if we have reach a user defined limit, we stop adding new strong classifier cascades.

## 5.5 Adaboost Testing:

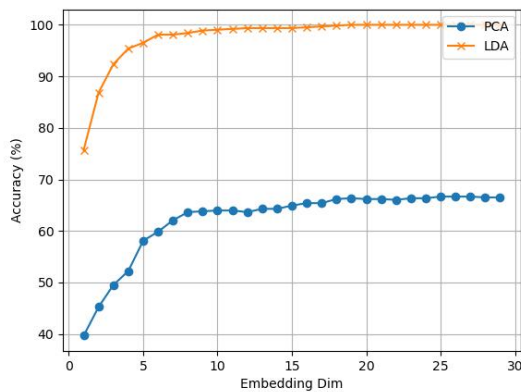
For testing Adaboost:

- I first loop through each of the cascades to calculate the predicted labels.
- For each cascade, I compute the true negative, true positive, false positive and false negative predictions and report the false positive rate and false negative rate.
- Additionally, to get the overall prediction accuracy, I keep track of the true negatives and false positives in an array, and I use a logical or operation over each cascade for both of these arrays to get the final predictions. This works because the initial predictions are only accurate on a subset of the data (due to the data removal process used during training). So we must take into account what each cascade predicts overall to get the final accuracy.


## 6 Results:

### 6.1 Classification Accuracy as a function of p:

Included below are graphs of the accuracy of each technique with respect to p, the dimension of the feature vector reduced representation.




(a) PCA and LDA with Yuyang accuracy graph




(b) Autoencoder accuracy graph

## 6.2 UMAP plots for PCA, LDA and Autoencoder:


Please note that, for this section, I only plot a subset of the overall testing data in the UMAP plots due to the requirement to balance out my classes before running the UMAP fit transform. Therefore, when testing my PCA output for low values of  $p$  some classes were not present in the testing data and are therefore not included in the UMAP plot to allow other classes to be printed normally. I did this by only printing classes that have at least 11 samples in the testing data (only for done PCA).




(a) PCA UMAP display of train embeddings with  $p=3, 8, 16$




(b) PCA UMAP display of test embeddings with  $p=3, 8, 16$




(a) LDA UMAP display of train embeddings with  $p=3, 8, 16$



(b) LDA UMAP display of test embeddings with  $p=3, 8, 16$




(a) Autoencoder UMAP display of train embeddings with  $p=3, 8, 16$




(b) Autoencoder UMAP display of test embeddings with  $p=3, 8, 16$




### 6.3 Confusion Matrix for PCA, LDA and Autoencoder:



(a) Confusion matrix of PCA with p=3, 8, 16



(b) Confusion matrix of LDA with p=3, 8, 16



(c) Confusion matrix of AutoEncoder testing with p=3, 8, 16

### 6.4 Comparison:

#### 6.4.1 PCA vs LDA:

Using the computation trick, PCA was much faster than LDA with Yuyang for dimensionality reduction. This was due to the need to calculate the  $S_B$  and  $S_W$  parameters. However, as shown in the accuracy graphs, LDA had much higher accuracy than PCA. Not only did its accuracy rise much more sharply than PCA, but it also had a higher starting point, and a higher finishing accuracy of 100%. On the other hand, PCA was much less accurate than LDA with a maximum accuracy of around 67% over the embedding dimensions analyzed. It additionally suffered to class imbalance in the testing predictions compared to LDA which was able to more accurately span all classes. I also internally compared these results to sklearn's PCA fit transform function and found that the results were consistent with mine making me believe that this issue is inherent to PCA for this dataset with 64 by 64 images.

### 6.4.2 Autoencoders vs PCA & LDA:

The autoencoder outperformed PCA and LDA on almost all metrics. Its accuracy was much higher across the board, and the autoencoder was able to reach 100% accuracy with an embedding dimension as low as 9, while LDA required an embedding dimension of 18 to do the same. Additionally, the UMAP displayed embeddings show much greater separation compared to PCA and LDA which would imply that the retained features in the latent space are more discriminative compared to PCA and LDA. The only downside of the autoencoder was the need for a GPU to train the model efficiently, and the fact that the training still took much longer than LDA and PCA even when training such a small model on an Nvidia A5000 with a large batch size. Techniques such as early stopping could have been utilized to reduce the training time once the loss flattened out. Additionally, reporting a validation loss with the training loss would have helped to find the optimal stopping time before the model overfit to the training data.

## 6.5 AdaBoost results:

### 6.5.1 Training Results:

The following data are the results of training my AdaBoost cascades with up to 5 weak classifiers per cascade and up to 3 cascades. While I am aware that the instructions asked for the data to be represented graphically, I have tabulated it instead as I believe it more clearly displays the results.

Overall my results showed:

- Accuracy: **96.7045**
- Best Strong Classifier's False Negative Rate: 0.08
- Best Strong Classifier's False Positive Rate: 0.05

Cascade ID	False Negative Rate	False Positive Rate	Final Accuracy
1	0.07	0.04	0.89
2	0.01	0.08	0.91
3	0.00	0.07	0.93

Table 1: False positive rate at each cascade

Weak Classifier ID	Feature	Threshold	Polarity	Error	Alpha	Accuracy
1	6993	-1.099	-1	0.139	1.827	0.861
2	1500	0.227	-1	0.226	1.231	0.861
3	5694	-0.041	-1	0.281	0.940	0.886
4	6747	-0.015	1	0.282	0.933	0.880
5	5697	0.040	1	0.325	0.731	0.893
1	769	0.137	1	0.104	2.150	0.894
2	6347	0.095	1	0.287	0.909	0.894
3	4720	0.000	-1	0.300	0.848	0.894
4	637	0.243	1	0.302	0.840	0.913
5	4267	-0.073	-1	0.302	0.838	0.906
1	7252	-0.319	-1	0.076	2.495	0.924
2	5758	-0.054	-1	0.279	0.949	0.924
3	4718	-0.005	-1	0.299	0.850	0.924
4	133	0.243	1	0.284	0.924	0.928
5	2851	0.125	-1	0.284	0.927	0.934

Table 2: Parameters for each weak classifier in the cascade and the corresponding classification accuracy

### 6.5.2 Testing Results:

Overall accuracy of 96.7045%.

Cascade ID	False Negative Rate	False Positive Rate
1	0.95	0.05
2	0.13	0.87
3	0.08	0.92

Table 3: False positive rate at each cascade

As you can see, due to our training strategy of removing data after each cascade, the different cascades were specialized at separate different types of data. In this way, my first cascade was very powerful at detecting true negatives as it already detected 417 out of the 440 true negatives through the first pass. However, the first cascade only classified 23 out of the 440 possible true positives. On the other hand, the final cascade was very strong at separating out the remaining positive samples with 404 out of 440 positive samples detected, and much weaker at detecting negative samples. This is reflected in the shifting false positive and false negative rates as well.

## 7 Source Code Listing:

```

1  # %%
2  import os
3  import numpy as np
4  import torch
5  import umap
6  from PIL import Image
7  from torch.utils.data import Dataset, DataLoader
8  from torchvision import transforms
9  import matplotlib.pyplot as plt
10 from sklearn.metrics import confusion_matrix
11 import seaborn as sns
12
13 # %%
14 #####
15 # Change these
16 p = 3 # [3, 8, 16]
17 training = False
18 TRAIN_DATA_PATH = '/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/train/'
19 EVAL_DATA_PATH = '/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/test/'
20 LOAD_PATH = f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/model_{p}.pt"
21 OUT_PATH = '/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/exp/'
22 #####
23
24 # %%
25 class DimReducerBuilder(Dataset):
26     def __init__(self, path, option):
27         self.path = path
28         self.image_list = [f for f in os.listdir(path) if f.endswith('.png')]
29         self.label_list = [int(f.split('_')[0]) for f in self.image_list]
30         self.len = len(self.image_list)
31         if option == "bw":
32             self.aug = transforms.Compose([
33                 transforms.Resize((64, 64)),
34                 transforms.Grayscale(num_output_channels=1),
35                 transforms.ToTensor(),
36             ])
37         else:
38             self.aug = transforms.Compose([
39                 transforms.Resize((64, 64)),
40                 transforms.ToTensor(),
41             ])
42
43     def __len__(self):
44         return self.len
45
46     def __getitem__(self, index):
47         fn = os.path.join(self.path, self.image_list[index])
48         x = Image.open(fn).convert('RGB')
49         x = self.aug(x)
50
51     # Flatten and normalize the image data:

```

```

52     img_vec = torch.reshape(x, shape=(x.shape[0], -1))
53     img_vecn = img_vec / torch.norm(img_vec, dim=1, keepdim=True)
54     # img_vecn = img_vec / img_vec.shape[1]
55     return {'img_vecn': img_vecn.squeeze(), 'y': self.label_list[index]}
56
57
58 # %% [markdown]
59 # ### Task 1: PCA
60
61 # %%
62 from sklearn.decomposition import PCA
63
64 # %%
65 def get_id_of_nearest_embedding(training_set, probe_embedding, num_classes=30,
66                                num_samples=21):
67     index_to_class_id = np.repeat(np.arange(num_classes), num_samples)
68     training_set = np.reshape(training_set.copy(), newshape=(training_set.shape[0]*
69 training_set.shape[1], -1))
70     # We first calculate the euclidean distance between the probe and all trained
71     embeddings:
72     distances = np.linalg.norm(training_set - probe_embedding, axis=1)
73     nearest_neighbor_idx = np.argmin(distances)
74
75     # Then we return the index with the smallest distance
76     return index_to_class_id[nearest_neighbor_idx]
77
78 # %%
79 def get_PCA_vecs(img_vec, p, speedup=True):
80     overall_mean = np.mean(img_vec, axis=0)
81
82     # Subtract by the mean:
83     img_meaned = img_vec - overall_mean
84
85     # PCA:
86     # Return p eigenvectors from the covariance matrix:
87     # Calculate the covariance matrix:
88     if speedup == False:
89         print("IMg_meaningn", img_meaned.shape)
90         _, _, Vt = np.linalg.svd(img_meaned @ img_meaned.T, full_matrices=False)
91
92         W_p = img_meaned.T @ Vt
93
94         # Normalize the principal components
95         W_p = W_p / np.linalg.norm(W_p, axis=1, keepdims=True)
96
97         # Extract the top-p right singular vectors (principal components)
98         W_phat = W_p[:, :p] # Shape is (CHW, p)
99     else:
100         C = img_meaned @ img_meaned.T # Shape is (B, B)
101
102         # Take the eig vecs of this, they are already in reverse order:
103         _, eigvec_c = np.linalg.eigh(C)
104
105         W_p = img_meaned.T @ eigvec_c # Shape is (CHW, p)
106         W_phat = W_p / np.linalg.norm(W_p, axis=1, keepdims=True) # Normalized
107         W_phat = W_phat[:, -p:]
108
109     # Project the image vectors into the p dimensional space
110     pca = img_meaned @ W_phat
111     return pca
112
113 # %% [markdown]
114 # ### Task 1: LDA
115
116 # %%
117 def calculate_global_parameters_for_LDA(train_loader, num_classes=30, feature_dim=4096):
118     num_classes = 30
119     overall_mean = 0
120     class_means = np.zeros((num_classes, feature_dim))
121     num_samples = 0
122     class_sample_counts = np.zeros((num_classes))
123
124     for batch in train_loader:

```

```

122     img_vecn = batch["img_vecn"].numpy()
123     labels = batch["y"] # Shape is (B,)
124
125     # Accumulate means:
126     overall_mean += np.sum(img_vecn, axis=0) # shape is: (4096) -> (CHW)
127     num_samples += img_vecn.shape[0]
128
129     # Set up the dataset wide information required
130     for label in np.unique(labels):
131         # idx of class_means = label + 1
132         class_samples = img_vecn[labels == label]
133         class_sample_counts[label - 1] += class_samples.shape[0]
134         class_means[label - 1] += np.sum(class_samples, axis=0)
135
136     # Finalize the means:
137     overall_mean = np.array(overall_mean / num_samples) # Shape: (4096,)
138     for label in range(num_classes):
139         class_means[label] /= class_sample_counts[label] # Shape: (4096,)
140
141     # Compute S_B, the outer product for the between-class scatter:
142     S_B = np.zeros((feature_dim, feature_dim)) # Shape is (4096, 4096)
143     for i in range(num_classes):
144         mean_diff = np.expand_dims(class_means[i] - overall_mean, axis=1) # Shape is
(4096, 1)
145         S_B += mean_diff @ mean_diff.T
146
147     # Compute Within-Class Scatter Matrix (S_W)
148     S_W = np.zeros((feature_dim, feature_dim)) # Shape is (4096, 4096)
149     for batch in train_loader:
150         img_vecn = batch["img_vecn"].numpy()
151         labels = batch["y"] # Shape is (B,)
152
153         # Set up the dataset wide information required
154         for label in np.unique(labels):
155             # idx of class_means = label + 1
156             class_samples = img_vecn[labels == label]
157
158             for sample in class_samples:
159                 diff = np.expand_dims((sample - class_means[label - 1]), axis=1) #
Shape: (4096, 1)
160                 S_W += diff @ diff.T # Outer product shape is (4096, 4096)
161     return S_B, S_W
162
163 # %%
164 def get_projection_matrix(S_W, S_B, p, lda_option, num_labels=30):
165     if lda_option == "YUYANG":
166         # Retain eigvecs where the values are not close to 0
167         eigvals, eigvecs = np.linalg.eigh(S_B)
168         idx = eigvals > 1e-6 # Filter eigenvalues
169         top_eigvals, top_eigvecs = eigvals[idx], eigvecs[:, idx]
170
171         # Normalize the eigenvectors:
172         sb_eigvecn = top_eigvecs / np.linalg.norm(top_eigvecs, axis=1, keepdims=True) #
Shapes is (CHW, K_Y)
173
174         eig_val_mat = np.diag(top_eigvals) #np.eye(num_labels - 1) * top_eigvals
175         # We can then construct a low dimensional projection of S_B with sb_eigvecn
176         D_B = np.sqrt(np.linalg.inv(eig_val_mat))
177         Z = np.dot(sb_eigvecn, D_B)
178
179         # Use eigendecomposition to diagonalize Z
180         _, U = np.linalg.eigh(Z.T @ S_W @ Z)
181         # Get the top eigenvectors and normalize them
182         U_top = U[:, -p:]
183         U_topn = U_top / np.linalg.norm(U_top, axis=1, keepdims=True)
184
185         # Generate the projection matrix
186         proj_mat = (U_topn.T @ Z.T).T
187     else:
188         # LDA Optimization
189         # Get the eigenvalue/vectors of S_W^-1 S_B
190         _, eigvecs = np.linalg.eig(np.linalg.inv(S_W) @ S_B)
191

```

```

192         # Get the top eigvecs since np already sorts them
193         proj_mat = eigvecs[:, -p:] # Columns are the eigenvectors
194
195     return proj_mat
196
197 # %% [markdown]
198 # ### Nearest neighbor classifier:
199
200 # %%
201 # Run through training dataset and create the mean embedding for all the images
202 # belonging to that class
203 def train_classifier(train_loader, lda_proj_mat, dim_reducer, p, num_classes=30):
204     class_embs = [[] for _ in range(num_classes)]
205
206     for batch in train_loader:
207         img_vecn = batch["img_vecn"].numpy()
208         labels = batch["y"] # Shape is (B,)
209
210         if img_vecn.shape[0] >= p:
211             if dim_reducer == "PCA":
212                 embs = get_PCA_vecs(img_vecn, p)
213
214             elif dim_reducer == "LDA":
215                 embs = img_vecn @ lda_proj_mat # Shape is (B, p)
216
217             else:
218                 raise ValueError("Wrong input type: dim_reducer should be PCA or LDA")
219
220             # Train Classifier embeddings
221             for label in np.unique(labels):
222                 # idx of class_means = label + 1
223                 class_embedding = embs[labels == label]
224
225                 for sample in class_embedding:
226                     class_embs[label - 1].append(sample)
227
228     return np.array(class_embs).astype(np.float32)
229
230 # %%
231 def run_testing_script(test_loader, lda_proj_mat, class_embs, dim_reducer, num_classes
232 =30, num_samples=21):
233     predicted_label = []
234     true_label = []
235     test_embs_list = [[] for _ in range(num_classes)]
236
237     for batch in test_loader:
238         img_vecn = batch["img_vecn"].numpy()
239         labels = batch["y"] # Shape is (B,)
240
241         if img_vecn.shape[0] >= p:
242             if dim_reducer == "PCA":
243                 embs = get_PCA_vecs(img_vecn, p)
244
245             elif dim_reducer == "LDA":
246                 # img_vecn = img_vecn / np.linalg.norm(img_vecn, axis=1, keepdims=True) #
247                 # Shapes is (B, CHW)
248                 # The projection matrix is just the top eigenvectors?
249                 embs = img_vecn @ lda_proj_mat # Shape is (B, p)
250
251             else:
252                 raise ValueError("Wrong input type: dim_reducer should be PCA or LDA")
253
254             # Compare nearest embeddings to get predicted label
255             for embedding, label in zip(embs, labels):
256                 embedding = np.array(np.expand_dims(embedding, axis=0), dtype=np.float32
257 )
258                 index = get_id_of_nearest_embedding(class_embs, embedding, num_classes,
259 num_samples)
260
261                 test_embs_list[index.item()].append(embedding)
262                 predicted_label.append(index.item() + 1)
263                 true_label.append(label)

```

```

260     return np.array(predicted_label, dtype=np.float32), np.array(true_label, dtype=np.
261         float32), test_embs_list
262 # %%
263 accuracy_list_lda = []
264 # %%
265 # batch_size = 630
266 # num_classes = 30
267 # dim_reducer = "LDA"
268 # lda_option = "YUYANG"
269 # train_loader = DataLoader(dataset=DimReducerBuilder(TRAIN_DATA_PATH, option="bw"),
270     batch_size=batch_size, shuffle=False)
271 # S_B, S_W = calculate_global_parameters_for_LDA(train_loader, num_classes=num_classes,
272     feature_dim=4096)
273 # for p in range(3, 32):
274 #     if p % 24 == 0:
275 #         print(p)
276 #         lda_proj_mat = get_projection_matrix(S_B=S_B, S_W=S_W, p=p, lda_option=lda_option)
277 #         # Run through test set and find the nearest embedding and assign that label to the
278 #         image
279 #         class_embs = train_classifier(train_loader, dim_reducer=dim_reducer, p=p,
280             num_classes=num_classes, lda_proj_mat=lda_proj_mat)
281 #         test_loader = DataLoader(dataset=DimReducerBuilder(EVAL_DATA_PATH, option="bw"),
282             batch_size=batch_size, shuffle=False)
283 #         pred_labels, true_labels, test_embs = run_testing_script(test_loader, lda_proj_mat
284             , class_embs, dim_reducer=dim_reducer)
285 #         accuracy = np.count_nonzero(pred_labels == true_labels) / len(pred_labels)
286 #         accuracy_list_lda.append(np.round(accuracy * 100, 4))
287 # %%
288 p = 16 # [3, 8, 16]
289 batch_size = 630
290 num_classes = 30
291 # %%
292 dim_reducer = "PCA"
293 lda_option = "YUYANG" # "YUYANG"
294 train_loader = DataLoader(dataset=DimReducerBuilder(TRAIN_DATA_PATH, option="bw"),
295     batch_size=batch_size, shuffle=False)
296 # %%
297 lda_proj_mat = None
298 if dim_reducer == "LDA":
299     # Compute S_W and S_B:
300     S_B, S_W = calculate_global_parameters_for_LDA(train_loader, num_classes=num_classes
301         , feature_dim=4096)
302     # Also decides whether we compute YUYANG or not:
303     lda_proj_mat = get_projection_matrix(S_B=S_B, S_W=S_W, p=p, lda_option=lda_option)
304 # %%
305 # Run through test set and find the nearest embedding and assign that label to the image
306 class_embs = train_classifier(train_loader, dim_reducer=dim_reducer, p=p, num_classes=
307     num_classes, lda_proj_mat=lda_proj_mat)
308 # %%
309 test_loader = DataLoader(dataset=DimReducerBuilder(EVAL_DATA_PATH, option="bw"),
310     batch_size=batch_size, shuffle=False)
311 pred_labels, true_labels, test_embs = run_testing_script(test_loader, lda_proj_mat,
312     class_embs, dim_reducer=dim_reducer)
313 accuracy = np.count_nonzero(pred_labels == true_labels) / len(pred_labels)
314 np.round(accuracy * 100, 4)
315 # %% [markdown]
316 # ### Graph the embeddings:
317 # %%
318 test_embs_to_print = []
319 for sample in test_embs:
320     if len(sample) >= 10:

```

```

321         test_embs_to_print.append(sample)
322
323     # %%
324     # Match the shapes of the train and test embeddings:
325     min_num_training_class_samples, min_num_test_class_samples = np.inf, np.inf
326     for train_sample, test_sample in zip(class_embs, test_embs_to_print):
327         if len(train_sample) < min_num_training_class_samples:
328             min_num_training_class_samples = len(train_sample)
329         if len(test_sample) < min_num_test_class_samples:
330             min_num_test_class_samples = len(test_sample)
331
332     print("Minimum training samples per class", min_num_training_class_samples)
333     print("Minimum testing samples per class", min_num_test_class_samples)
334
335     graph_train_embs = np.zeros((num_classes, min_num_training_class_samples, p))
336     graph_test_embs = np.zeros((num_classes, min_num_test_class_samples, p))
337
338     for i, (train_sample, test_sample) in enumerate(zip(class_embs, test_embs_to_print)):
339         test_sample = np.squeeze(test_sample)
340         graph_train_embs[i] = np.array(train_sample[:min_num_training_class_samples], dtype=
341             np.float32)
342         graph_test_embs[i] = np.array(test_sample[:min_num_test_class_samples], dtype=np.
343             float32)
344
345     # Reshape embeddings and generate labels
346     train_embeddings = graph_train_embs.reshape(-1, p)
347     train_labels = np.repeat(np.arange(num_classes), min_num_training_class_samples)
348
349     test_embeddings = graph_test_embs.reshape(-1, p)
350     test_labels = np.repeat(np.arange(num_classes), min_num_test_class_samples)
351
352     # Reduce to 2D with UMAP
353     umap_reducer = umap.UMAP(n_components=2)
354     train_umap = umap_reducer.fit_transform(train_embeddings)
355     test_umap = umap_reducer.transform(test_embeddings)
356
357     # %%
358     # Plot training data with different colors for each class
359     plt.figure(figsize=(8, 6))
360     scatter = plt.scatter(train_umap[:, 0], train_umap[:, 1], c=train_labels, cmap='tab20',
361         alpha=0.7)
362     plt.colorbar(scatter, ticks=range(num_classes), label="Class Label")
363     plt.title("Training Data: UMAP Embeddings")
364     plt.axis("off")
365     plt.savefig(f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/PCA/UMAP/train_umap_{p}
366         }.jpg")
367     plt.close()
368
369     # Plot test data with predicted labels
370     plt.figure(figsize=(8, 6))
371     scatter = plt.scatter(test_umap[:, 0], test_umap[:, 1], c=test_labels, cmap='tab20',
372         alpha=0.7)
373     plt.colorbar(scatter, ticks=range(num_classes), label="Ground Truth Label")
374     plt.title("Test Data: UMAP Embeddings with Predicted Labels")
375     plt.axis("off")
376     plt.savefig(f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/PCA/UMAP/test_umap_{p}.
377         jpg")
378     plt.close()
379
380     # Generate and plot the confusion matrix:
381     cm = confusion_matrix(true_labels, pred_labels)
382
383     # Create a heatmap using Seaborn
384     plt.figure(figsize=(16, 16))
385     sns.heatmap(cm, annot=True, fmt='.0f', cmap='Blues')
386
387     # Add labels and title
388     plt.ylabel('Actual Class')
389     plt.xlabel('Predicted Class')
390     plt.axis("off")
391     plt.savefig(f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/PCA/Confusion_Mat/
392         conf_mat_{p}.jpg")

```



```

387 plt.close()
388
389 # %%
390 # Generate x-axis values (e.g., epoch numbers or iteration indices)
391 x_values = list(range(1, len(accuracy_list) + 1))
392 plt.plot(x_values, accuracy_list, marker='o', linestyle='-', label='PCA')
393 plt.plot(x_values, accuracy_list_lda, marker='x', linestyle='-', label='LDA')
394 plt.legend(loc="upper right")
395 plt.xlabel("Embedding Dim")
396 plt.ylabel("Accuracy (%)")
397 plt.grid(True)
398
399 plt.savefig("/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/Accuracy/PCA.jpg")
400
401 # %%
402 import os
403 import numpy as np
404 import torch
405 from torch import nn, optim
406 import umap
407 from PIL import Image
408 from torch.autograd import Variable
409 from torch.utils.data import Dataset, DataLoader
410 from torchvision import transforms
411 import matplotlib.pyplot as plt
412 from sklearn.metrics import confusion_matrix
413 import seaborn as sns
414
415 # %%
416 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
417 device
418
419 # %% [markdown]
420 # ### Accuracies:
421 #
422 # p = 3: 0%
423 #
424
425 # %%
426 class DataBuilder(Dataset):
427     def __init__(self, path, option):
428         self.path = path
429         self.image_list = [f for f in os.listdir(path) if f.endswith('.png')]
430         self.label_list = [int(f.split('_')[0]) for f in self.image_list]
431         self.len = len(self.image_list)
432         if option == "bw":
433             self.aug = transforms.Compose([
434                 transforms.Resize((64, 64)),
435                 transforms.Grayscale(num_output_channels=1),
436                 transforms.ToTensor(),
437             ])
438         else:
439             self.aug = transforms.Compose([
440                 transforms.Resize((64, 64)),
441                 transforms.ToTensor(),
442             ])
443
444     def __len__(self):
445         return self.len
446
447     def __getitem__(self, index):
448         fn = os.path.join(self.path, self.image_list[index])
449         x = Image.open(fn).convert('RGB')
450         x = self.aug(x)
451
452         return {'x': x, 'y': self.label_list[index]}
453
454 # %%
455 def get_id_of_nearest_embedding(training_set, probe_embedding):
456     # We first calculate the euclidean distance between the probe and all trained
457     embeddings:
458     distances = np.linalg.norm(training_set - probe_embedding, axis=1)

```

```

459
460     # Then we return the index with the smallest distance
461     return np.argmin(distances)
462
463 # %%
464 def train(epoch, vae_loss, model, optimizer, trainloader):
465     model.train()
466     train_loss = 0
467
468     for batch_idx, data in enumerate(trainloader):
469         optimizer.zero_grad()
470         input = data["x"].to(device)
471         mu, logvar = model.encode(input)
472         z = model.reparameterize(mu, logvar)
473         xhat = model.decode(z)
474         loss = vae_loss(xhat, input, mu, logvar)
475         loss.backward()
476         train_loss += loss.item()
477         optimizer.step()
478
479     print('==> Epoch: {} Average loss: {:.4f}'.format(
480         epoch, train_loss / len(trainloader.dataset)))
481     return model
482 class VaeLoss(nn.Module):
483     def __init__(self):
484         super(VaeLoss, self).__init__()
485         self.mse_loss = nn.MSELoss(reduction="sum")
486
487     def forward(self, xhat, x, mu, logvar):
488         loss_MSE = self.mse_loss(xhat, x)
489         loss_KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
490         return loss_MSE + loss_KLD
491
492 # %%
493 class Autoencoder(nn.Module):
494     def __init__(self, encoded_space_dim):
495         super().__init__()
496         self.encoded_space_dim = encoded_space_dim
497         ### Convolutional section
498         self.encoder_cnn = nn.Sequential(
499             nn.Conv2d(3, 8, 3, stride=2, padding=1),
500             nn.LeakyReLU(True),
501             nn.Conv2d(8, 16, 3, stride=2, padding=1),
502             nn.LeakyReLU(True),
503             nn.Conv2d(16, 32, 3, stride=2, padding=1),
504             nn.LeakyReLU(True),
505             nn.Conv2d(32, 64, 3, stride=2, padding=1),
506             nn.LeakyReLU(True)
507         )
508         ### Flatten layer
509         self.flatten = nn.Flatten(start_dim=1)
510         ### Linear section
511         self.encoder_lin = nn.Sequential(
512             nn.Linear(4 * 4 * 64, 128),
513             nn.LeakyReLU(True),
514             nn.Linear(128, encoded_space_dim * 2)
515         )
516         self.decoder_lin = nn.Sequential(
517             nn.Linear(encoded_space_dim, 128),
518             nn.LeakyReLU(True),
519             nn.Linear(128, 4 * 4 * 64),
520             nn.LeakyReLU(True)
521         )
522         self.unflatten = nn.Unflatten(dim=1,
523                                     unflattened_size=(64, 4, 4))
524         self.decoder_conv = nn.Sequential(
525             nn.ConvTranspose2d(64, 32, 3, stride=2,
526                               padding=1, output_padding=1),
527             nn.BatchNorm2d(32),
528             nn.LeakyReLU(True),
529             nn.ConvTranspose2d(32, 16, 3, stride=2,
530                               padding=1, output_padding=1),
531             nn.BatchNorm2d(16),

```

```

532         nn.LeakyReLU(True),
533         nn.ConvTranspose2d(16, 8, 3, stride=2,
534                             padding=1, output_padding=1),
535         nn.BatchNorm2d(8),
536         nn.LeakyReLU(True),
537         nn.ConvTranspose2d(8, 3, 3, stride=2,
538                             padding=1, output_padding=1)
539     )
540
541     def encode(self, x):
542         x = self.encoder_cnn(x)
543         x = self.flatten(x)
544         x = self.encoder_lin(x)
545         mu, logvar = x[:, :self.encoded_space_dim], x[:, self.encoded_space_dim:]
546         return mu, logvar
547
548     def decode(self, z):
549         x = self.decoder_lin(z)
550         x = self.unflatten(x)
551         x = self.decoder_conv(x)
552         x = torch.sigmoid(x)
553         return x
554
555     @staticmethod
556     def reparameterize(mu, logvar):
557         std = logvar.mul(0.5).exp_()
558         eps = Variable(std.data.new(std.size()).normal_())
559         return eps.mul(std).add_(mu)
560
561
562 # %% [markdown]
563 # ### VAE Training:
564
565 # %%
566 #####
567 # Change these
568 p = 8
569 batch_size = 24
570 training = False
571 TRAIN_DATA_PATH = '/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/train/'
572 EVAL_DATA_PATH = '/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/test/'
573 LOAD_PATH = f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/exp/model_{p}.pt"
574 OUT_PATH = '/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/exp/'
575 #####
576
577 # %%
578 train_loader = DataLoader(dataset=DataBuilder(TRAIN_DATA_PATH, option="bw"), batch_size=
    batch_size, shuffle=True)
579
580
581 # %%
582 accuracy_list = []
583
584 # %%
585 p = 16
586 LOAD_PATH = f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/exp/model_{p}.pt"
587 train_loader = DataLoader(dataset=DataBuilder(TRAIN_DATA_PATH, option="bw"), batch_size=
    batch_size, shuffle=True)
588 training=False
589 model = Autoencoder(p).to(device)
590
591 if training:
592     epochs = 100
593     log_interval = 1
594     trainloader = DataLoader(
595         dataset=DataBuilder(TRAIN_DATA_PATH, option=""),
596         batch_size=32,
597     )
598     optimizer = optim.Adam(model.parameters(), lr=1e-3)
599     vae_loss = VaeLoss()
600     for epoch in range(1, epochs + 1):
601         model = train(epoch, vae_loss, model, optimizer, trainloader)
602     torch.save(model.state_dict(), os.path.join(OUT_PATH, f'model_{p}.pt'))

```

```

603 else:
604     trainloader = DataLoader(
605         dataset=DataBuilder(TRAIN_DATA_PATH, option=""),
606         batch_size=1,
607     )
608     model.load_state_dict(torch.load(LOAD_PATH))
609     model.eval()
610
611     X_train, y_train = [], []
612     for batch_idx, data in enumerate(trainloader):
613         mu, logvar = model.encode(data['x'].to(device))
614         z = mu.detach().cpu().numpy().flatten()
615         X_train.append(z)
616         y_train.append(data['y'].item())
617     X_train = np.stack(X_train)
618     y_train = np.array(y_train)
619
620     testloader = DataLoader(
621         dataset=DataBuilder(EVAL_DATA_PATH, option=""),
622         batch_size=1,
623     )
624     X_test, y_test = [], []
625     for batch_idx, data in enumerate(testloader):
626         mu, logvar = model.encode(data['x'].to(device))
627         z = mu.detach().cpu().numpy().flatten()
628         X_test.append(z)
629         y_test.append(data['y'].item())
630     X_test = np.stack(X_test)
631     y_test = np.array(y_test)
632
633
634 # %%
635 train_embs = [[] for _ in range(30)]
636 test_embs = [[] for _ in range(30)]
637
638 for i, (train_emb, train_label, test_emb, test_label) in enumerate(zip(X_train, y_train,
639     X_test, y_test)):
640     train_embs[train_label - 1].append(train_emb)
641     test_embs[test_label - 1].append(test_emb)
642
643 # %%
644 train_embs = np.array(train_embs, dtype=np.float32)
645 test_embs = np.array(test_embs, dtype=np.float32)
646
647 # %%
648 # Now that all the array is reordered, I can add them to the faiss search
649 num_classes, num_samples, embedding_dim = train_embs.shape
650 flattened_train_embs = train_embs.reshape(-1, embedding_dim)
651 print("Flattened", flattened_train_embs.shape)
652
653 # Create a mapping from flattened indices to class IDs
654 index_to_class_id = np.repeat(np.arange(num_classes), num_samples)
655 print("indices", index_to_class_id.shape)
656
657
658
659 # %%
660 true_labels = []
661 pred_labels = []
662 for test_emb, test_label in zip(X_test, y_test):
663     search_emb = np.array(np.expand_dims(test_emb, axis=0), dtype=np.float32)
664     index = get_id_of_nearest_embedding(flattened_train_embs, search_emb)
665     true_labels.append(test_label)
666     pred_labels.append(index_to_class_id[index.item()] + 1)
667 true_labels = np.array(true_labels)
668 pred_labels = np.array(pred_labels)
669 accuracy = np.count_nonzero(pred_labels == true_labels) / len(pred_labels)
670 print("Accuracy: ", np.round(accuracy * 100, 4), "\n")
671 accuracy_list.append(np.round(accuracy * 100, 4))
672
673 # %%
674 # Generate x-axis values (e.g., epoch numbers or iteration indices)

```

```

675 x_values = list(range(1, len(accuracy_list) + 1))
676 plt.plot(x_values, accuracy_list, marker='o', linestyle='--', label='Accuracy')
677 plt.grid(True)
678
679 plt.savefig("/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/Accuracy/Autoencoder.jpg
        ")
680
681 # %%
682 num_classes, num_samples, embedding_dim = train_embs.shape
683 flattened_train_embs = train_embs.reshape(-1, embedding_dim)
684 train_labels = np.repeat(np.arange(num_classes), num_samples)
685
686 num_classes, num_samples, embedding_dim = test_embs.shape
687 flattened_test_embs = test_embs.reshape(-1, embedding_dim)
688 test_labels = np.repeat(np.arange(num_classes), num_samples)
689
690 # %%
691 # Reduce to 2D with UMAP
692 umap_reducer = umap.UMAP(n_components=2)
693 train_umap = umap_reducer.fit_transform(flattened_train_embs)
694 test_umap = umap_reducer.transform(flattened_test_embs)
695
696 # %%
697 # Plot training data with different colors for each class
698 plt.figure(figsize=(8, 6))
699 scatter = plt.scatter(train_umap[:, 0], train_umap[:, 1], c=train_labels, cmap='tab20',
        alpha=0.7)
700 plt.colorbar(scatter, ticks=range(num_classes), label="Class Label")
701 plt.title("Training Data: UMAP Embeddings")
702 plt.axis("off")
703 plt.savefig(f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/AutoEncoder/UMAP/
        train_umap_{p}.jpg")
704 plt.close()
705
706 # Plot test data with predicted labels
707 plt.figure(figsize=(8, 6))
708 scatter = plt.scatter(test_umap[:, 0], test_umap[:, 1], c=test_labels, cmap='tab20',
        alpha=0.7)
709 plt.colorbar(scatter, ticks=range(num_classes), label="Ground Truth Label")
710 plt.title("Test Data: UMAP Embeddings with Predicted Labels")
711 plt.axis("off")
712 plt.savefig(f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/AutoEncoder/UMAP/
        test_umap_{p}.jpg")
713 plt.close()
714
715 # Generate and plot the confusion matrix:
716 cm = confusion_matrix(true_labels, pred_labels)
717
718 # Create a heatmap using Seaborn
719 plt.figure(figsize=(16, 16))
720 sns.heatmap(cm, annot=True, fmt='.0f', cmap='Blues')
721
722 # Add labels and title
723 plt.ylabel('Actual Class')
724 plt.xlabel('Predicted Class')
725 plt.axis("off")
726 plt.savefig(f"/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/Results/AutoEncoder/
        Confusion_Mat/conf_mat_{p}.jpg")
727 plt.close()
728
729 # %%
730
731 # %%
732 import os
733 import numpy as np
734 import torch
735 import torch.nn.functional as F
736 from PIL import Image
737 from torch.utils.data import Dataset, DataLoader
738 from torchvision import transforms
739 import matplotlib.pyplot as plt
740
741 # %%

```

```

742 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
743 device
744
745 # %%
746 def apply_haar_filter(img_bw, haar_size):
747     if haar_size % 2 == 1:
748         # Odd -> add 1
749         # Else this is already the largest even number > 4 sigma
750         haar_size += 1
751     haar_dx = np.vstack((-1*np.ones((haar_size, 1)), np.ones((haar_size, 1)) ))
752     haar_dy = torch.tensor(-1*haar_dx.copy().T)
753     haar_dx = torch.tensor(haar_dx)
754
755     # haar_dx = torch.tensor([[ -1] * haar_size + [1] * haar_size], dtype=torch.float32).
756     # unsqueeze(0).unsqueeze(0)
757     # haar_dy = -haar_dx.transpose(2, 3)
758
759     # dx = cv2.filter2D(img_bw, -1, haar_dx)
760     # dy = cv2.filter2D(img_bw, -1, haar_dy)
761
762     dx = F.conv2d(img_bw, haar_dx, padding='same')
763     dy = F.conv2d(img_bw, haar_dy, padding='same')
764
765     return np.hstack((dx, dy))
766
767 def apply_haar_filter(img_bw, haar_size):
768     haar_dx_np = np.vstack((-1 * np.ones((haar_size, 1)), np.ones((haar_size, 1))))
769     haar_dy_np = -haar_dx_np.T
770
771     # Convert Haar kernels to PyTorch tensors with proper shape
772     haar_dx = torch.tensor(haar_dx_np, dtype=torch.float32).unsqueeze(0).unsqueeze(0) #
773     # Shape: (1, 1, H, W)
774     haar_dy = torch.tensor(haar_dy_np, dtype=torch.float32).unsqueeze(0).unsqueeze(0) #
775     # Shape: (1, 1, H, W)
776
777     # Ensure img_bw has the correct dimensions (batch_size, channels, height, width)
778     if len(img_bw.shape) == 3:
779         img_bw = img_bw.unsqueeze(0) # Add batch dimension if missing
780
781     # Apply Haar filters using F.conv2d
782     dx = F.conv2d(img_bw, haar_dx, padding='same')
783     dy = F.conv2d(img_bw, haar_dy, padding='same')
784     return torch.hstack((dx, dy))
785
786 # %%
787 class DataBuilder(Dataset):
788     def __init__(self, path, option=False):
789         self.path = path
790         self.image_list = [f for f in os.listdir(path) if f.endswith('.png')]
791         self.len = len(self.image_list)
792         self.aug = transforms.Compose([
793             transforms.Resize((64, 64)),
794             transforms.Grayscale(num_output_channels=1),
795             transforms.ToTensor(),
796         ])
797
798     def __len__(self):
799         return self.len
800
801     def __getitem__(self, index):
802         fn = os.path.join(self.path, self.image_list[index])
803         x = Image.open(fn).convert('RGB')
804         x = self.aug(x)
805
806         low_feature_vector = [torch.squeeze(torch.reshape(x, shape=(x.shape[0], -1)))]
807         for haar_size in range(2, x.shape[1], 2):
808             haar_img = apply_haar_filter(x, haar_size)
809             pooled_img = F.avg_pool2d(haar_img, kernel_size=(haar_size, haar_size))
810             pooled_img_flat = torch.squeeze(torch.reshape(pooled_img, shape=(pooled_img.
811 shape[0], -1)))
812             low_feature_vector.append(pooled_img_flat)
813         feature_vecs = torch.hstack(low_feature_vector)

```

```

811         return feature_vecs
812
813
814 # %%
815 pos_train_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/train/positive/"
816 neg_train_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/train/negative/"
817 pos_test_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/test/negative/"
818 neg_test_path = "/mnt/cloudNAS3/Adubois/Classes/ECE661/HW10/test/negative/"
819 num_pos_train = len(os.listdir(pos_train_path))
820 num_neg_train = len(os.listdir(neg_train_path))
821 num_pos_test = len(os.listdir(pos_test_path))
822 num_neg_test = len(os.listdir(neg_test_path))
823
824 # %%
825 pos_train_loader = DataLoader(dataset=DataBuilder(pos_train_path), batch_size=
    num_pos_train, shuffle=True)
826 neg_train_loader = DataLoader(dataset=DataBuilder(neg_train_path), batch_size=
    num_neg_train, shuffle=True)
827 pos_test_loader = DataLoader(dataset=DataBuilder(pos_test_path), batch_size=num_pos_test
    , shuffle=True)
828 neg_test_loader = DataLoader(dataset=DataBuilder(neg_test_path), batch_size=num_neg_test
    , shuffle=True)
829
830 # %%
831 for pos_train, neg_train, pos_test, neg_test in zip(pos_train_loader, neg_train_loader,
    pos_test_loader, neg_test_loader):
832     print("Pos Train: ", pos_train.shape)
833     print("Neg Train: ", neg_train.shape)
834     # Create a matrix of all training images, positive or negative:
835     train_imgs = np.vstack((pos_train, neg_train))
836     train_labels = np.hstack(( np.ones(pos_train.shape[0]), -1*np.ones(neg_train.shape
    [0])) )
837
838     # Randomize the order of the training images. Maintain a constant pairing of label
    to img though
839     shuffle_indices = np.random.permutation(train_imgs.shape[0])
840     train_imgs = train_imgs[shuffle_indices]
841     train_labels = train_labels[shuffle_indices]
842
843     # Do the same for the testing images:
844     test_imgs = np.vstack((pos_test, neg_test))
845     test_labels = np.hstack(( np.ones(pos_test.shape[0]), -1*np.ones(neg_test.shape[0])
    ))
846     shuffle_indices = np.random.permutation(test_imgs.shape[0])
847     test_imgs = test_imgs[shuffle_indices]
848     test_labels = test_labels[shuffle_indices]
849
850     print("Train imgs", train_imgs.shape)
851     print("Test imgs", test_imgs.shape)
852
853 # %%
854 class WeakClassifier():
855     def __init__(self):
856         # These are the terms we need to calculate for the weak-classifier:
857         self.best_feature = None
858         self.best_threshold = None
859         self.best_polarity = None
860         self.min_error = np.inf
861
862     def get_params(self, imgs, labels, weight_mat):
863         # Normalize each weight matrix:
864         weight_mat = weight_mat / np.sum(weight_mat)
865
866         # We need to loop through each feature in the img matrix. Each feature is
    counted as a column in that matrix:
867         for feature_idx in range(imgs.shape[1]):
868             # Extract current feature (the column)
869             features = imgs[:, feature_idx]
870
871             # Sort the feature, weights, and labels
872             sorted_indices = np.argsort(features)
873             sorted_features = features[sorted_indices]
874             sorted_weights = weight_mat[sorted_indices]

```

```

875         sorted_labels = labels[sorted_indices]
876
877         # We need to use multiplication here to preserve the original shape
878         # However, we don't want the opposite labels to affect the cumulative summ
879         S_plus = np.cumsum(sorted_weights * (sorted_labels == 1))
880         S_minus = np.cumsum(sorted_weights * (sorted_labels == -1))
881         T_plus = np.sum(sorted_weights * (sorted_labels == 1))
882         T_minus = np.sum(sorted_weights * (sorted_labels == -1))
883
884         # Calculate the polarity errors:
885         e_1 = S_plus + T_minus - S_minus
886         e_neg1 = S_minus + T_plus - S_plus
887
888         # Calculate classification error:
889         for i, feature in enumerate(sorted_features):
890             # Since Error = min(e_1, e_neg1) we will always compute both and keep
the trailing
891             # minimum along both calculations
892             if e_1[i] < self.min_error:
893                 self.min_error = e_1[i]
894                 self.best_feature = feature_idx
895                 self.best_threshold = feature
896                 self.best_polarity = 1
897             if e_neg1[i] < self.min_error:
898                 self.min_error = e_neg1[i]
899                 self.best_feature = feature_idx
900                 self.best_threshold = feature
901                 self.best_polarity = -1
902         return (self.best_feature, self.best_threshold, self.best_polarity, self.
min_error)
903
904
905 # %%
906 class ClassifierCascade():
907     def __init__(self):
908         self.classifier_list = []
909         self.alpha_list = []
910         self.cascades = []
911         self.max_num_classifiers_per_cascade = 5
912         self.max_cascades = 3
913
914     def run_strong_classifier(self, imgs, labels):
915         # I associate a uniform initial weight with each image initially:
916         weight_mat = np.ones(imgs.shape[0], dtype=np.float32) / imgs.shape[0]
917
918         # Define the cascade:
919         for classifier_idx in range(self.max_num_classifiers_per_cascade):
920             # Every new iteration, we add in a new weak classifier until we have reached
the maximum number, or they have the correct accuracy.
921             self.classifier_list.append(WeakClassifier().get_params(imgs, labels,
weight_mat))
922             feature, threshold, polarity, error = self.classifier_list[classifier_idx]
923
924             # Update algorithm parameters
925             beta = error / (1 - error)
926             alpha = np.log(1 / beta)
927             self.alpha_list.append(alpha)
928
929             print(f"Weak classifier id {(classifier_idx + 1):.3f} feature: {feature:.3f
}, threshold: {threshold:.3f}, polarity: {polarity:.3f}, error: {error:.3f}, alpha:
{alpha:.3f}")
930
931             # Update weights accordingly:
932             # You need to find your predictions (feature vs threshold feature value)
933             # However, also make sure to take into account the polarity
934             pred_labels = np.where((imgs[:, feature] * polarity) >= (threshold *
polarity), 1, -1)
935             wrong_preds = pred_labels != labels
936
937             # Multiply by 1 where predictions are incorrect, else by beta.
938             # Also normalize weights to make sure they sum to 1
939             weight_mat = weight_mat * np.where(wrong_preds, 1, beta)
940             weight_mat /= np.sum(weight_mat)

```



```

941         # Calculate the accuracy of all weak classifiers so far:
942         # To do so we first need to get the predictions by passing the input through
943         all the weak classifiers
944         # We can then get the sign of this predictions to assign it to a class label
945         (1 or -1)
946         predictions_so_far = np.zeros_like(labels)
947         # Get the predictions for each classifier, alpha is a weight factor for how
948         much each classifier contributes
949         for (f, th, p, _), alpha in zip(self.classifier_list, self.alpha_list):
950             predictions_so_far += alpha * np.where((imgs[:, f] * p) >= (th * p), 1,
951             -1)
952         final_pred_labels = np.sign(predictions_so_far)
953
954         # Calculate accuracy
955         accuracy = np.count_nonzero(final_pred_labels == labels) / len(
956         final_pred_labels)
957         print(f"Iteration {classifier_idx + 1}: Accuracy = ", np.round(accuracy, 3))
958
959         # Evaluate if there are enough classifiers:
960         if accuracy > 1.00:
961             break
962         return self.classifier_list, self.alpha_list
963
964     def run_cascades(self, imgs, labels):
965         # This function will run multiple cascades until the false positive rate reaches
966         0
967         for cascade_idx in range(self.max_cascades):
968             # We first train a strong classifier
969             classifier_params, alphas = ClassifierCascade().run_strong_classifier(imgs,
970             labels)
971             self.cascades.append((classifier_params, alphas))
972
973             # We then need to evaluate the most recent cascade as we did before
974             predictions_so_far = np.zeros_like(labels)
975             for (f, th, p, _), alpha in zip(classifier_params, alphas):
976                 predictions_so_far += alpha * np.where((imgs[:, f] * p) >= (th * p), 1,
977                 -1)
978             cascade_pred_labels = np.sign(predictions_so_far)
979
980             # false_positives = np.mean((cascade_pred_labels == 1) & (labels == -1))
981             false_positives = np.count_nonzero((cascade_pred_labels == 1) & (labels ==
982             -1)) / len(cascade_pred_labels)
983             false_negatives = np.count_nonzero((cascade_pred_labels == -1) & (labels ==
984             1)) / len(cascade_pred_labels)
985             accuracy = np.count_nonzero(cascade_pred_labels == labels) / len(
986             cascade_pred_labels)
987             print(f"Cascade id: {cascade_idx + 1}, False positive rate: {false_positives
988             :.2f}, False negative rate: {false_negatives:.2f} Final Accuracy: {accuracy:.2f}")
989
990             # Now that we know the false positive rate, we can either terminate, or keep
991             going by removing correctly labeled negatives:
992             if false_positives + false_negatives < 0.01:
993                 break
994             # Remove all the correctly classified negative images from the dataset
995             idx_to_keep = (labels == 1) | ((cascade_pred_labels == 1) & (labels == -1))
996             imgs = imgs[idx_to_keep]
997             labels = labels[idx_to_keep]
998
999             # We also need to stop if there are no more imgs left (not removing any)
1000             if len(idx_to_keep) == len((labels == 1)):
1001                 break
1002         return self.cascades
1003
1004     def test_cascade(self, imgs, labels):
1005         total_num_images = imgs.shape[0]
1006         final_true_negatives = np.zeros_like(labels)
1007         final_true_positives = np.zeros_like(labels)
1008
1009         # Get the accuracy:
1010         for cascade_idx, (classifier_params, alphas) in enumerate(self.cascades):
1011             # Get the predictions on the test dataset:

```

```

1001     predictions_so_far = np.zeros_like(labels)
1002     for (f, th, p, _), alpha in zip(classifier_params, alphas):
1003         predictions_so_far += alpha * np.where((imgs[:, f] * p) >= (th * p), 1,
1004         -1)
1005         cascade_pred_labels = np.sign(predictions_so_far)
1006         # Compute performance metrics for the current cascade
1007         true_negative_mask = (cascade_pred_labels == -1) & (labels == -1)
1008         true_positive_mask = (cascade_pred_labels == 1) & (labels == 1)
1009         final_true_negatives = np.logical_or(final_true_negatives,
1010         true_negative_mask)
1011         final_true_positives = np.logical_or(final_true_positives,
1012         true_positive_mask)
1013         false_positives = np.count_nonzero((cascade_pred_labels == 1) & (labels ==
1014         -1))
1015         false_negatives = np.count_nonzero((cascade_pred_labels == -1) & (labels ==
1016         1))
1017         true_negatives = np.count_nonzero(true_negative_mask)
1018         true_positives = np.count_nonzero(true_positive_mask)
1019         tot_negative = np.count_nonzero(labels == -1)
1020         tot_positive = np.count_nonzero(labels == 1)
1021
1022         print(f"True Positives: {true_positives}, True Negatives: {true_negatives},
1023         Total Image Count {total_num_images}")
1024         false_positive_rate = false_positives / (tot_negative)
1025         false_negative_rate = false_negatives / (tot_positive)
1026         accuracy = (true_positives + true_negatives) / len(cascade_pred_labels)
1027         percent_of_dataset_kept = len(cascade_pred_labels) / total_num_images
1028         print(f"Cascade id: {cascade_idx + 1}, False positive rate: {
1029         false_positive_rate:.2f}, False negative rate: {false_negative_rate:.2f}, Accuracy:
1030         {np.round(accuracy*100, 4)}, over {np.round(percent_of_dataset_kept*100, 4)} of the
1031         imgs.")
1032
1033         # Print final accuracy:
1034         final_tp = np.count_nonzero(final_true_positives)
1035         final_tn = np.count_nonzero(final_true_negatives)
1036         final_accuracy = (final_tp + final_tn) / total_num_images
1037         print(f"Final Accuracy: {np.round(final_accuracy*100, 4)}")
1038
1039 # %%
1040 classifier_cascade = ClassifierCascade()
1041 cascade_params = classifier_cascade.run_cascades(train_imgs.copy(), train_labels.copy())
1042
1043 # %%
1044 classifier_cascade.test_cascade(test_imgs.copy(), test_labels.copy())
1045
1046 # %%

```